

Introduction to Computer Systems

Homework

Vincent Lee

ihalbmond@gmail.com

2017年9月16日

目录

1 A Tour of Computer Systems	6
2 Representing and Manipulating Information	7
2.53	7
2.60	7
2.61	7
2.65	8
2.86	9
2.95	9
2.97	10
3 Machine-Level Representation of Programs	11
3.58	11
3.59	11
3.60	11
3.61	11
3.62	12
3.63	12
3.64	12
3.65	12
3.66	12
3.67	12
3.68	12
3.69	12
4 Processor Architecture	13
4.43	13
4.44	13
4.45	13
4.46	18
4.47	20
4.48	20
4.49	20

目录	3
4.50	20
4.51	31
4.52	31
4.53	31
4.54	32
4.55	33
4.56	33
4.57	33
4.58	35
5 Optimizing Program Performance	36
5.13	36
5.14	37
6 The Memory Hierarchy	38
6.23	38
7 Linking	39
7.8	39
7.9	39
7.12	39
7.13	40
8 Exceptional Control Flow	41
8.9	41
8.10	41
8.11	41
8.12	42
8.13	42
8.14	42
8.15	42
8.16	42
8.17	42
8.18	42
8.19	43

目录	4
8.20	44
8.21	44
8.22	45
8.25	46
8.26	47
9 Virtual Memory	48
9.11	48
9.12	49
9.13	50
9.14	51
9.15	52
9.16	52
9.17	53
9.18	56
9.19	58
9.20	59
10 System-Level I/O	60
10.6	60
10.7	60
10.8	61
10.9	62
10.10	63
11 Network Programming	64
11.6	64
11.8	65
12 Concurrent Programming	66
12.16	66
12.17	67
12.18	68
12.22	68
12.23	69

目录	5
12.24	69
12.25	70
12.26	71

1 A Tour of Computer Systems

2 Representing and Manipulating Information

2.53

```
1 #define POS_INFINITY 1e400
2 #define NEG_INFINITY (-POS_INFINITY)
3 #define NEG_ZERO (-1.0/POS_INFINITY)
```

2.60

suppose we number the bytes in a w-bit word from 0 (less significant) to w/8-1 (most significant). write code for the followign c function, which will return an unsigned value in which byte i of argument x has been replaced by byte b:

```
1 unsigned replace_byte (unsigned x, int i, unsigned char b)
2 {
3     unsigned char *a = (unsigned char*) (&x);
4     a[i] = b;
5     return x;
6 }
```

2.61

Write C expressions that evaluate to 1 when the following conditions are true, and to 0 when they are false. Assume x is of type int.

- (A) Any bit of x equals 1.
- (B) Any bit of x equals 0.
- (C) Any bit in the least significant byte of x equals 1.
- (D) Any bit in the most significant byte of x equals 0.

Your code should follow the bit-level integer coding rules (page 120), with the additional restriction that you may not use equality (==) or inequality (!=) tests.

```

1  !!x
2  !!(~x)
3  !(x << ( (sizeof(int) - 1) << 3 ) )
4  !(~(x >> ( (sizeof(int) - 1) << 3 ) ) )

```

2.65

Write code to implement the following function:

```

1  /* Return 1 when x contains an odd number of 1s; 0 otherwise. Assume w=32. */
2  int odd_ones(unsigned x);

```

Your function should follow the bit-level integer coding rules (page 120), except that you may assume that data type `int` has $w = 32$ bits. Your code should contain a total of at most 12 arithmetic, bit-wise, and logical operations.

```

1  /*
2      using divide and conquer algorithm
3  */
4  int odd_ones(unsigned x)
5  {
6      x ^= x >> 16;
7      x ^= x >> 8;
8      x ^= x >> 4;
9      return 0x69966696 >> x & 1;
10 }

```

2.86

Intel-compatible processors also support an “extended precision” floating-point format with an 80-bit word divided into a sign bit, $k = 15$ exponent bits, a single integer bit, and $n = 63$ fraction bits. The integer bit is an explicit copy of the implied bit in the IEEE floating-point representation. That is, it equals 1 for normalized values and 0 for denormalized values. Fill in the following table giving the approximate values of some “interesting” numbers in this format:

Description	Extended precision	
	value	Decimal
Smallest positive denormalized	2^{-16445}	3.6452×10^{-4951}
Smallest positive normalized	2^{-16382}	3.3621×10^{-4932}
Largest normalized	$(2 - 2^{-63}) \times 2^{16383}$	1.1897×10^{4932}

2.95

Following the bit-level floating-point coding rules, implement the function with the following prototype:

```
1 /* Compute 0.5*f. If f is NaN, then return f. */
2 float_bits float_half(float_bits f);
```

For floating-point number f , this function computes $0.5 * f$. If f is NaN, your function should simply return f . Test your function by evaluating it for all 2^{32} values of argument f and comparing the result to what would be obtained using your machine’s floating-point operations.

```
1 float_bits float_half(unsigned uf) {
2     unsigned char exp_bit = uf >> 23; // unsigned int to unsigned char, overflow
3     if (exp_bit == 0xff) // f is NaN
4         return uf;
5     if (exp_bit <= 1) // exp_bit == 0 or exp_bit == 1
6         return (uf >> 31 << 30) + (uf >> 1) + (uf & (uf >> 1) & 1); // round to even
7     return uf - 0x800000;
8 }
```

2.97

Following the bit-level floating-point coding rules, implement the function with

```
1 the following prototype:
2 /* Compute (float) i */
3 float_bits float_i2f(int i);
```

For argument *i*, this function computes the bit-level representation of (float) *i*. Test your function by evaluating it for all 2^{32} values of argument *i* and comparing the result to what would be obtained using your machine's floating-point operations.

```
1 float_bits float_i2f(int x) {
2     int sign_bit = x & 0x80000000;
3     int exp_bit = 0x4f800000; // upper bound of exp
4     int carry = 0;
5     if (sign_bit) x = -x;
6     if (x == 0) return 0;
7     while (x > 0)
8     {
9         exp_bit -= 0x800000; // 2^23
10        x <<= 1;
11    }
12    carry = ((x & 0x80) && (x & 0x17f)); // round to even
13    return sign_bit + exp_bit + (x >> 8) + carry;
14 }
```

3 Machine-Level Representation of Programs

3.58

```
1 long decode2(long x, long y, long z)
2 {
3     y -= z;
4     x *= y;
5     long rax = y;
6     rax = rax << 63 >> 63;
7     rax ^= x;
8     return rax;
9 }
```

3.59

3.60

(A) x in %rdi, n in %esi, result in %rax, mask in %rdx.

(B) $result = 0$, $mask = 1$

(C) $mask \neq 0$

(D) $mask = mask \ll n$

(E) $result = x \& mask$

(F)

```
long loop(long x, int n)
{
    long result = 0;
    long mask;
    for(mask = 1; mask != 0; mask = mask << n) {
        result |= x & mask;
    }
    return result;
}
```

3.61

```
1 long cread_alt(long *xp) {  
2     long t = 0;  
3     return *(xp ? xp : &t);  
4 }
```

3.62

3.63

3.64

$$R = 7, S = 5, T = 13$$

3.65

3.66

3.67

3.68

$$A = 9, B = 5$$

3.69

$$CNT = 7$$

```
1 typedef struct {  
2     int first;  
3     a_struct a[CNT];  
4     int last;  
5 } b_struct;  
6  
7 typedef struct {  
8     long idx;  
9     long x[4];  
10 } a_struct;
```

4 Processor Architecture

4.43

(A) In light of analysis done in Problem 4.6, does this code sequence correctly describe the behavior of the instruction `pushl %esp`? Explain.

(a) This code pushes `%esp - 4` into stack, while `"pushl %esp"` pushes `%esp` into stack.

(b) This code would change conditional code.

(B) How could you rewrite the code sequence so that it correctly describes both the cases where `REG` is `%esp` as well as any other register?

```
1    movl    REG,    -4(%esp)
2    leal    -4(%esp), %esp
```

4.44

(A) In light of analysis done in Problem 4.7, does this code sequence correctly describe the behavior of the instruction `popl %esp`? Explain.

(a) This code changes `%esp` to `(%esp) + 4`, while `"popl %esp"` changes `%esp` to `(%esp)`

(b) This code would change conditional code.

(B) How could you rewrite the code sequence so that it correctly describes both the cases where `REG` is `%esp` as well as any other register?

```
1    leal    4(%esp), %esp
2    movl    -4(%esp), REG
```

4.45

(A) Write and test a C version that references the array elements with pointers, rather than using array indexing.

```

1 void bubble_a(int *data, int count) {
2     int i, last;
3     for (last = count-1; last > 0; last--) {
4         int *p_last = data + last;
5         for (int *p1 = data; p1 != p_last; ++p1) {
6             p2 = p1 + 1;                                /* compare two adjacent pointers */
7             if (*p2 < *p1) {
8                 int temp = *p2; *p2 = *p1; *p1 = temp;
9             }
10        }
11    }
12 }
13
14 /*
15     Test Code Part
16 */
17 int a[] = {7, 6, 8, 3, 0};
18
19 int main() {
20     int i;
21     bubble_a(a, 5);
22     for(i = 0; i < 4; i++)
23         if(a[i] > a[i+1])
24             return 1;
25     return 0;
26 }

```

- (B) Write and test a Y86 program consisting of the function and test code. You may find it useful to pattern your implementation after IA32 code generated by compiling your C code. Although pointer comparisons are normally done using unsigned arithmetic, you can use signed arithmetic for this exercise.

```

1     .pos 0
2 init:
3     irmovl Stack, %esp
4     irmovl Stack, %ebp
5     call Main
6     halt
7
8 bubble_a:
9     pushl    %edi
10    pushl    %esi
11    pushl    %ebx
12    mrmovl   20(%esp), %edi

```

```

13      irmovl $1, %eax
14      rrmovl %edi, %esi
15      subl %eax, %esi
16      mrmovl 16(%esp), %eax
17      andl %esi, %esi
18      jle L1
19      rrmovl %eax, %edi
20      rrmovl %esi, %ebx
21      addl %ebx, %ebx
22      addl %ebx, %ebx
23      addl %eax, %ebx
24      jmp L3
25 L7:
26      mrmovl 4(%eax), %edx
27      mrmovl (%eax), %ecx
28      pushl %edx
29      subl %ecx, %edx
30      popl %edx
31      jge L4
32      rmmovl %ecx, 4(%eax)
33      rmmovl %edx, (%eax)
34 L4:
35      pushl %ebx
36      irmovl $4, %ebx
37      addl %ebx, %eax
38      popl %ebx
39      pushl %eax
40      subl %ebx, %eax
41      popl %eax
42      jne L7
43 L6:
44      pushl %eax
45      irmovl $4, %eax
46      subl %eax, %ebx
47      irmovl $1, %eax
48      subl %eax, %esi
49      popl %eax
50      je L1
51 L3:
52      andl %esi, %esi
53      jle L6
54      rrmovl %edi, %eax
55      jmp L7
56 L1:
57      popl %ebx
58      popl %esi

```

```

59     popl    %edi
60     ret
61
62 Main:
63     irmovl  $8, %eax
64     subl    %eax, %esp
65     irmovl  $5, %eax
66     rmmovl  %eax, 4(%esp)
67     irmovl  $a, %eax
68     rmmovl  %eax, (%esp)
69     call    bubble_a
70     irmovl  $8, %eax
71     addl    %eax, %esp
72     xorl    %eax, %eax
73     ret
74
75     .align 4
76 a:
77     .long   7
78     .long   6
79     .long   8
80     .long   3
81     .long   0
82
83     .pos 0x1000
84 Stack:

```

This is the output in terminal:

```

1 Stopped in 235 steps at PC = 0x11.  Status 'HLT', CC Z=1 S=0 O=0
2
3 Changes to registers:
4 %ecx:  0x00000000  0x00000001
5 %esp:  0x00000000  0x00001000
6 %ebp:  0x00000000  0x00001000
7
8 Changes to memory:
9 0x00dc: 0x00000007  0x00000000
10 0x00e0: 0x00000006  0x00000003
11 0x00e4: 0x00000008  0x00000006
12 0x00e8: 0x00000003  0x00000007
13 0x00ec: 0x00000000  0x00000008
14 0x0fe0: 0x00000000  0x000000e0
15 0x0ff0: 0x00000000  0x000000d1
16 0x0ff4: 0x00000000  0x000000dc
17 0x0ff8: 0x00000000  0x00000005
18 0x0ffc: 0x00000000  0x00000011

```

Look at the line 9 to line 13, the array was sorted.

4.46

```

1      .pos 0
2  init:
3      irmovl Stack, %esp
4      irmovl Stack, %ebp
5      call Main
6      halt
7
8  bubble_a:
9      pushl  %edi
10     pushl  %esi
11     pushl  %ebx
12     mrmovl 20(%esp), %edi
13     irmovl $1, %eax
14     rrmovl %edi, %esi
15     subl   %eax, %esi
16     mrmovl 16(%esp), %eax
17     andl   %esi, %esi
18     jle L1
19     rrmovl %eax, %edi
20     rrmovl %esi, %ebx
21     addl   %ebx, %ebx
22     addl   %ebx, %ebx
23     addl   %eax, %ebx
24     jmp L3
25 L7:
26     mrmovl 4(%eax), %edx
27     mrmovl (%eax), %ecx
28     pushl  %ebx
29     pushl  %edx
30     subl   %ecx, %edx
31     popl   %edx
32     cmovl  %ebx, %edx
33     rmmovl %edx, 4(%eax)
34     rmmovl %ecx, (%eax)
35     irmovl $4, %ebx
36     addl   %ebx, %eax
37     popl   %ebx
38     pushl  %eax
39     subl   %ebx, %eax
40     popl   %eax
41     jne L7
42 L6:
43     pushl  %eax
44     irmovl $4, %eax
45     subl   %eax, %ebx

```

```
46      irmovl $1, %eax
47      subl  %eax, %esi
48      popl  %eax
49      je    L1
50 L3:
51      andl  %esi, %esi
52      jle   L6
53      rrmovl %edi, %eax
54      jmp   L7
55 L1:
56      popl  %ebx
57      popl  %esi
58      popl  %edi
59      ret
60
61
62 Main:
63      irmovl $8, %eax
64      subl  %eax, %esp
65      irmovl $5, %eax
66      rmmovl %eax, 4(%esp)
67      irmovl $a, %eax
68      rmmovl %eax, (%esp)
69      call  bubble_a
70      irmovl $8, %eax
71      addl  %eax, %esp
72      xorl  %eax, %eax
73      ret
74
75      .align 4
76 a:
77      .long 7
78      .long 6
79      .long 8
80      .long 3
81      .long 0
82
83      .pos 0x1000
84 Stack:
```

4.47

See the table below for detail.

4.48

See the table below for detail.

Stage	Instruction	
	<i>iaddl V, rB</i>	<i>leave</i>
Fetch	$icode : ifun \leftarrow M_1[PC]$ $rA : rB \leftarrow M_1[PC + 1]$ $ValC \leftarrow M_4[PC + 2]$ $ValP \leftarrow PC + 6$	$icode : ifun \leftarrow M_1[PC]$ $ValP \leftarrow PC + 1$
Decode	$ValB \leftarrow R[rB]$	$ValA \leftarrow R[\%ebp]$ $ValB \leftarrow R[\%ebp]$
Execute	$ValE \leftarrow ValB + ValC$ $SetCC$	$ValE \leftarrow ValB + 4$
Memory		$ValM \leftarrow M_4[ValA]$
Write-back	$R[rB] \leftarrow ValE$	$R[\%esp] \leftarrow ValE$ $R[\%ebp] \leftarrow ValM$
PC-update	$PC \leftarrow ValP$	$PC \leftarrow ValP$

4.49

See *archlab partC* for detail.

4.50

```

1 #!/* $begin seq-all-hcl */
2
3 #####
4
5 #   HCL Description of Control for Single Cycle Y86 Processor SEQ   #
6
7 #   Copyright (C) Randal E. Bryant, David R. O'Hallaron, 2010      #
8

```

```

9 #####
10
11
12
13 ## Your task is to implement the iaddl and leave instructions
14
15 ## The file contains a declaration of the icodes
16
17 ## for iaddl (IIADDL) and leave (ILEAVE).
18
19 ## Your job is to add the rest of the logic to make it work
20
21
22
23 #####
24
25 #      C Include's.  Don't alter these                                #
26
27 #####
28
29
30
31 quote '#include <stdio.h>'
32
33 quote '#include "isa.h"'
34
35 quote '#include "sim.h"'
36
37 quote 'int sim_main(int argc, char *argv[]);'
38
39 quote 'int gen_pc(){return 0;}'
40
41 quote 'int main(int argc, char *argv[])'
42
43 quote '    {plusmode=0;return sim_main(argc,argv);}'
44
45
46
47 #####
48
49 #      Declarations.  Do not change/remove/delete any of these      #
50
51 #####
52
53
54

```

```

55 ##### Symbolic representation of Y86 Instruction Codes #####
56
57 intsig INOP      'I_NOP'
58
59 intsig IHALT     'I_HALT'
60
61 intsig IRRMOVL   'I_RRMOVL'
62
63 intsig IIRMOVL   'I_IRMOVL'
64
65 intsig IRMMOVL   'I_RMMOVL'
66
67 intsig IMRMOVL   'I_MRMOVL'
68
69 intsig IOPL      'I_ALU'
70
71 intsig IJXX      'I_JMP'
72
73 intsig ICALL      'I_CALL'
74
75 intsig IRET       'I_RET'
76
77 intsig IPUSHL     'I_PUSHL'
78
79 intsig IPOPL      'I_POPL'
80
81 # Instruction code for iaddl instruction
82
83 intsig IIADDL     'I_IADDL'
84
85 # Instruction code for leave instruction
86
87 intsig ILEAVE     'I_LEAVE'
88
89
90
91 ##### Symbolic representations of Y86 function codes #####
92
93 intsig FNONE      'F_NONE'          # Default function code
94
95
96
97 ##### Symbolic representation of Y86 Registers referenced explicitly #####
98
99 intsig RESP       'REG_ESP'          # Stack Pointer
100

```

```

101 intsig REBP      'REG_EBP'      # Frame Pointer
102
103 intsig RNONE     'REG_NONE'     # Special value indicating "no register"
104
105
106
107 ##### ALU Functions referenced explicitly #####
108
109 intsig ALUADD     'A_ADD'        # ALU should add its arguments
110
111
112
113 ##### Possible instruction status values #####
114
115 intsig SAOK 'STAT_AOK'          # Normal execution
116
117 intsig SADR 'STAT_ADR'          # Invalid memory address
118
119 intsig SINS 'STAT_INS'          # Invalid instruction
120
121 intsig SHLT 'STAT_HLT'          # Halt instruction encountered
122
123
124
125 ##### Signals that can be referenced by control logic #####
126
127
128
129 ##### Fetch stage inputs #####
130
131 intsig pc 'pc'                  # Program counter
132
133 ##### Fetch stage computations #####
134
135 intsig imem_icode 'imem_icode'   # icode field from instruction memory
136
137 intsig imem_ifun  'imem_ifun'    # ifun field from instruction memory
138
139 intsig icode      'icode'         # Instruction control code
140
141 intsig ifun       'ifun'          # Instruction function
142
143 intsig rA         'ra'            # rA field from instruction
144
145 intsig rB         'rb'            # rB field from instruction
146

```

```

147 intsig valC 'valc'          # Constant from instruction
148
149 intsig valP 'valp'          # Address of following instruction
150
151 boolsig imem_error 'imem_error' # Error signal from instruction memory
152
153 boolsig instr_valid 'instr_valid' # Is fetched instruction valid?
154
155
156
157 ##### Decode stage computations #####
158
159 intsig valA 'vala'          # Value from register A port
160
161 intsig valB 'valb'          # Value from register B port
162
163
164
165 ##### Execute stage computations #####
166
167 intsig vale 'vale'          # Value computed by ALU
168
169 boolsig Cnd 'cond'          # Branch test
170
171
172
173 ##### Memory stage computations #####
174
175 intsig valM 'valm'          # Value read from memory
176
177 boolsig dmem_error 'dmem_error' # Error signal from data memory
178
179
180
181
182
183 #####
184
185 # Control Signal Definitions. #
186
187 #####
188
189
190
191 ##### Fetch Stage #####
192

```



```

193
194
195 # Determine instruction code
196
197 int icode = [
198
199     imem_error: INOP;
200
201     1: imem_icode;      # Default: get from instruction memory
202
203 ];
204
205
206
207 # Determine instruction function
208
209 int ifun = [
210
211     imem_error: FNONE;
212
213     1: imem_ifun;      # Default: get from instruction memory
214
215 ];
216
217
218
219 bool instr_valid = icode in
220
221     { INOP, IHALT, IRRMOVL, IIRMOVL, IRMMOVL, IMRMOVL,
222
223         IOPL, IJXX, ICALL, IRET, IPUSHL, IPOPL, IIADDL };
224
225
226
227 # Does fetched instruction require a regid byte?
228
229 bool need_regids =
230
231     icode in { IRRMOVL, IOPL, IPUSHL, IPOPL,
232
233         IIRMOVL, IRMMOVL, IMRMOVL, IIADDL };
234
235
236
237 # Does fetched instruction require a constant word?
238

```

```

239 bool need_valC =
240
241     icode in { IIRMOVL, IRMMOVL, IMRMOVL, IJXX, ICALL, IIADDL };
242
243
244
245 ##### Decode Stage #####
246
247
248
249 ## What register should be used as the A source?
250
251 int srcA = [
252
253     icode in { IRRMOVL, IRMMOVL, IOPL, IPUSHL } : rA;
254
255     icode in { IPOPL, IRET } : RESP;
256
257     1 : RNONE; # Don't need register
258
259 ];
260
261
262
263 ## What register should be used as the B source?
264
265 int srcB = [
266
267     icode in { IOPL, IRMMOVL, IMRMOVL, IIADDL } : rB;
268
269     icode in { IPUSHL, IPOPL, ICALL, IRET } : RESP;
270
271     1 : RNONE; # Don't need register
272
273 ];
274
275
276
277 ## What register should be used as the E destination?
278
279 int dstE = [
280
281     icode in { IRRMOVL } && Cnd : rB;
282
283     icode in { IIRMOVL, IOPL, IIADDL } : rB;
284

```

```

285     icode in { IPUSHL, IPOPL, ICALL, IRET } : RESP;
286
287     1 : RNONE; # Don't write any register
288
289 ];
290
291
292
293 ## What register should be used as the M destination?
294
295 int dstM = [
296
297     icode in { IMRMOVL, IPOPL } : rA;
298
299     1 : RNONE; # Don't write any register
300
301 ];
302
303
304
305 ##### Execute Stage #####
306
307
308
309 ## Select input A to ALU
310
311 int aluA = [
312
313     icode in { IRRMOVL, IOPL } : valA;
314
315     icode in { IIRMOVL, IRMMOVL, IMRMOVL, IIADDL } : valC;
316
317     icode in { ICALL, IPUSHL } : -4;
318
319     icode in { IRET, IPOPL } : 4;
320
321     # Other instructions don't need ALU
322
323 ];
324
325
326
327 ## Select input B to ALU
328
329 int aluB = [
330

```

```

331     icode in { IRMMOVL, IMRMOVL, IOPL, ICALL,
332
333             IPUSHL, IRET, IPOPL, IIADDL } : valB;
334
335     icode in { IRRMOVL, IIRMOVL } : 0;
336
337     # Other instructions don't need ALU
338
339 ];
340
341
342
343 ## Set the ALU function
344
345 int alufun = [
346
347     icode == IOPL : ifun;
348
349     1 : ALUADD;
350
351 ];
352
353
354
355 ## Should the condition codes be updated?
356
357 bool set_cc = icode in { IOPL, IIADDL };
358
359
360
361 ##### Memory Stage #####
362
363
364
365 ## Set read control signal
366
367 bool mem_read = icode in { IMRMOVL, IPOPL, IRET };
368
369
370
371 ## Set write control signal
372
373 bool mem_write = icode in { IRMMOVL, IPUSHL, ICALL };
374
375
376

```

```

377 ## Select memory address
378
379 int mem_addr = [
380     icode in { IRMMOVL, IPUSHL, ICALL, IMRMOVL } : valE;
381     icode in { IPOPL, IRET } : valA;
382     # Other instructions don't need address
383 ];
384
385 ## Select memory input data
386
387 int mem_data = [
388     # Value from register
389     icode in { IRMMOVL, IPUSHL } : valA;
390     # Return PC
391     icode == ICALL : valP;
392     # Default: Don't write anything
393 ];
394
395 ## Determine instruction status
396
397 int Stat = [
398     imem_error || dmem_error : SADR;
399     !instr_valid: SINS;
400     icode == IHALT : SHLT;
401     1 : SAOK;
402 ];

```

```
423
424
425 ##### Program Counter Update #####
426
427
428
429 ## What address should instruction be fetched at
430
431
432
433 int new_pc = [
434
435     # Call. Use instruction constant
436
437     icode == ICALL : valC;
438
439     # Taken branch. Use instruction constant
440
441     icode == IJXX && Cnd : valC;
442
443     # Completion of RET instruction. Use value from stack
444
445     icode == IRET : valM;
446
447     # Default: Use incremented PC
448
449     1 : valP;
450
451 ];
452
453 /* $end seq-all-hcl */
```

4.51

If instruction follows too closely after one that writes register, slow it down, Hold instruction in decode, Dynamically inject nop into execute stage.

Source Registers srcA and srcB of current instruction in decode stage.

Destination Registers dstE and dstM fields. Instructions in execute, memory, and write-back stages.

Special Case

1. Don't stall for register ID 15 (0xF).
2. Don't stall for failed conditional move (use e_dstE instead of E_dstE).

4.52

See *archlab partC* for detail.

4.53

See *archlab partC* for detail.

4.54

Because jump instruction doesn't use ALU. We could transport valC signal to Memory Stage by ALU :

$$ValE := 0 + ValC$$

Here is the difference between the original file and the modified file :

```

1
2
3 diff pipe-nt.hcl pipe-nt-origin.hcl
4
5 142,144c142
6 <
7 <   M_icode == IJXX && M_ifun != UNCOND && M_Cnd : M_valE;      # changed
8 <
9 ---
10 >   M_icode == IJXX && !M_Cnd : M_valA;
11 188,191c186
12 <
13 <   f_icode in { ICALL } : f_valC;
14 <   f_icode in { IJXX } && f_ifun == UNCOND : f_valC;    # changed
15 <
16 ---
17 >   f_icode in { IJXX, ICALL } : f_valC;
18 255c250
19 <   E_icode in { IIRMOVL, IRMMOVL, IMRMOVL,      IJXX } : E_valC;    # added IJXX
20 ---
21 >   E_icode in { IIRMOVL, IRMMOVL, IMRMOVL } : E_valC;
22 265c260
23 <   E_icode in { IRRMOVL, IIRMOVL,      IJXX } : 0;      # added IJXX
24 ---
25 >   E_icode in { IRRMOVL, IIRMOVL } : 0;
26 351c346
27 <   (E_icode == IJXX && E_ifun != UNCOND && e_Cnd) ||
28 ---
29 >   (E_icode == IJXX && !e_Cnd) ||
30 362c357
31 <   (E_icode == IJXX && E_ifun != UNCOND && e_Cnd) ||
32 ---
33 >   (E_icode == IJXX && !e_Cnd) ||

```

4.55

similar to 4.54

4.56**4.57**

(A) the condition of *load forwarding* :

```
1 D_icode in { IRMMOVL, IPUSHL } && E_dstM == d_srcA
```

(B) Here is the difference between the original file and the modified file :

```
1 diff pipe-lf.hcl pipe-lf-origin.hcl
2
3 274,277d273
4 <
5 < # added
6 < M_icode in { IMRMOVL, IPOPL } && E_icode in { IRMMOVL, IPUSHL } && M_dstM == E_srcA :
   m_valM;
7 <
8 336,340c332
9 < (
10 <     E_icode in { IMRMOVL, IPOPL } && E_dstM in { d_srcA, d_srcB } # load/use
11 <     && !(D_icode in { IRMMOVL, IPUSHL } && E_dstM == d_srcA ) # load/forwarding
12 < )
13 < ||
14 ---
15 > 0 ||
16 349,353c341
17 < (
18 <     E_icode in { IMRMOVL, IPOPL } && E_dstM in { d_srcA, d_srcB } # load/use
19 <     && !(D_icode in { IRMMOVL, IPUSHL } && E_dstM == d_srcA ) # load/
   forwarding
20 < )
21 < ;
22 ---
23 > 0;
24 371,375c359
25 < (
26 <     E_icode in { IMRMOVL, IPOPL } && E_dstM in { d_srcA, d_srcB } # load/use
27 <     && !(D_icode in { IRMMOVL, IPUSHL } && E_dstM == d_srcA ) # load/
   forwarding
28 < )
29 < ;
```

```
30 ---  
31 > 0;
```

4.58

program performance can be greatly enhanced if the compiler is able to generate code using conditional data transfers rather than conditional control transfers.

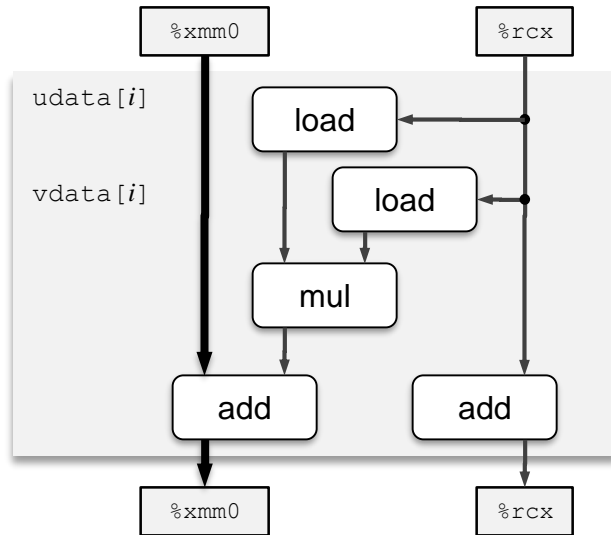
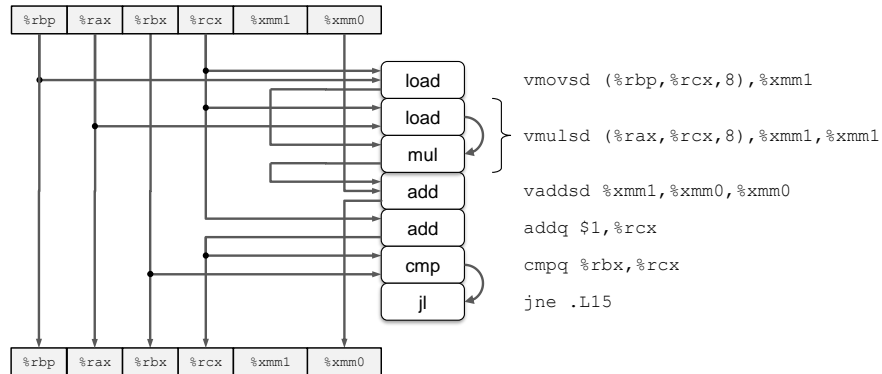
in page 530, CS:APP 2e:

```
1 // Version 1.
2 for (i = 0; i < n; i++) {
3     if (a[i] > b[i]) {
4         int t = a[i];
5         a[i] = b[i];
6         b[i] = t;
7     }
8 }
9 // Version 2.
10 for (i = 0; i < n; i++) {
11     int min = a[i] < b[i] ? a[i] : b[i];
12     int max = a[i] < b[i] ? b[i] : a[i];
13     a[i] = min;
14     b[i] = max;
15 }
```

5 Optimizing Program Performance

5.13

(A) In the figure below, *critical path* is emphasized by using bold.



(B) The addition operator of double, *latency* = 3.0

- (C) The addition operator of int, *latency* = 1.0
- (D) floating-point addition is in the *critical path* while floating-point multiplication isn't.

With 2 functional units capable of performing floating-point multiplication, the processor can potentially sustain a rate of 2 operations per cycle. The next multiplication launches before the end of the previous multiplication.

5.14

- (A) With 4 functional units capable of performing integer addition, the processor can potentially sustain a rate of 4 operations per cycle. Unfortunately, the need to read elements from memory creates an additional throughput bound. The 2 load units limit the processor to reading at most 2 data values per clock cycle. So processor can't achieve a CPE less than 1.00.
- (B) With 1 functional unit capable of performing floating-point addition, the processor can potentially sustain a rate of 1 operation per cycle. the performance is restricted by the Capacity of floating-point addition.

6 The Memory Hierarchy

6.23

7 Linking

7.8

- (A) (a) $REF(main.1) \rightarrow DEF(main.1)$
 (b) $REF(main.2) \rightarrow DEF(main.2)$
- (B) Here we have two weak definitions of `x`, so the symbol resolution in this case is UNKNOWN (Rule 3). But, notice that *double* is larger than *int*. It's usually behave as below:
- (a) $REF(main.1) \rightarrow DEF(main.2)$
 (b) $REF(main.2) \rightarrow DEF(main.2)$
- (C) There are two strong definitions of `x` (Rule 1), so this is an *ERROR*.

7.9

When this program is compiled and executed on a Linux system, it prints the string "0x48" and terminates normally, even though p2 never initializes variable main.

Can you explain this?

Because of Rule 2, the strong symbol associated with the function `main` in `foo6.o` overrides the weak symbol associated with the variable `main` in `bar6.o`.

Thus, the reference to variable `main` in `bar6` resolves to the value of symbol `main`, which in this case is the address of the first byte of function `main`. This byte contains the hex value `0x48`, which is the binary encoding of `pushq %rbp`, the first instruction in procedure `main`

7.12

- (A) *0xa*
 (B) *0x22*

7.13

- (A) *How many object files are contained in the versions of `libc.a` and `libm.a` on your system?*

`libc.a` has 1082 members and `libm.a` has 373 members.

- (B) *Does `gcc -O2` produce different executable code than `gcc -O2 -g`?*

The code in the `.text` section is identical, whether a program is compiled using `-g` or not. The difference is that the `-O2 -g` object file contains debugging info in the `.debug` section, while the `-O2` version does not.

- (C) *What shared libraries does the gcc driver on your system use?*

On our system, the gcc driver uses the standard C library (`libc.so.6`) and the dynamic linker (`ld-linux.so.2`).

8 Exceptional Control Flow

8.9

Process pair	Concurrent?
AB	N
AC	Y
AD	Y
BC	Y
BD	Y
CD	Y

8.10

In this chapter, we have introduced some functions with unusual call and return behaviors: `setjmp`, `longjmp`, `execve`, and `fork`. Match each function with one of the following behaviors:

- (A) Called once, returns twice.
- (B) Called once, never returns.
- (C) Called once, returns one or more times.

function	behavior
<code>setjmp</code>	C
<code>longjmp</code>	B
<code>execve</code>	B
<code>fork</code>	A

8.11

8.12

8

8.13

any of the following sequences represents a possible output:

$x = 4$	$x = 4$	$x = 2$
$x = 3$	$x = 2$	$x = 4$
$x = 2$	$x = 3$	$x = 3$

8.14

3

8.15

5

8.16

counter = 2

8.17

there are only three possible outcomes (each column is an outcome):

Hello	Hello	Hello
1	1	0
Bye	0	1
0	Bye	Bye
2	2	2
Bye	Bye	Bye

8.18

A C E

8.19

$$2^n$$

8.20

```
1 #include "csapp.h"
2
3 int main(int argc, char **argv) {
4     if (!getenv("COLUMNS"))
5     {
6         setenv("COLUMNS", 80, 1);
7         Execve("/bin/ls", argv, environ);
8         unsetenv("COLUMNS");
9     }
10    else
11    {
12        Execve("/bin/ls", argv, environ);
13    }
14    exit(0);
15 }
```

8.21

1. *bac*
2. *abc*

8.22

Write your own version of the Unix system function

```
1 int mysystem(char *command);
```

The `mysystem` function executes `command` by calling `"/bin/sh -c command"`, and then returns after `command` has completed. If `command` exits normally (by calling the `exit` function or executing a return statement), then `mysystem` returns the command exit status. For example, if `command` terminates by calling `exit(8)`, then `system` returns the value 8. Otherwise, if `command` terminates abnormally, then `mysystem` returns the status returned by the shell.

```
1 int mysystem(char *command) {
2     pid_t pid;
3     int status;
4
5     if (!(pid = Fork())) { /* child */
6         char *argv[4];
7         argv[0] = "sh";
8         argv[1] = "-c";
9         argv[2] = command;
10        argv[3] = NULL;
11        execve("/bin/sh", argv, environ);
12        exit(-1); /* control should never reach here */
13    }
14
15    /*
16     * parent
17     * In fact, we could use sigsuspend() instead of tight loop, which is better.
18     */
19    while (1) {
20        if (waitpid(pid, &status, 0) > 0) {
21            if (WIFEXITED(status))
22                return WEXITSTATUS(status);
23            else
24                return status;
25        }
26        else if (errno != EINTR) /* restart waitpid if interrupted */
27            return -1;
28    }
29 }
```

8.25

```
1 #include "csapp.h"
2 #include <unistd.h>
3
4 static sigjmp_buf env;
5
6 /* SIGALRM handler */
7 static void handler(int sig) {
8     siglongjmp(env, 1);
9     return;
10 }
11
12 char *tfgets(char *s, int size, FILE *stream) {
13     static handler_t *old_handler;
14     switch (sigsetjmp(env, 1)) {
15         case 0: {
16             old_handler = Signal(SIGALRM, handler);
17             Alarm(5);
18             Fgets(s, size, stream); /* return user input */
19
20             /*
21              * Signal(SIGALRM, SIG_IGN) is WRONG, the previous handler may not be SIG_IGN;
22              */
23             Signal(SIGALRM, old_handler); /* restore the old signal handler */
24             return s;
25         }
26         case 1: {
27             Signal(SIGALRM, old_handler); /* restore the old signal handler */
28             return NULL; /* return NULL if fgets times out */
29         }
30     }
31 }
32
33 int main() {
34     static char buf[1500012786];
35     while (1) {
36         if (tfgets(buf, sizeof(buf), stdin) != NULL)
37             printf("read: %s\n", buf);
38         else
39             printf("timed out\n");
40     }
41     exit(0);
42 }
```

8.26

plz see *tshlab* for detail.

9 Virtual Memory

9.11

(A) Virtual address format

13	12	11	10	9	8	7	6	5	4	3	2	1	0
0	0	0	0	1	0	0	1	1	1	1	1	0	0

(B) Address translation

Parameter	Value
VPN	0x9
TLB index	0x1
TLB tag	0x2
TLB hit?	N
Page fault?	N
PPN	0x17

(C) Physical address format

11	10	9	8	7	6	5	4	3	2	1	0
0	1	0	1	1	1	1	1	1	1	0	0

(D) Physical memory reference

Parameter	Value
Byte offset	0
Cache index	0xf
Cache tag	0x17
Cache hit?	N
Cache byte returned	-

9.12

(A) Virtual address format

13	12	11	10	9	8	7	6	5	4	3	2	1	0
0	0	0	0	1	1	1	0	1	0	1	0	0	1

(B) Address translation

Parameter	Value
VPN	0xe
TLB index	0x2
TLB tag	0x3
TLB hit?	N
Page fault?	N
PPN	0x11

(C) Physical address format

11	10	9	8	7	6	5	4	3	2	1	0
0	1	0	0	0	1	1	0	1	0	0	1

(D) Physical memory reference

Parameter	Value
Byte offset	0
Cache index	0xa
Cache tag	0x11
Cache hit?	N
Cache byte returned	-

9.13

(A) Virtual address format

13	12	11	10	9	8	7	6	5	4	3	2	1	0
0	0	0	0	0	0	0	1	0	0	0	0	0	0

(B) Address translation

Parameter	Value
VPN	0x1
TLB index	0x1
TLB tag	0x0
TLB hit?	N
Page fault?	Y
PPN	-

(C) N/A

(D) N/A

9.14

Given an input file `hello.txt` that consists of the string "Hello, world!\n", write a C program that uses `mmap` to change the contents of `hello.txt` to "Jello, world!\n".

```
1 #include "csapp.h"
2
3 /*
4  * mmapwrite - uses mmap to modify a disk file
5  */
6 void mmapwrite(int fd, int len)
7 {
8     char *bufp;
9
10    /* bufp = Mmap(NULL, len, PROT_READ|PROT_WRITE, MAP_SHARED, fd, 0); */
11    bufp = Mmap(NULL, len, PROT_READ|PROT_WRITE, MAP_PRIVATE, fd, 0);
12    bufp[0] = 'J';
13 }
14
15 /* mmapwrite driver */
16 int main(int argc, char **argv)
17 {
18     int fd;
19     struct stat stat;
20
21    /* check for required command line argument */
22    if (argc != 2) {
23        printf("usage: %s <filename>\n", argv[0]);
24        exit(0);
25    }
26
27    /* open the input file and get its size */
28    fd = Open(argv[1], O_RDWR, 0);
29    fstat(fd, &stat);
30    mmapwrite(fd, stat.st_size);
31    exit(0);
32 }
```

9.15

Determine the block sizes and header values that would result from the following sequence of *malloc* requests. Assumptions:

- (1) The allocator maintains double-word alignment, and uses an implicit free list with the block format from Figure 9.35.
- (2) Block sizes are rounded up to the nearest multiple of 8 bytes.

Request	Block size	Block header
malloc(3)	8	0x9
malloc(11)	16	0x11
malloc(20)	24	0x19
malloc(21)	32	0x21

9.16

Determine the minimum block size for each of the following combinations of alignment requirements and block formats. Assumptions: Explicit free list, 4-byte *pred* and *succ* pointers in each free block, zero-sized payloads are not allowed, and headers and footers are stored in 4-byte words.

Alignment	Allocated block	Free block	Minimum block size
Single word	Header and footer	Header and footer	16
Single word	Header, but no footer	Header and footer	16
Double word	Header and footer	Header and footer	16
Double word	Header, but no footer	Header and footer	16

9.17

Develop a version of the allocator in Section 9.9.12 that performs a next-fit search instead of a first-fit search.

I defined a global roving pointer (*void *rover*) that points initially to the front of the list.

```

1  /*
2   * If NEXT_FIT defined use next fit search, else use first-fit search
3   */
4  #define NEXT_FIT

```

```

1  /* Global variables */
2  static char *heap_listp = 0; /* Pointer to first block */
3  #ifndef NEXT_FIT
4  static char *rover;          /* Next fit rover */
5  #endif
6
7  /*
8   * mm_init - Initialize the memory manager
9   */
10 int mm_init(void)
11 {
12     /* Create the initial empty heap */
13     if ((heap_listp = mem_sbrk(4*WSIZE)) == (void *)-1) //line:vm:mm:begininit
14         return -1;
15     PUT(heap_listp, 0); /* Alignment padding */
16     PUT(heap_listp + (1*WSIZE), PACK(DSIZE, 1)); /* Prologue header */
17     PUT(heap_listp + (2*WSIZE), PACK(DSIZE, 1)); /* Prologue footer */
18     PUT(heap_listp + (3*WSIZE), PACK(0, 1)); /* Epilogue header */
19     heap_listp += (2*WSIZE); //line:vm:mm:endinit
20
21 #ifndef NEXT_FIT
22     rover = heap_listp;
23 #endif
24
25     /* Extend the empty heap with a free block of CHUNKSIZE bytes */
26     if (extend_heap(CHUNKSIZE/WSIZE) == NULL)
27         return -1;
28     return 0;
29 }

```

```

1  /*
2   * find_fit - Find a fit for a block with asize bytes
3   */
4  static void *find_fit(size_t asize)
5  {
6  #ifdef NEXT_FIT
7      /* Next fit search */
8      char *oldrover = rover;
9
10     /* Search from the rover to the end of list */
11     for ( ; GET_SIZE(HDRP(rover)) > 0; rover = NEXT_BLK(P(rover)))
12         if (!GET_ALLOC(HDRP(rover)) && (asize <= GET_SIZE(HDRP(rover))))
13             return rover;
14
15     /* search from start of list to old rover */
16     for (rover = heap_listp; rover < oldrover; rover = NEXT_BLK(P(rover)))
17         if (!GET_ALLOC(HDRP(rover)) && (asize <= GET_SIZE(HDRP(rover))))
18             return rover;
19
20     return NULL; /* no fit found */
21 #else
22     /* First-fit search */
23     void *bp;
24
25     for (bp = heap_listp; GET_SIZE(HDRP(bp)) > 0; bp = NEXT_BLK(P(bp))) {
26         if (!GET_ALLOC(HDRP(bp)) && (asize <= GET_SIZE(HDRP(bp)))) {
27             return bp;
28         }
29     }
30     return NULL; /* No fit */
31 #endif
32 }

```

After coalescing a block, Make sure the rover isn't pointing into the free block that we just coalesced.

```

1 static void *coalesce(void *bp)
2 {
3     int prev_alloc = GET_ALLOC(FTRP(PREV_BLKPTR(bp)));
4     int next_alloc = GET_ALLOC(HDRP(NEXT_BLKPTR(bp)));
5     size_t size = GET_SIZE(HDRP(bp));
6
7     if (prev_alloc && next_alloc) { /* Case 1 */
8         return bp;
9     }
10
11    else if (prev_alloc && !next_alloc) { /* Case 2 */
12        size += GET_SIZE(HDRP(NEXT_BLKPTR(bp)));
13        PUT(HDRP(bp), PACK(size, 0));
14        PUT(FTRP(bp), PACK(size, 0));
15    }
16
17    else if (!prev_alloc && next_alloc) { /* Case 3 */
18        size += GET_SIZE(HDRP(PREV_BLKPTR(bp)));
19        PUT(FTRP(bp), PACK(size, 0));
20        PUT(HDRP(PREV_BLKPTR(bp)), PACK(size, 0));
21        bp = PREV_BLKPTR(bp);
22    }
23
24    else { /* Case 4 */
25        size += GET_SIZE(HDRP(PREV_BLKPTR(bp))) +
26            GET_SIZE(FTRP(NEXT_BLKPTR(bp)));
27        PUT(HDRP(PREV_BLKPTR(bp)), PACK(size, 0));
28        PUT(FTRP(NEXT_BLKPTR(bp)), PACK(size, 0));
29        bp = PREV_BLKPTR(bp);
30    }
31
32    #ifndef NEXT_FIT
33        /* Make sure the rover isn't pointing into the free block */
34        /* that we just coalesced */
35        if ((rover > (char *)bp) && (rover < NEXT_BLKPTR(bp)))
36            rover = bp;
37    #endif
38
39    return bp;
40 }

```

See *code/vm/malloc/mm.c* for detail.

9.18

The allocator in Section 9.9.12 requires both a header and a footer for each block in order to perform constant-time coalescing. Modify the allocator so that free blocks require a header and footer, but allocated blocks require only a header.

```

1 #define ALLOC_NO_FOOTER

```

```

1 /* Read the size and allocated fields from address p */
2 #define GET_SIZE(p) (GET(p) & ~0x7) //line:vm:mm:getsize
3 #define GET_ALLOC(p) (GET(p) & 0x1) //line:vm:mm:getalloc
4
5 #ifndef ALLOC_NO_FOOTER
6 #define GET_PREV_ALLOC(p) (GET(p) & 0x2)
7 #endif

```

```

1 /*
2  * mm_init - Initialize the memory manager
3  */
4 int mm_init(void) {
5     /* Create the initial empty heap */
6     if ((heap_listp = mem_sbrk(4*WSIZE)) == (void *)-1) //line:vm:mm:begininit
7         return -1;
8     PUT(heap_listp, 0); // Alignment padding */
9     PUT(heap_listp + (1*WSIZE), PACK(DSIZE, 1)); /* Prologue header */
10
11 #ifndef ALLOC_NO_FOOTER
12     PUT(heap_listp + (2*WSIZE), PACK(DSIZE, 0x3)); /* Prologue footer */
13 else
14     PUT(heap_listp + (2*WSIZE), PACK(DSIZE, 1)); /* Prologue footer */
15 #endif
16
17     PUT(heap_listp + (3*WSIZE), PACK(0, 1)); /* Epilogue header */
18     heap_listp += (2*WSIZE); //line:vm:mm:endinit
19
20     /* Extend the empty heap with a free block of CHUNKSIZE bytes */
21     if (extend_heap(CHUNKSIZE/WSIZE) == NULL) return -1; else return 0;
22 }

```

```

1  /*
2  * place - Place block of asize bytes at start of free block bp
3  *          and split if remainder would be at least minimum block size
4  */
5  static void place(void *bp, size_t asize)
6  {
7  #ifdef ALLOC_NO_FOOTER
8      size_t csize = GET_SIZE(HDRP(bp));
9
10     if ((csize - asize) >= (2*DSIZE)) {
11         PUT(HDRP(bp), PACK(asize, 1));
12         bp = NEXT_BLK(P(bp));
13         PUT(HDRP(bp), PACK(csize-asize, 0x2));
14         PUT(FTRP(bp), PACK(csize-asize, 0x2));
15     }
16     else {
17         PUT(HDRP(bp), PACK(csize, 1));
18         bp = NEXT_BLK(P(bp));
19         GET(HDRP(bp)) |= 0x2;
20         GET(FTRP(bp)) |= 0x2;
21     }
22 #else
23     size_t csize = GET_SIZE(HDRP(bp));
24
25     if ((csize - asize) >= (2*DSIZE)) {
26         PUT(HDRP(bp), PACK(asize, 1));
27         PUT(FTRP(bp), PACK(asize, 1));
28         bp = NEXT_BLK(P(bp));
29         PUT(HDRP(bp), PACK(csize-asize, 0));
30         PUT(FTRP(bp), PACK(csize-asize, 0));
31     }
32     else {
33         PUT(HDRP(bp), PACK(csize, 1));
34         PUT(FTRP(bp), PACK(csize, 1));
35     }
36 #endif
37 }

```

9.19

You are given three groups of statements relating to memory management and garbage collection below. In each group, only one statement is true. Your task is to indicate which statement is true.

- (1)
 - (a) In a buddy system, up to 50% of the space can be wasted due to internal fragmentation.
 - (b) The first-fit memory allocation algorithm is slower than the best-fit algorithm (on average).
 - (c) Deallocation using boundary tags is fast only when the list of free blocks is ordered according to increasing memory addresses.
 - (d) The buddy system suffers from internal fragmentation, but not from external fragmentation.
- (2)
 - (a) Using the first-fit algorithm on a free list that is ordered according to decreasing block sizes results in low performance for allocations, but avoids external fragmentation.
 - (b) For the best-fit method, the list of free blocks should be ordered according to increasing memory addresses.
 - (c) The best-fit method chooses the largest free block into which the requested segment fits.
 - (d) Using the first-fit algorithm on a free list that is ordered according to increasing block sizes is equivalent to using the best-fit algorithm.
- (3) Mark & sweep garbage collectors are called conservative if:
 - (a) They coalesce freed memory only when a memory request cannot be satisfied.
 - (b) They treat everything that looks like a pointer as a pointer.
 - (c) They perform garbage collection only when they run out of memory.
 - (d) They do not free memory blocks forming a cyclic list.

(1) a

(2) d

(3) b

9.20

Write your own version of malloc and free, and compare its running time and space utilization to the version of malloc provided in the standard C library.

See *Malloc Lab* for detail.

10 System-Level I/O

10.6

What is the output of the following program?

```
1 #include "csapp.h"
2
3 int main()
4 {
5     int fd1, fd2;
6     fd1 = Open("foo.txt", O_RDONLY, 0);
7     fd2 = Open("bar.txt", O_RDONLY, 0);
8     Close(fd2);
9     fd2 = Open("baz.txt", O_RDONLY, 0);
10    printf("fd2 = %d\n", fd2);
11    exit(0);
12 }
```

On entry, descriptors 0-2 are already open. The open function always returns the lowest possible descriptor, so the first two calls to open return descriptors 3 and 4. The call to the close function frees up descriptor 4, so the final call to open returns descriptor 4, and thus the output of the program is $fd2 = 4$.

10.7

Modify the cpfile program in Figure 10.5 so that it uses the Rio functions to copy standard input to standard output, MAXBUF bytes at a time.

```
1 #include "csapp.h"
2
3 int main(int argc, char **argv) {
4     int n;
5     char buf[MAXBUF];
6
7     while((n = Rio_readn(STDIN_FILENO, buf, MAXBUF)) != 0)
8         Rio_writen(STDOUT_FILENO, buf, n);
9     exit(0);
10 }
```

10.8

Write a version of the statcheck program in Figure 10.10, called fstatcheck, that takes a descriptor number on the command line rather than a file name.

Just invoke fstat instead of stat, Here is my solution:

```
1  #include "csapp.h"
2
3  int main (int argc, char **argv)
4  {
5      struct stat stat;
6      char *type, *readok;
7      int size;
8
9      if (argc != 2) {
10         fprintf(stderr, "usage: %s <fd>\n", argv[0]);
11         exit(0);
12     }
13     Fstat(atoi(argv[1]), &stat);
14     if (S_ISREG(stat.st_mode)) /* Determine file type */
15         type = "regular";
16     else if (S_ISDIR(stat.st_mode))
17         type = "directory";
18     else if (S_ISCHR(stat.st_mode))
19         type = "character device";
20     else
21         type = "other";
22
23     if ((stat.st_mode & S_IRUSR)) /* Check read access */
24         readok = "yes";
25     else
26         readok = "no";
27
28     size = stat.st_size; /* check size */
29
30     printf("type: %s, read: %s, size=%d\n", type, readok, size);
31
32     exit(0);
33 }
```

10.9

Consider the following invocation of the `fstatcheck` program from Problem 10.8:

```
1 unix> fstatcheck 3 < foo.txt
```

You might expect that this invocation of `fstatcheck` would fetch and display metadata for file `foo.txt`. However, when we run it on our system, it fails with a **“bad file descriptor.”** Given this behavior, fill in the pseudo-code that the shell must be executing between the `fork` and `execve` calls:

```
1 if (Fork() == 0) { /* Child */  
2     /* What code is the shell executing right here? */  
3     Execve("fstatcheck", argv, envp);  
4 }
```

Before the call to `execve`, the child process opens `foo.txt` as descriptor 3, redirects `stdin` to `foo.txt`, and then (here is the kicker) closes descriptor 3:

```
1 if (Fork() == 0) { /* child */  
2     fd = Open("foo.txt", O_RDONLY, 0); /* fd == 3 */  
3     Dup2(fd, STDIN_FILENO);  
4     Close(fd);  
5     Execve("fstatcheck", argv, envp);  
6 }
```

When `fstatcheck` begins running in the child, there are exactly three open files, corresponding to descriptors 0, 1, and 2, with descriptor 1 redirected to `foo.txt`.

10.10

Modify the cpfile program in Figure 10.5 so that it takes an optional command line argument infile. If infile is given, then copy infile to standard output; otherwise, copy standard input to standard output as before. The twist is that your solution must use the original copy loop (lines 9–11) for both cases. You are only allowed to insert code, and you are not allowed to change any of the existing code.

```
1 #include "csapp.h"
2
3 int main(int argc, char **argv) {
4     int n;
5     rio_t rio;
6     static char buf[MAXLINE];
7
8     if ((argc < 1) || (argc > 2) ) {
9         fprintf(stderr, "usage: %s <infile>\n", argv[0]);
10        exit(1);
11    }
12
13    if (argc == 2) {
14        int fd;
15        if ((fd = Open(argv[1], O_RDONLY, 0)) < 0) {
16            fprintf(stderr, "Couldn't read %s\n", argv[1]);
17            exit(1);
18        }
19        Dup2(fd, STDIN_FILENO);
20        Close(fd);
21    }
22
23    Rio_readinitb(&rio, STDIN_FILENO);
24    while((n = Rio_readlineb(&rio, buf, MAXLINE)) != 0)
25        Rio_writen(STDOUT_FILENO, buf, n);
26    exit(0);
27 }
```

11 Network Programming

11.6

(A) It's the same as our *echo server*.

```
(B) $ ./tiny 8000
 2 Accepted connection from (localhost, 40734)
 3 GET / HTTP/1.1
 4 Host: localhost:8000
 5 User-Agent: Mozilla/5.0 (X11; Ubuntu; Linux x86_64; rv:50.0) Gecko/20100101 Firefox/50.0
 6 Accept: text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8
 7 Accept-Language: en-US,en;q=0.5
 8 Accept-Encoding: gzip, deflate
 9 Connection: keep-alive
10 Upgrade-Insecure-Requests: 1
11
12 Response headers:
13 HTTP/1.0 200 OK
14 Server: Tiny Web Server
15 Connection: close
16 Content-length: 120
17 Content-type: text/html
18
19 Accepted connection from (localhost, 40736)
20 GET /godzilla.gif HTTP/1.1
21 Host: localhost:8000
22 User-Agent: Mozilla/5.0 (X11; Ubuntu; Linux x86_64; rv:50.0) Gecko/20100101 Firefox/50.0
23 Accept: */*
24 Accept-Language: en-US,en;q=0.5
25 Accept-Encoding: gzip, deflate
26 Referer: http://localhost:8000/
27 Connection: keep-alive
28
29 Response headers:
30 HTTP/1.0 200 OK
31 Server: Tiny Web Server
32 Connection: close
33 Content-length: 12155
34 Content-type: image/gif
```

(C) HTTP/1.1

(D) Host: 请求资源所在的域名IP以及端口号。

User-Agent: 代理（浏览器）。

Accept: 指定接受的介质类型。本例中为首选网页文件和图片文件。

Referer: 来源网页信息。本例中图片文件来源为home.html。

Accept-Encoding: 指定接受的编码方法，通常为压缩方法。本例中为首选gzip压缩格式。

Accept-Language: 指定接受的语言。本例中为首选中文。

Connection: 指定完成本次连接后是否断开连接。本例中为保持连接。

11.8

Modify *Tiny* so that it reaps *CGI* children inside a *SIGCHLD* handler instead of explicitly waiting for them to terminate.

Install a *SIGCHLD* handler in the main routine and delete the call to *wait* in *serve dynamic*.

```
1 void sigchld_handler(int sig)
2 {
3     pid_t pid;
4
5     while ((pid = waitpid(-1, NULL, 0)) > 0)
6         printf("Handler reaped child %d\n", (int)pid);
7
8     if (errno != ECHILD)
9         unix_error("waitpid error");
10
11     return;
12 }
```

12 Concurrent Programming

12.16

```
1 #include "csapp.h"
2
3 void *thread(void *vargp);
4
5 int main(int argc, char **argv)
6 {
7     pthread_t *tid;
8     int i, n;
9
10    if (argc != 2)
11    {
12        fprintf(stderr, "usage: %s <nthreads>\n", argv[0]);
13        exit(0);
14    }
15    n = atoi(argv[1]);
16    tid = Malloc(n * sizeof(pthread_t));
17
18    for (i = 0; i < n; i++)
19        Pthread_create(&tid[i], NULL, thread, NULL);
20
21    for (i = 0; i < n; i++)
22        Pthread_join(tid[i], NULL);
23    exit(0);
24 }
25
26 /* thread routine */
27 void *thread(void *vargp)
28 {
29    printf("Hello, world!\n");
30    return NULL;
31 }
```

12.17

This is the student's first introduction to the many synchronization problems that can arise in threaded programs.

- (A) The problem is that the main thread calls `exit` without waiting for the peer thread to terminate. The `exit` call terminates the entire process, including any threads that happen to be running. So the peer thread is being killed before it has a chance to print its output string.
- (B) We can fix the bug by replacing the `exit` function with either `pthread_exit`, which waits for outstanding threads to terminate before it terminates the process, or `pthread_join`, which explicitly reaps the peer thread.

12.18

(A) unsafe

(B) safe

(C) unsafe

12.22

Test your understanding of the select function by modifying the server in Figure 12.6 so that it echoes at most one text line per iteration of the main server loop.

```
1 #include "csapp.h"
2 void echo(int connfd);
3 void command(void);
4
5 int main(int argc, char **argv)
6 {
7     int listenfd, connfd, port;
8     socklen_t clientlen = sizeof(struct sockaddr_in);
9     struct sockaddr_in clientaddr;
10    fd_set read_set, ready_set;
11
12    if (argc != 2) {
13        fprintf(stderr, "usage: %s <port>\n", argv[0]);
14        exit(0);
15    }
16    port = atoi(argv[1]);
17    listenfd = Open_listenfd(port);
18
19    FD_ZERO(&read_set); /* Clear read set */
20    FD_SET(STDIN_FILENO, &read_set); /* Add stdin to read set */
21    FD_SET(listenfd, &read_set); /* Add listenfd to read set */
22
23
24    /*
25     * * rio init
26     */
27    rio_t rio;
28    Rio_readinitb(&rio, connfd);
29
```

```

30
31     while (1) {
32         ready_set = read_set;
33         Select(listenfd+1, &ready_set, NULL, NULL, NULL);
34         if (FD_ISSET(STDIN_FILENO, &ready_set))
35             command(); /* Read command line from stdin */
36         if (FD_ISSET(listenfd, &ready_set)) {
37             connfd = Accept(listenfd, (SA *)&clientaddr, &clientlen);
38
39             /*
40              * * echoes at most one text line per iteration.
41              */
42             size_t n;
43             static char buf[MAXLINE];
44
45             if ((n = Rio_readlineb(&rio, buf, MAXLINE)) != 0) {
46                 printf("server received %d bytes\n", n);
47                 Rio_writen(connfd, buf, n);
48             }
49
50             Close(connfd);
51         }
52     }
53 }
54
55 void command(void) {
56     char buf[MAXLINE];
57     if (!Fgets(buf, MAXLINE, stdin))
58         exit(0); /* EOF */
59     printf("%s", buf); /* Process the input command */
60 }

```

12.23

12.24

The functions in the Rio I/O package (Section 10.5) are thread-safe. Are they reentrant as well?

- (A) Each of the buffered Rio functions is passed a pointer to a buffer, and then operates exclusively on this buffer and local stack variables. If

they are invoked properly by the calling function, such that none of the buffers are shared, then they are reentrant.

- (B) We can fix the bug by replacing the `exit` function with either `pthread_exit`, which waits for outstanding threads to terminate before it terminates the process, or `pthread_join`, which explicitly reaps the peer thread.

12.25

In the prethreaded concurrent echo server in Figure 12.28, each thread calls the `echo_cnt` function (Figure 12.29). Is `echo_cnt` thread-safe? Is it reentrant? Why or why not?

The `echo_cnt` function is thread-safe because (a) It protects accesses to the shared global `byte_cnt` with a mutex, and (b) All of the functions that it calls, such as `rio_readlineb` and `rio_writen`, are thread-safe. However, because of the shared variable, `echo_cnt` is not reentrant.

12.26

```

1  /*
2   * gethostbyname_ts - A thread-safe wrapper for gethostbyname
3   */
4  #include "csapp.h"
5
6  static sem_t mutex; /* protects calls to gethostbyname */
7
8  static void init_gethostbyname_ts(void) {
9      Sem_init(&mutex, 0, 1);
10 }
11
12 struct hostent *gethostbyname_ts(char *hostname) {
13     struct hostent *sharedp, *unsharedp;
14
15     unsharedp = Malloc(sizeof(struct hostent));
16     P(&mutex);
17     sharedp = gethostbyname(hostname);
18     *unsharedp = *sharedp; /* copy shared struct to private struct */
19     V(&mutex);
20     return unsharedp;
21 }
22
23 int main(int argc, char **argv) {
24     char **pp; struct in_addr addr; struct hostent *hostp;
25
26     if (argc != 2) {
27         fprintf(stderr, "usage: %s <hostname>\n", argv[0]);
28         exit(0);
29     }
30
31     init_gethostbyname_ts();
32     hostp = gethostbyname_ts(argv[1]);
33     if (hostp) {
34         printf("official hostname: %s\n", hostp->h_name);
35         for (pp = hostp->h_aliases; *pp != NULL; pp++)
36             printf("alias: %s\n", *pp);
37         for (pp = hostp->h_addr_list; *pp != NULL; pp++) {
38             addr.s_addr = *((unsigned int *)*pp);
39             printf("address: %s\n", inet_ntoa(addr));
40         }
41     }
42     else
43         printf("host %s not found\n", argv[1]);
44 }

```
