

Processor Architecture II: PIPE: Pipelined Implementation

Introduction to Computer Systems
11th Lecture, Oct 24, 2016

Instructors:

Xiangqun Chen , Junlin Lu

Guangyu Sun , Xuetao Guan

Shiliang Zhang

Pipeline Part 1: Overview

- **General Principles of Pipelining**
 - Goal
 - Difficulties
- **Creating a Pipelined Y86 Processor**
 - Rearranging SEQ
 - Inserting pipeline registers
 - Problems with data and control hazards

Real-World Pipelines: Car Washes

Sequential



Parallel



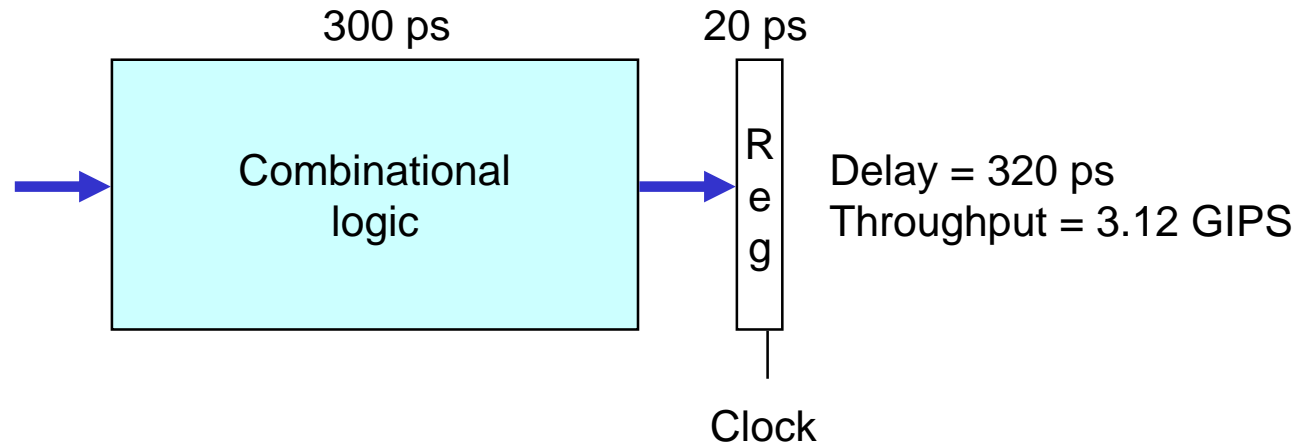
Pipelined



■ Idea

- Divide process into independent stages
- Move objects through stages in sequence
- At any given times, multiple objects being processed

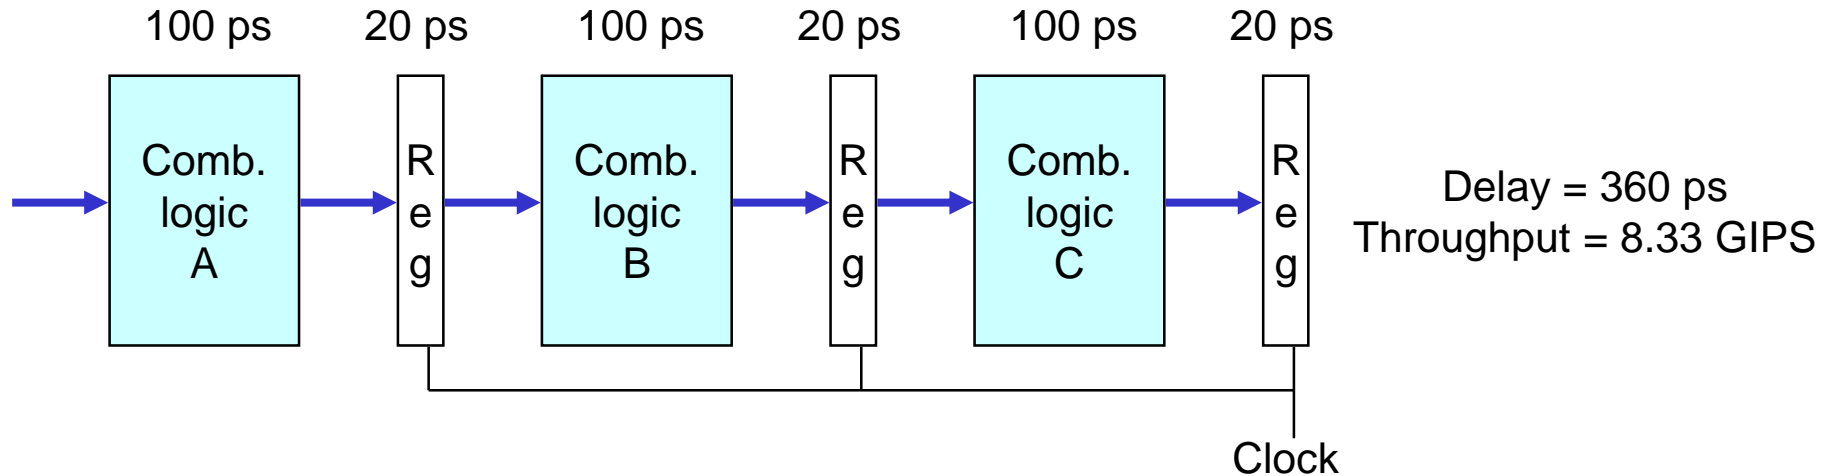
Computational Example



■ System

- Computation requires total of 300 picoseconds
- Additional 20 picoseconds to save result in register
- Must have clock cycle of at least 320 ps

3-Way Pipelined Version

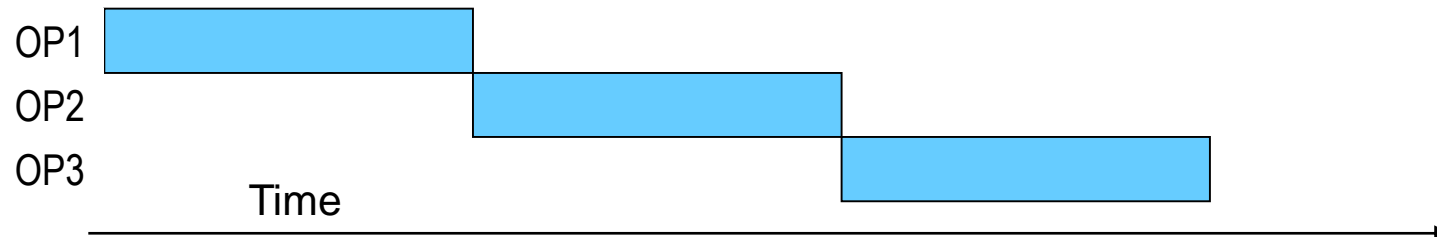


■ System

- Divide combinational logic into 3 blocks of 100 ps each
- Can begin new operation as soon as previous one passes through stage A.
 - Begin new operation every 120 ps
- Overall latency increases
 - 360 ps from start to finish

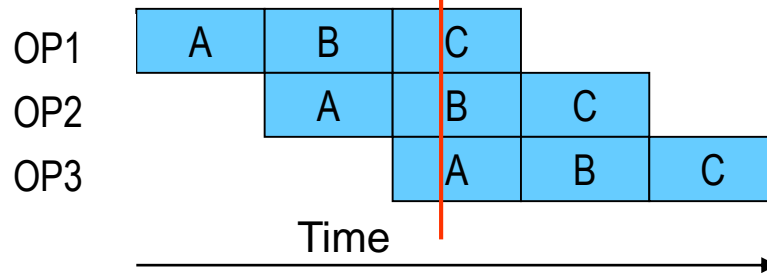
Pipeline Diagrams

■ Unpipelined



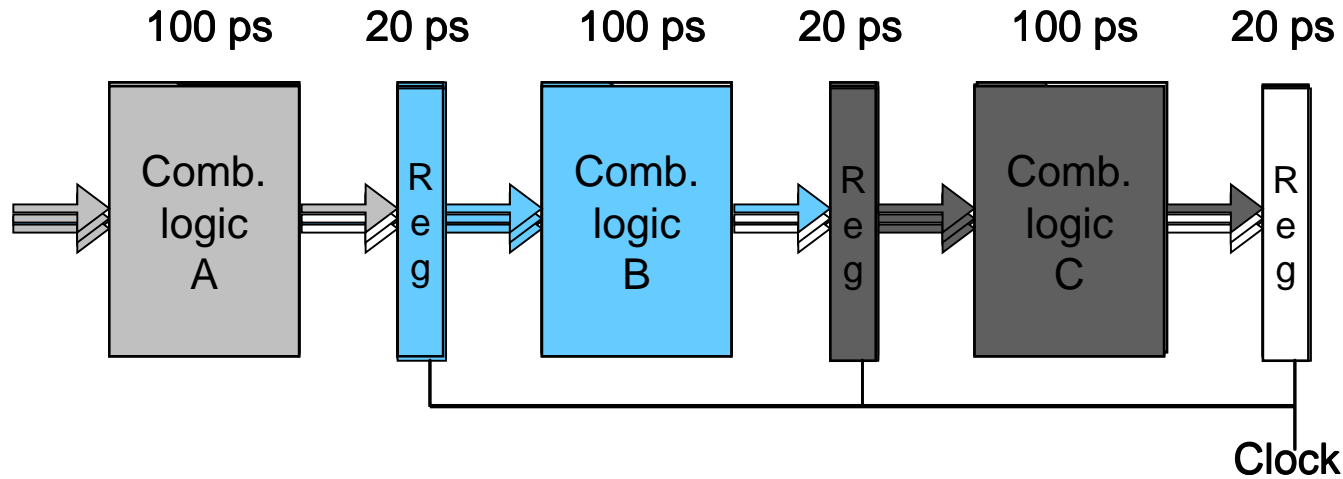
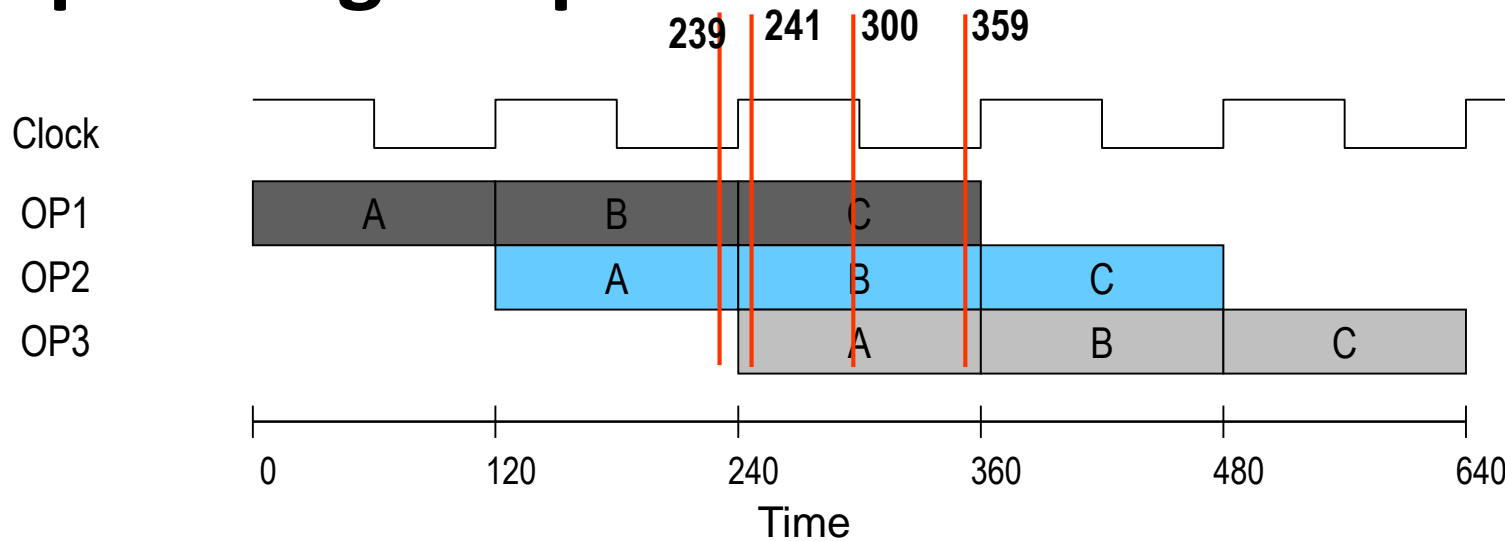
- Cannot start new operation until previous one completes

■ 3-Way Pipelined

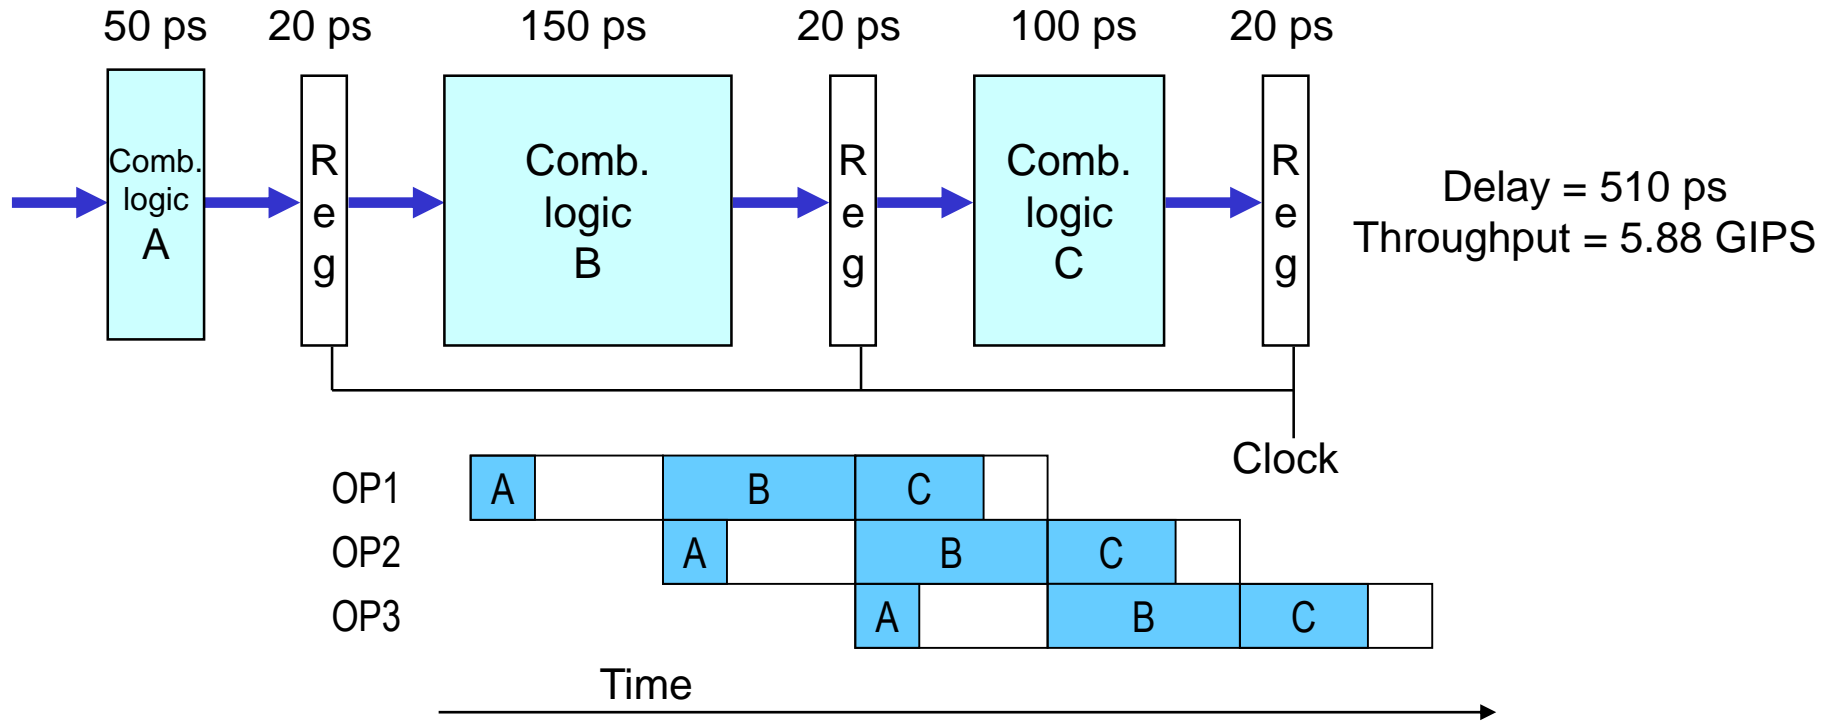


- Up to 3 operations in process simultaneously

Operating a Pipeline

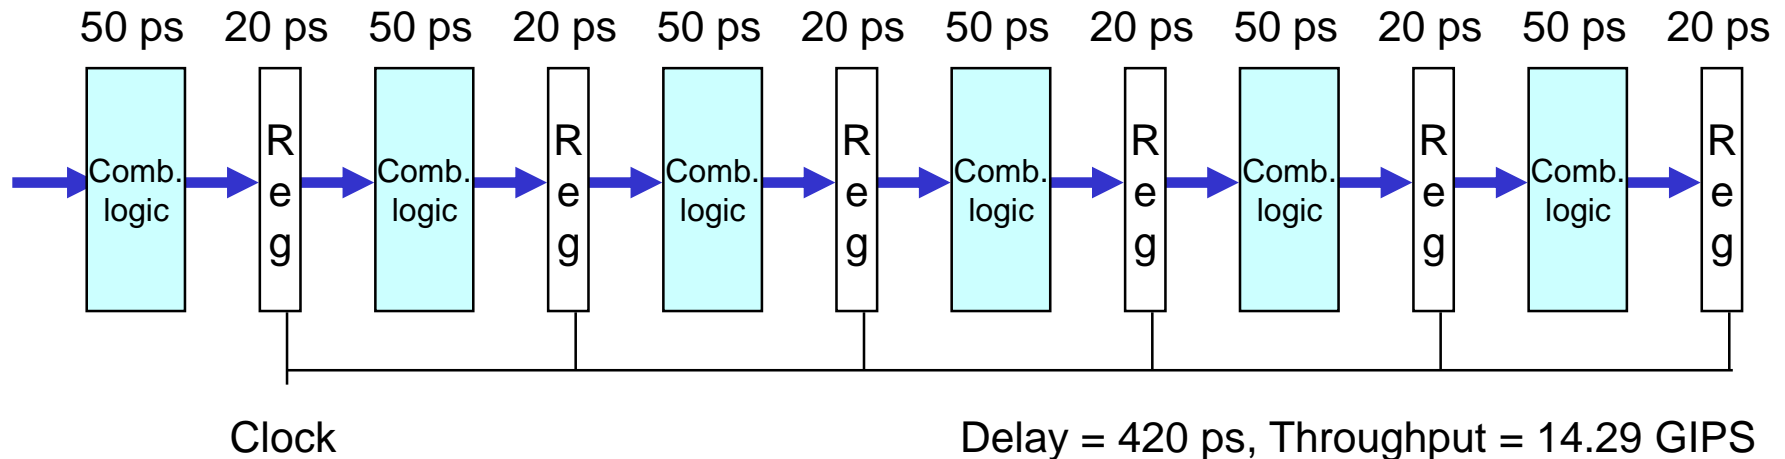


Limitations: Nonuniform Delays



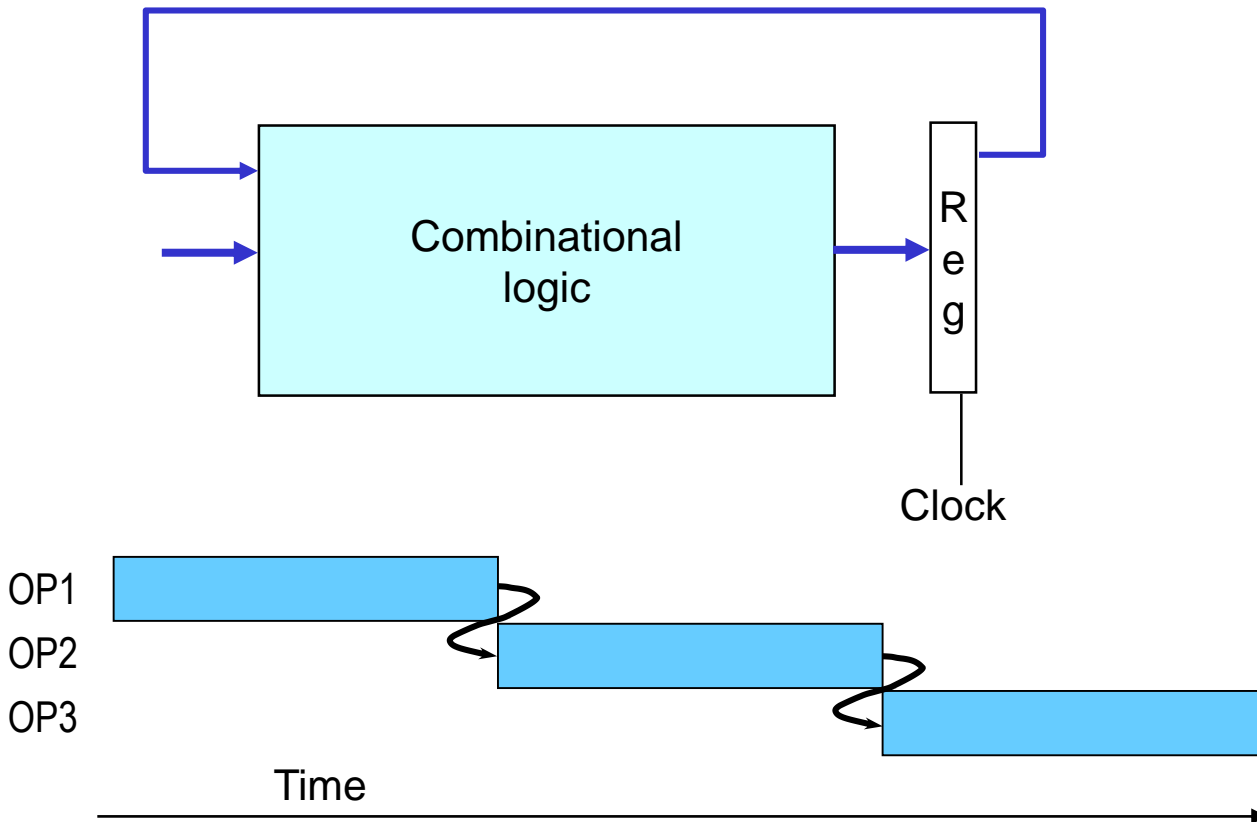
- Throughput limited by slowest stage
- Other stages sit idle for much of the time
- Challenging to partition system into balanced stages

Limitations: Register Overhead



- As try to deepen pipeline, overhead of loading registers becomes more significant
- Percentage of clock cycle spent loading register:
 - 1-stage pipeline: 6.25%
 - 3-stage pipeline: 16.67%
 - 6-stage pipeline: 28.57%
- High speeds of modern processor designs obtained through very deep pipelining

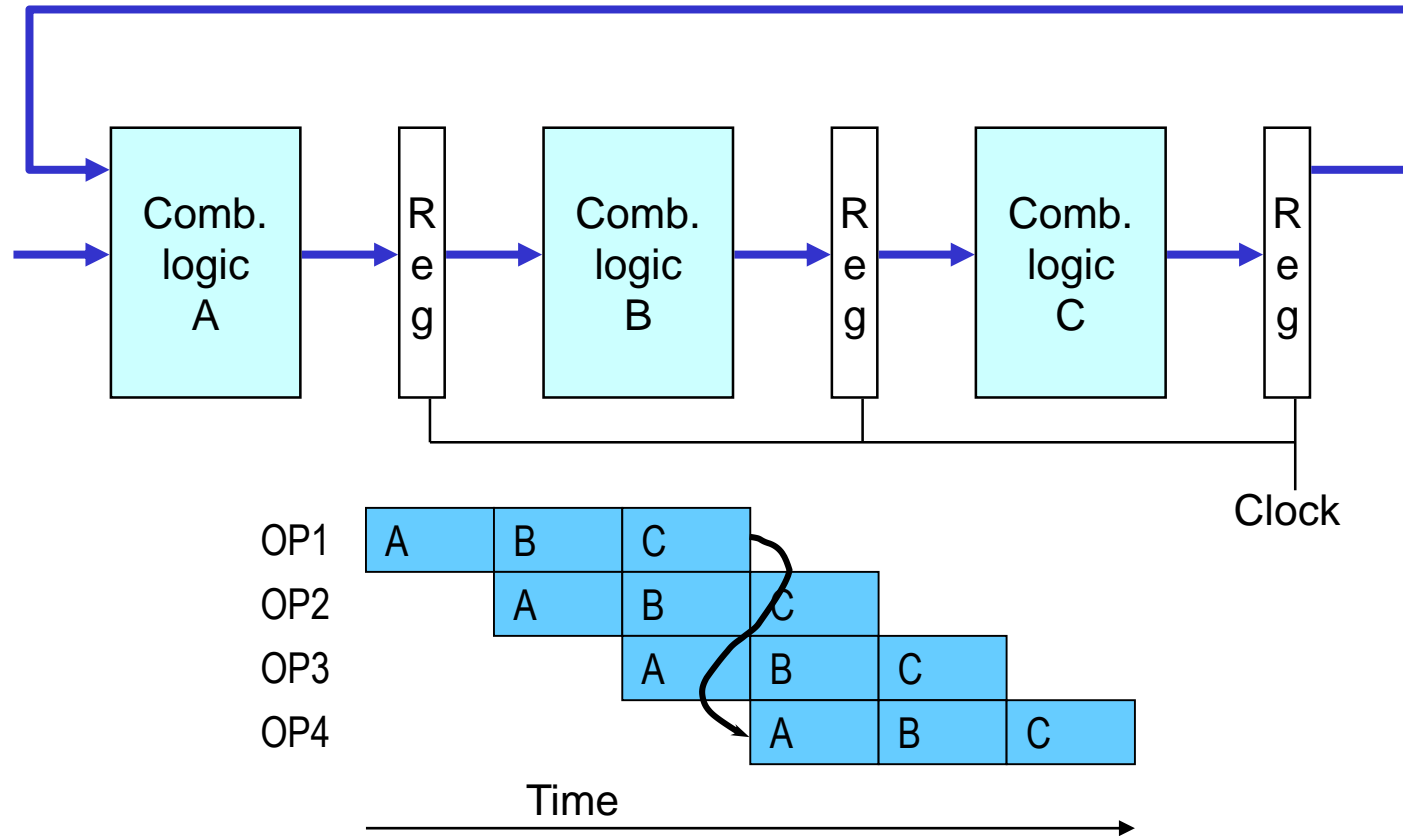
Data Dependencies



■ System

- Each operation depends on result from preceding one

Data Hazards



- Result does not feed back around in time for next operation
- Pipelining has changed behavior of system

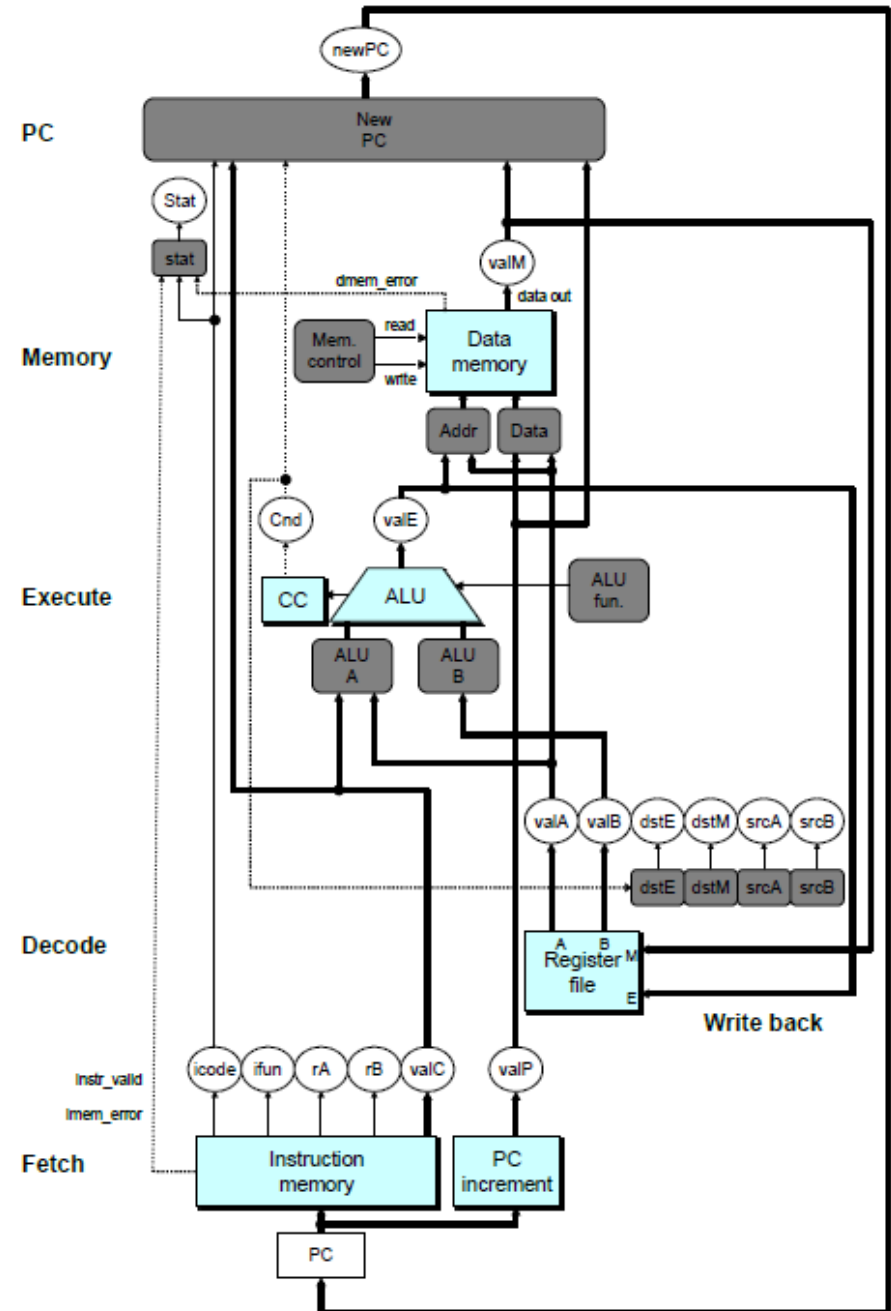
Data Dependencies in Processors

```
1    irmovl $50, %eax
2    addl %eax, %ebx
3    mrmovl 100(%ebx), %edx
```

- **Result from one instruction used as operand for another**
 - Read-after-write (RAW) dependency
- **Very common in actual programs**
- **Must make sure our pipeline handles these properly**
 - Get correct results
 - Minimize performance impact

SEQ Hardware

- Stages occur in sequence
- One operation in process at a time



SEQ+ Hardware

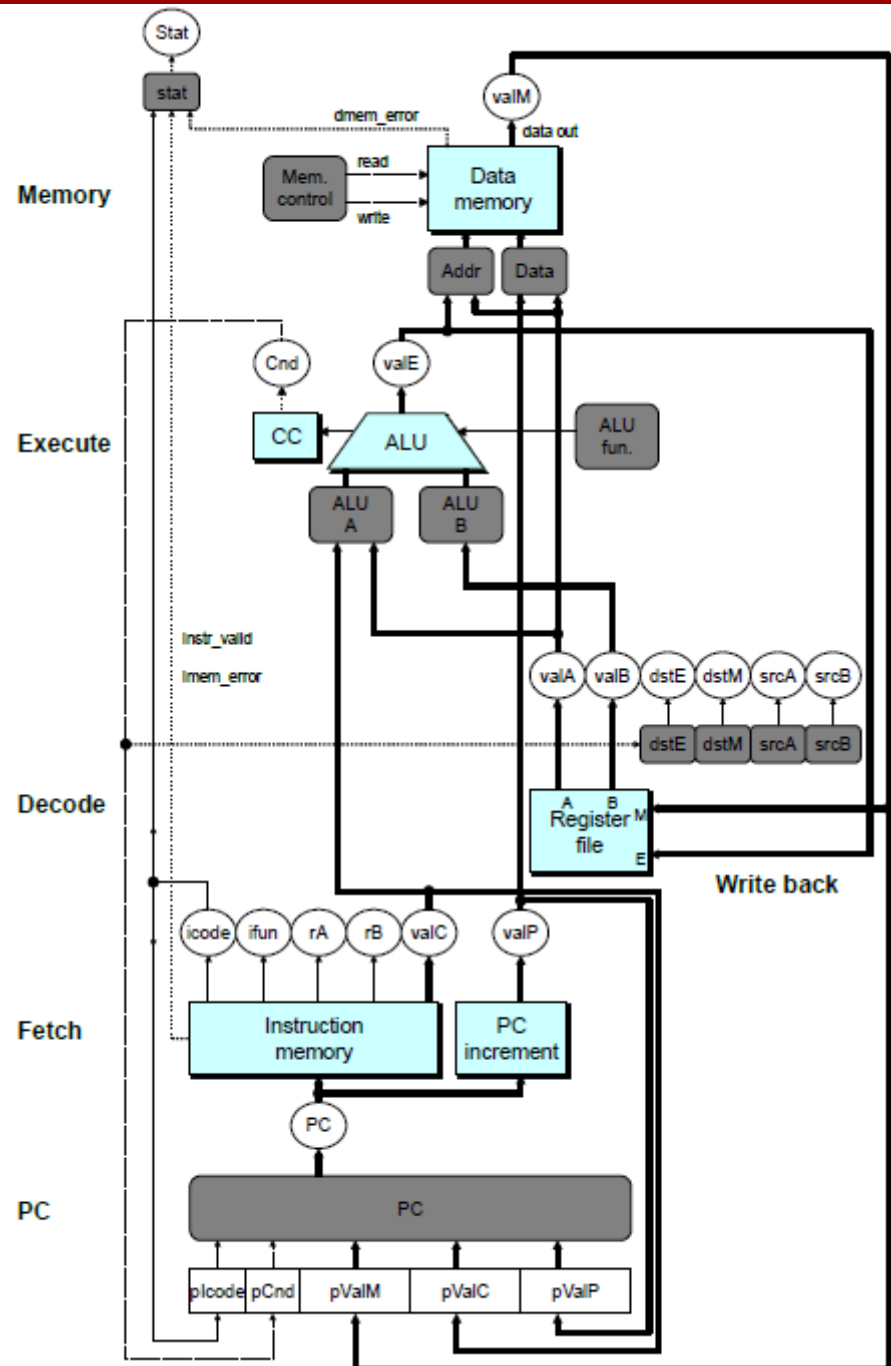
- Still sequential implementation
- Reorder PC stage to put at beginning

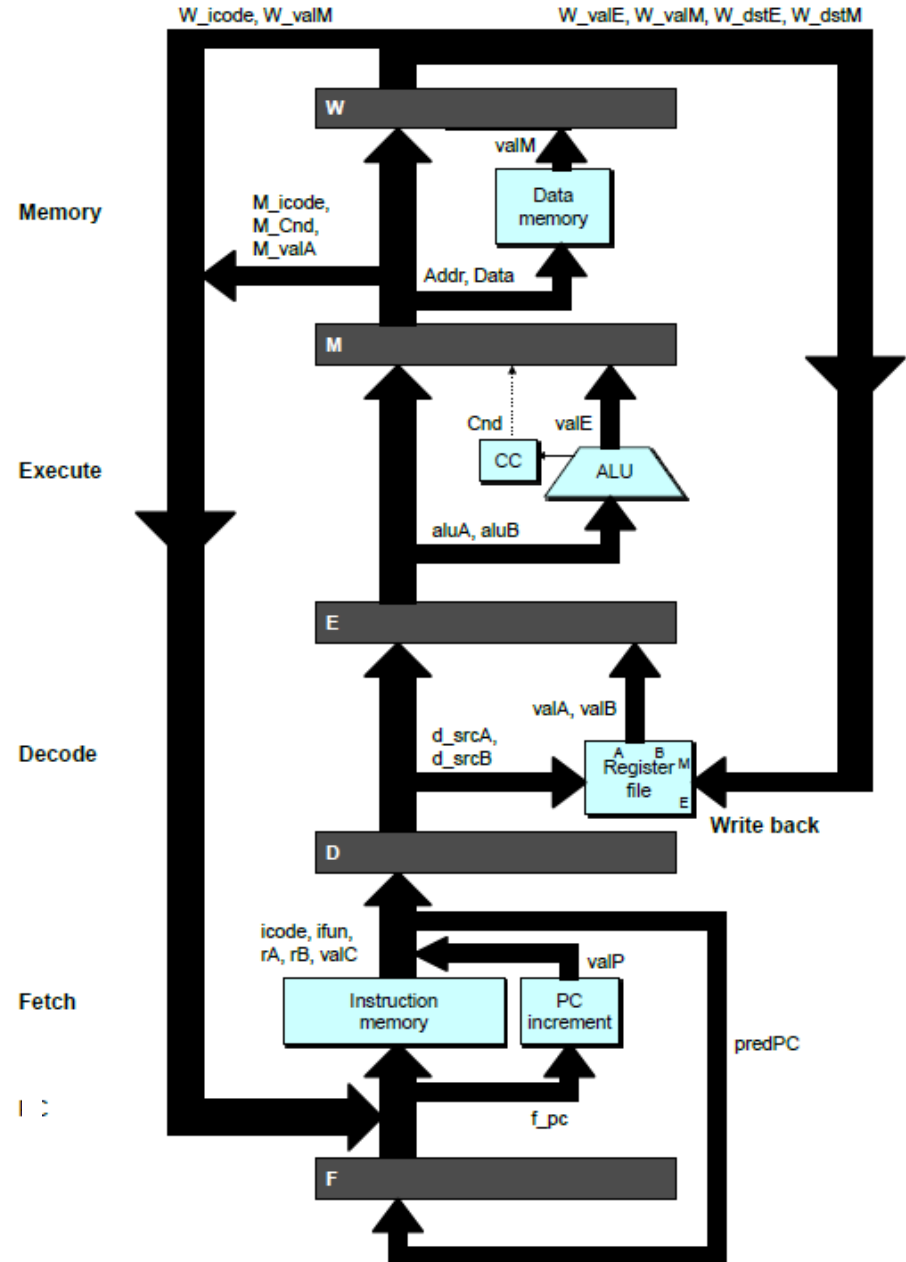
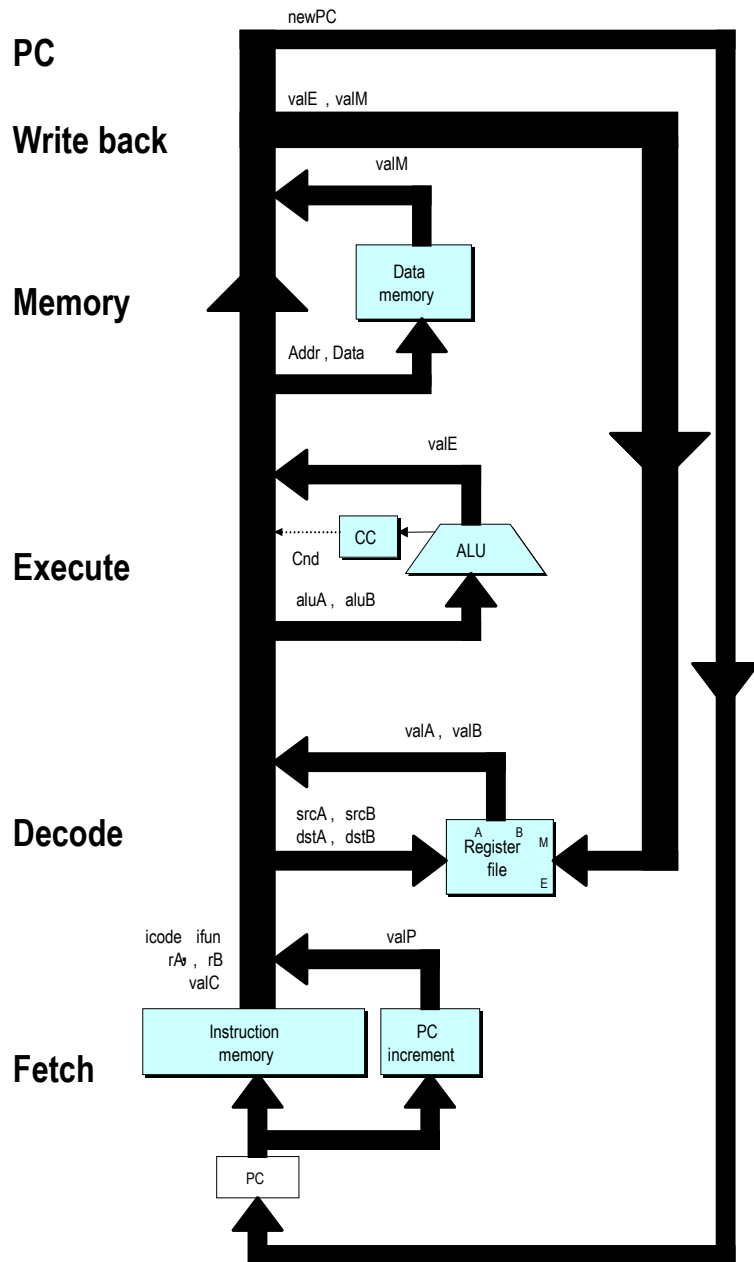
PC Stage

- Task is to select PC for current instruction
- Based on results computed by previous instruction

Processor State

- PC is no longer stored in register
- But, can determine PC based on other stored information





Pipeline Stages

■ Fetch

- Select current PC
- Read instruction
- Compute incremented PC

■ Decode

- Read program registers

■ Execute

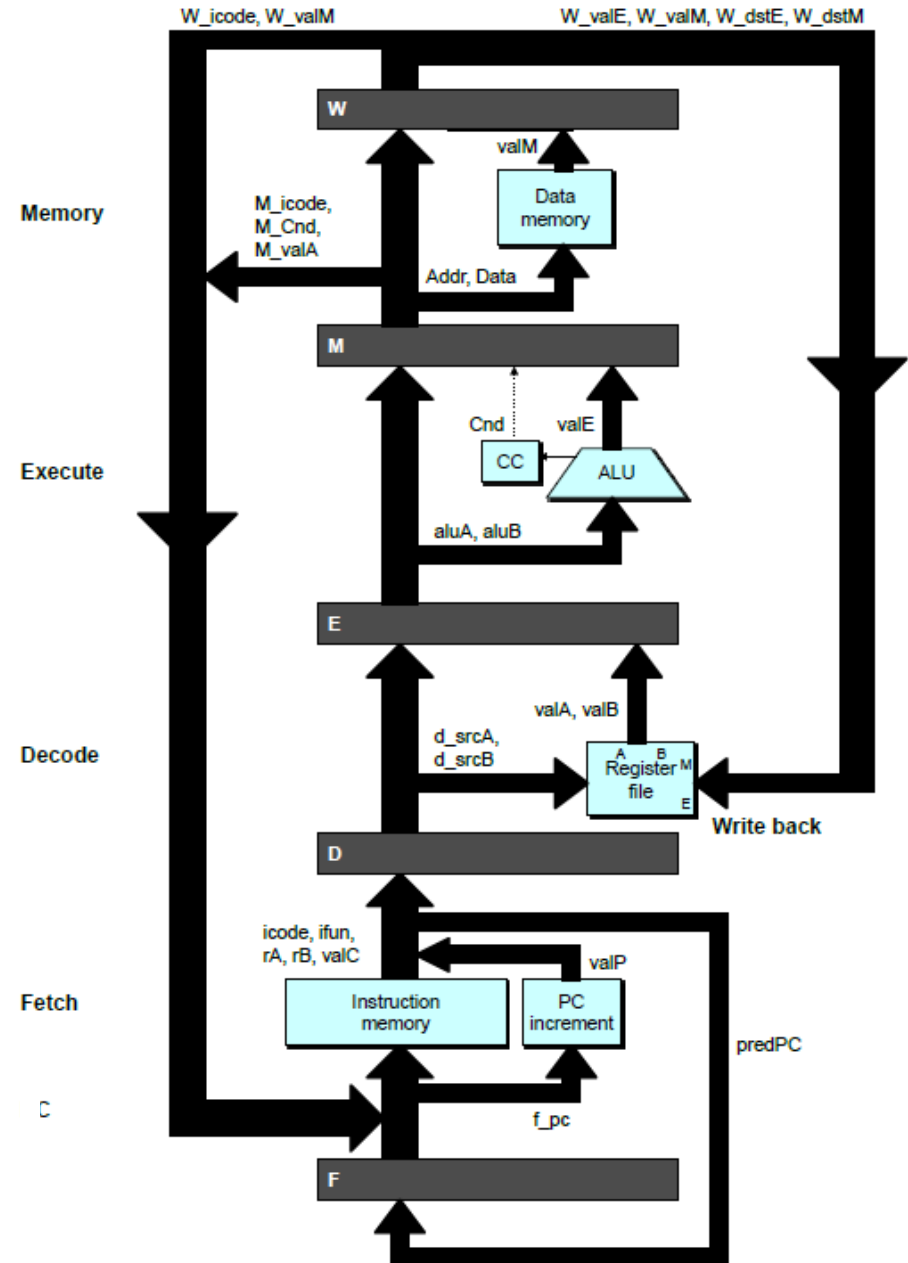
- Operate ALU

■ Memory

- Read or write data memory

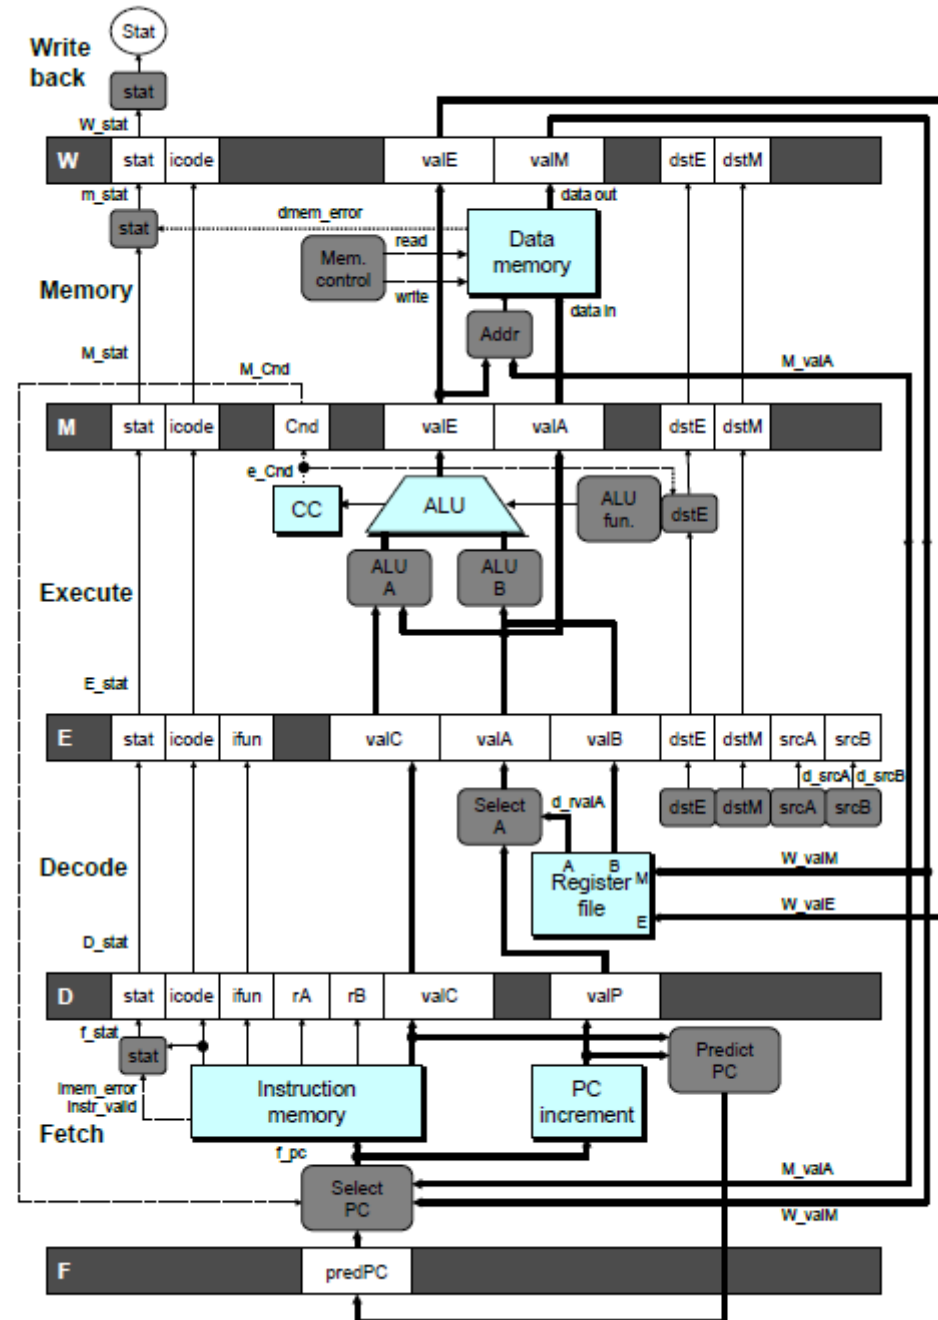
■ Write Back

- Update register file



PIPE- Hardware

- Pipeline registers hold intermediate values from instruction execution
- **Forward (Upward) Paths**
 - Values passed from one stage to next
 - Cannot jump past stages
 - e.g., valC passes through decode



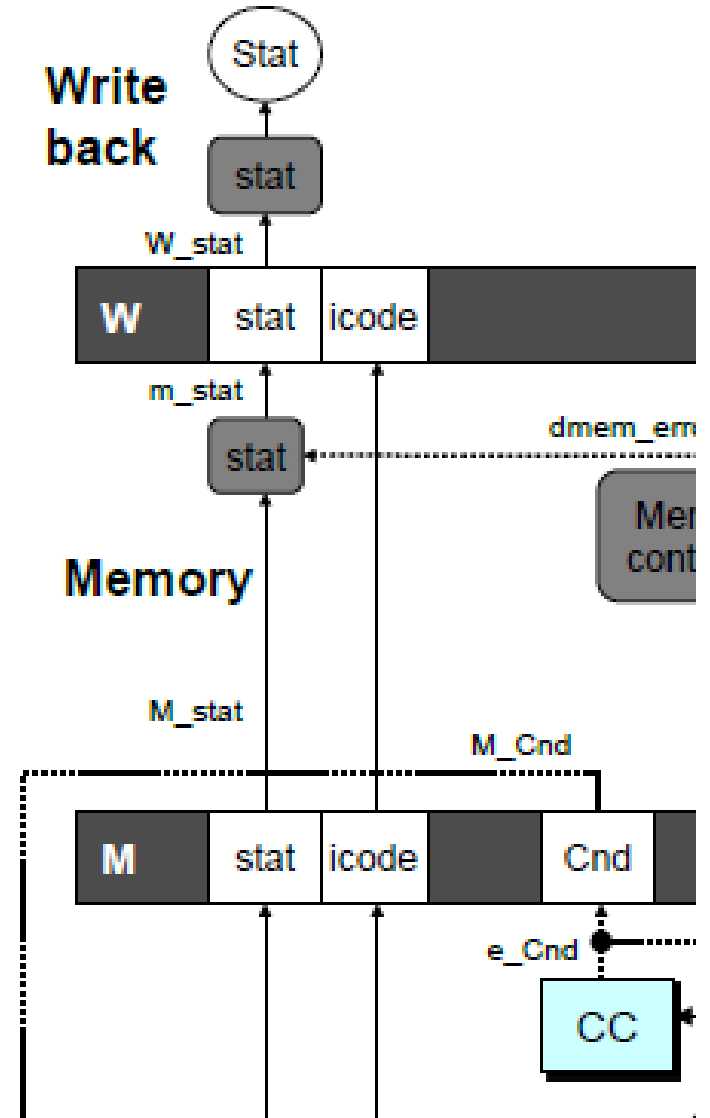
Signal Naming Conventions

■ S_Field

- Value of Field held in stage S pipeline register

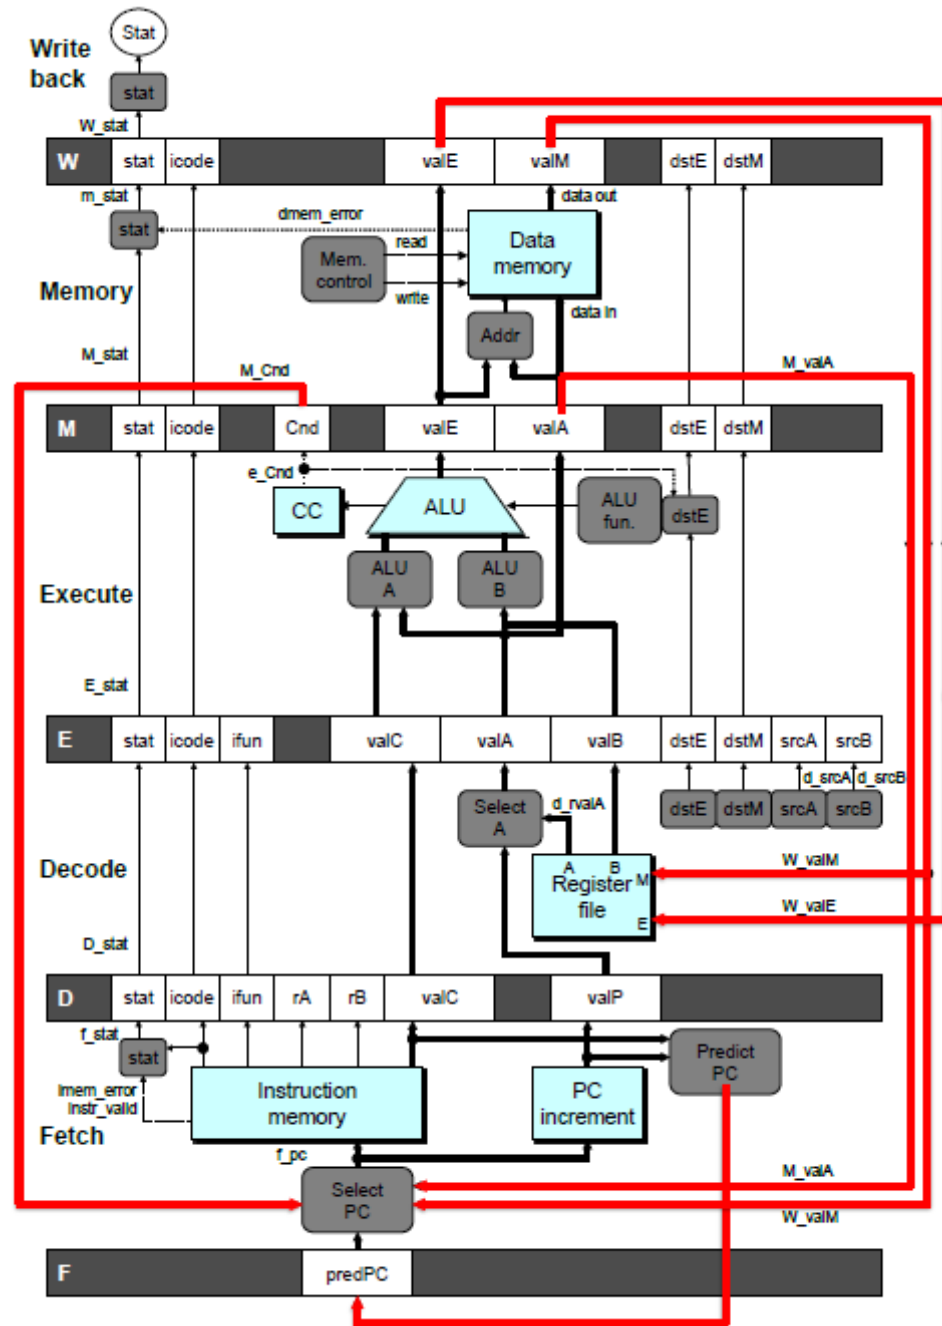
■ s_Field

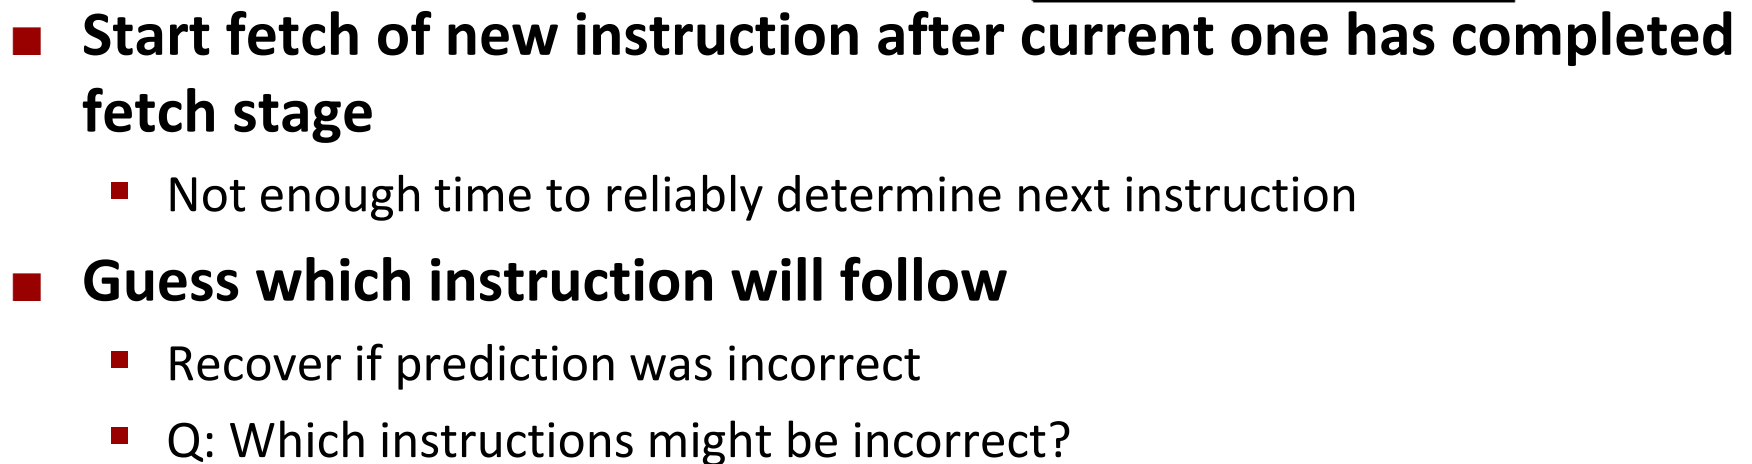
- Value of Field computed in stage S



Feedback Paths

- **Predicted PC**
 - Guess value of next PC
- **Branch information**
 - Jump taken/not-taken
 - Fall-through or target address
- **Return point**
 - Read from memory
- **Register updates**
 - To register file write ports





Our Prediction Strategy

■ Instructions that Don't Transfer Control

- Predict next PC to be valP
- Always reliable

■ Call and Unconditional Jumps

- Predict next PC to be valC (destination)
- Always reliable

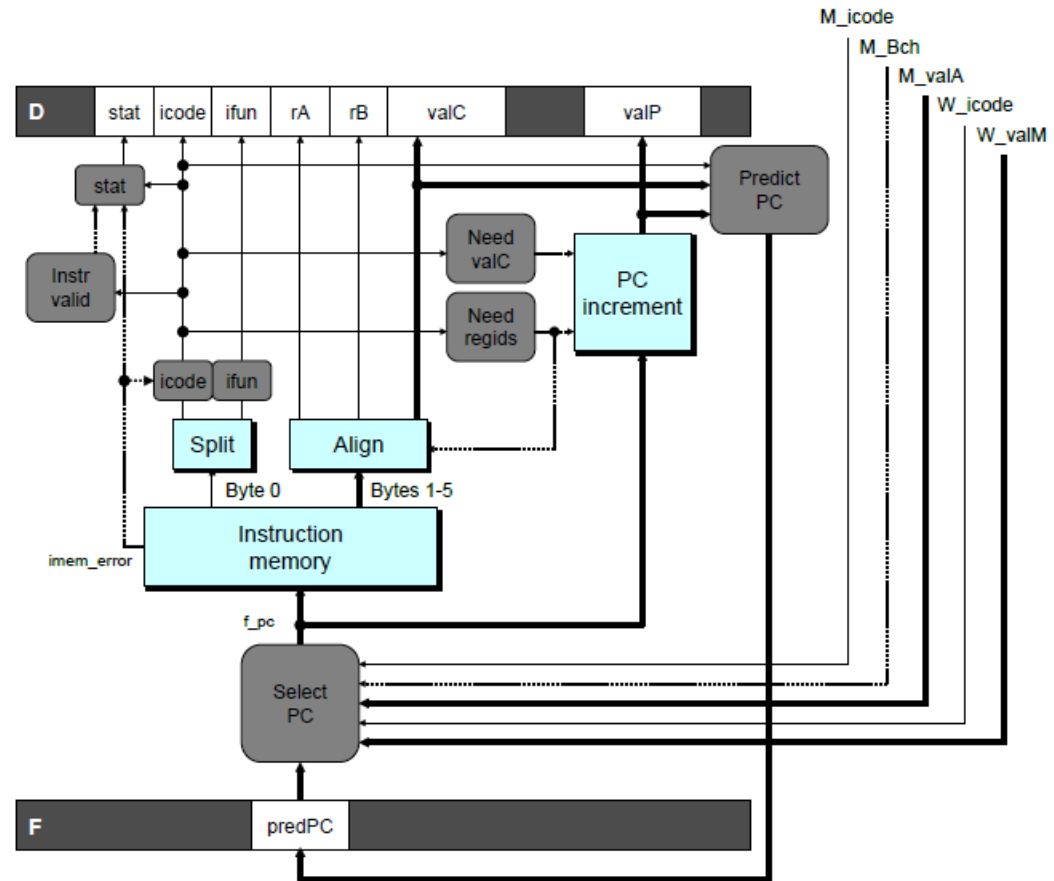
■ Conditional Jumps

- Predict next PC to be valC (destination)
- Only correct if branch is taken
 - Typically right 60% of time

■ Return Instruction

- Don't try to predict

Recovering from PC Misprediction



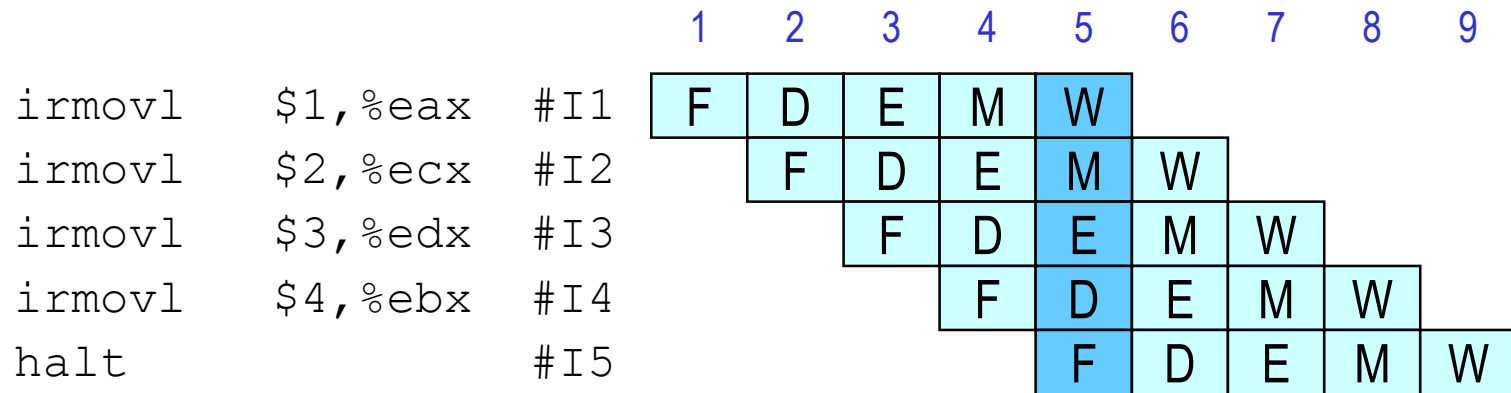
■ Mispredicted Jump

- Will see branch condition flag once instruction reaches memory stage
- Can get fall-through PC from valA (value M_valA)

■ Return Instruction

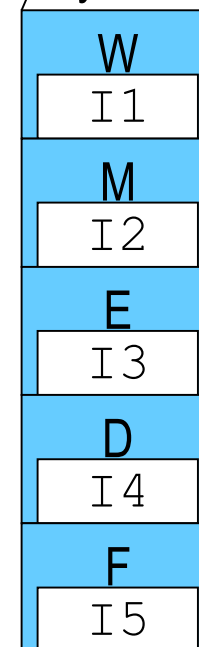
- Will get return PC when `ret` reaches write-back stage (W_valM)

Pipeline Demonstration



Cycle 5

■ File: demo-basic.ys



Data Dependencies: 3 Nop's

demo-h3.y_s

0x000: irmovl \$10,%edx

0x006: irmovl \$3,%eax

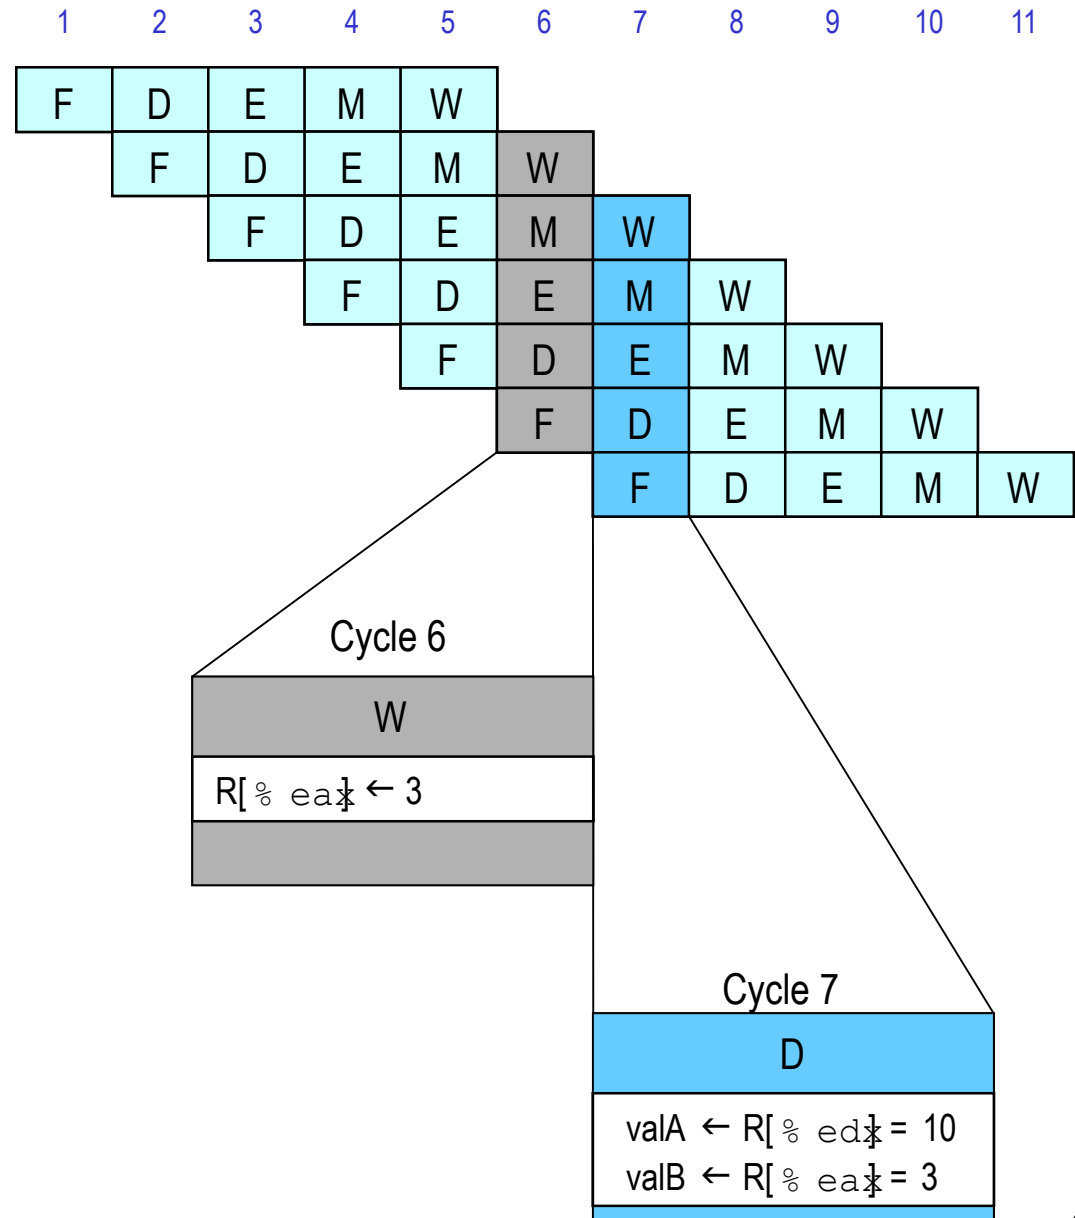
0x00c: nop

0x00d: nop

0x00e: nop

0x00f: addl %edx,%eax

0x011: halt



Data Dependencies: 2 Nop's

demo-h2.y

0x000: irmovl \$10,%edx

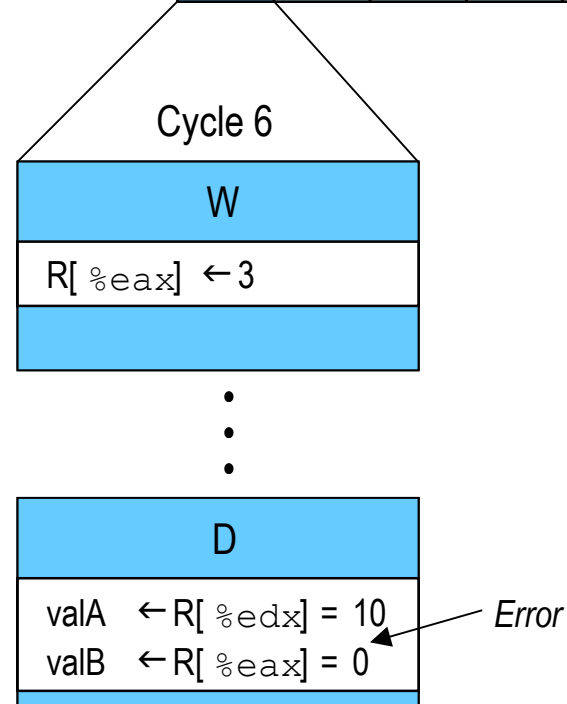
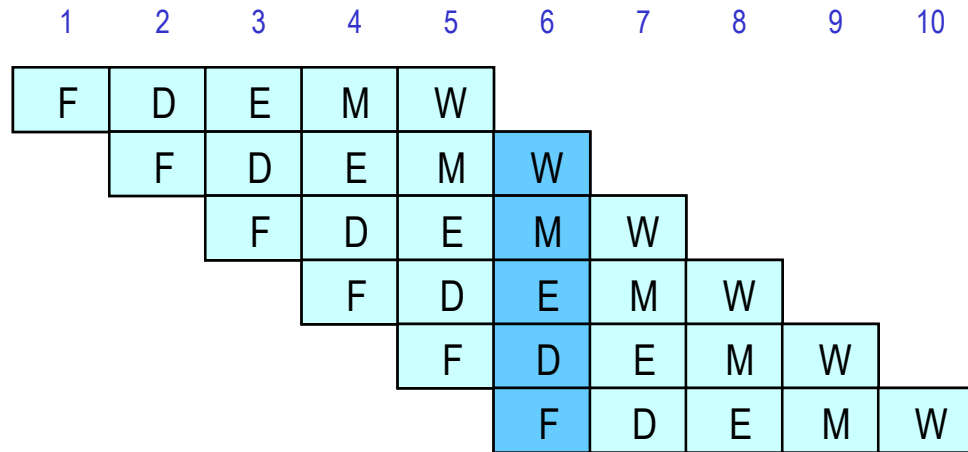
0x006: irmovl \$3,%eax

0x00c: nop

0x00d: nop

0x00e: addl %edx,%eax

0x010: halt



Data Dependencies: 1 Nop

demo-h1.y

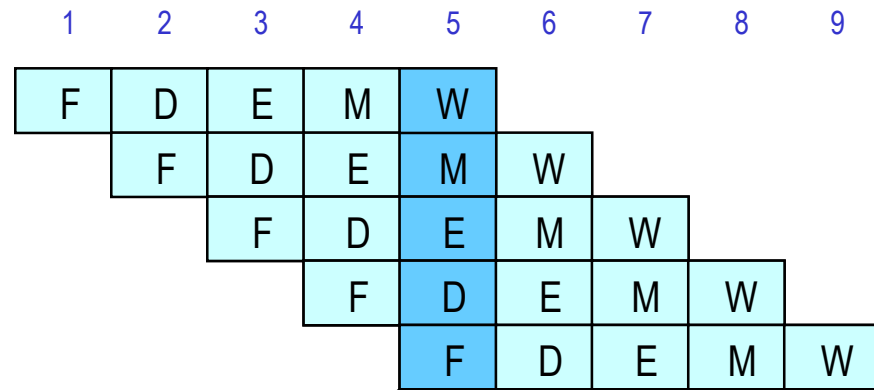
0x000: irmovl \$10,%edx

0x006: irmovl \$3,%eax

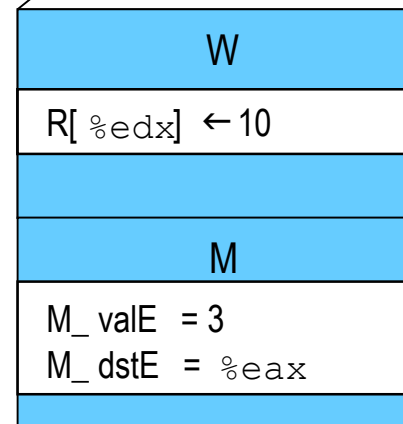
0x00c: nop

0x00d: addl %edx,%eax

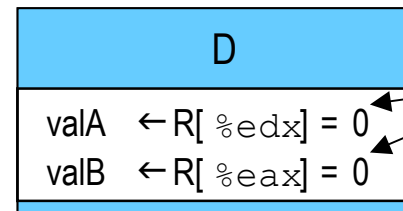
0x00f: halt



Cycle 5



•
•
•



Error

Data Dependencies: No Nop

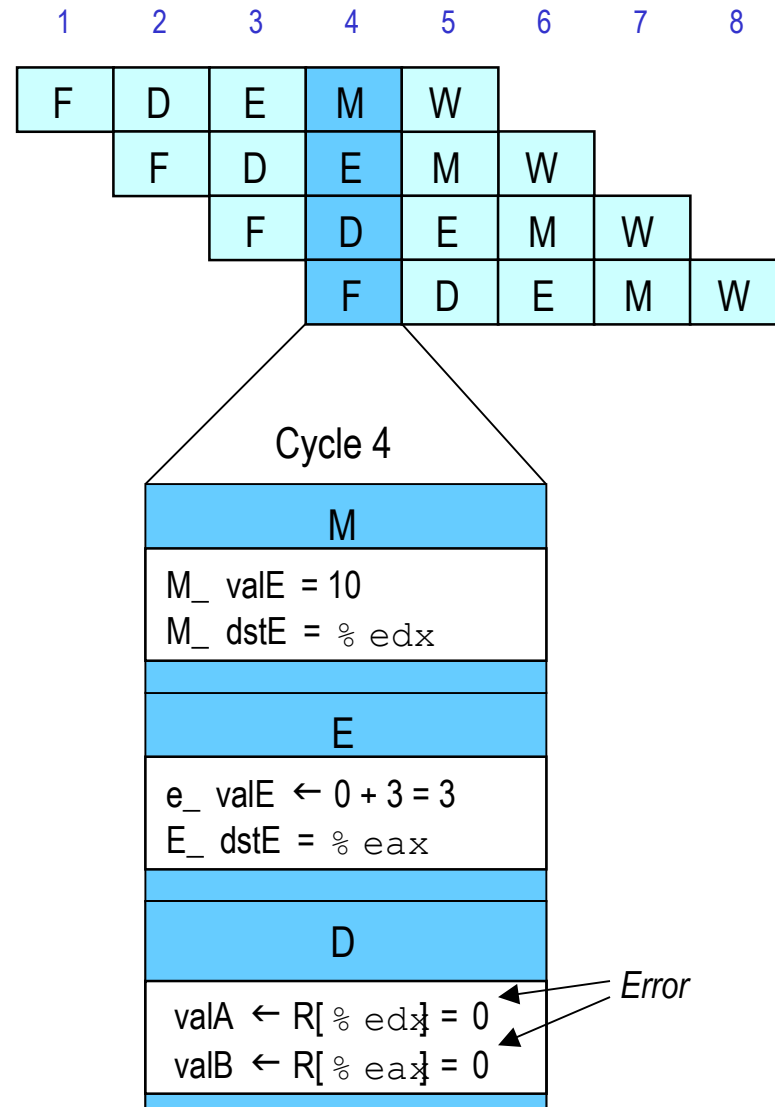
```
# demo-h0.ys
```

```
0x000: irmovl $10,%edx
```

```
0x006: irmovl $3,%eax
```

```
0x00c: addl %edx,%eax
```

```
0x00e: halt
```



Branch Misprediction Example

demo-j.ys

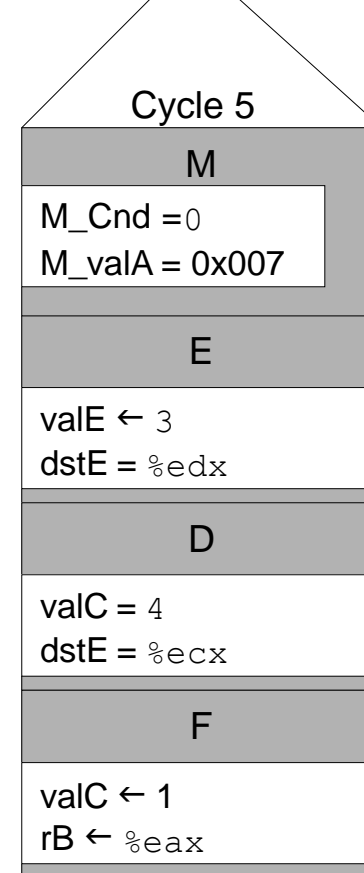
```
0x000:    xorl %eax,%eax
0x002:    jne  t                # Not taken
0x007:    irmovl $1, %eax      # Fall through
0x00d:    nop
0x00e:    nop
0x00f:    nop
0x010:    halt
0x011:  t:  irmovl $3, %edx      # Target(Should not execute)
0x017:    irmovl $4, %ecx      # Should not execute
0x01d:    irmovl $5, %edx      # Should not execute
```



Branch Misprediction Trace

# demo-j	1	2	3	4	5	6	7	8	9
0x000: xorl %eax,%eax	F	D	E	M	W				
0x002: jne t # Not taken		F	D	E	M	W			
0x011: t: irmovl \$3, %edx # Target			F	D	E	M	W		
0x017: irmovl \$4, %ecx # Target+1				F	D	E	M	W	
0x007: irmovl \$1, %eax # Fall Through					F	D	E	M	W

- Incorrectly execute two instructions at branch target



demo-ret.ys

Return Example

```

0x000:    irmovl Stack,%esp    # Initialize stack pointer
0x006:    nop                 # Avoid hazard on %esp
0x007:    nop
0x008:    nop
0x009:    call p              # Procedure call
0x00e:    irmovl $5,%esi      # Return point
0x014:    halt
0x020:    .pos 0x20
0x020: p:  nop                # procedure
0x021:    nop
0x022:    nop
0x023:    ret
0x024:    irmovl $1,%eax       # Should not be executed
0x02a:    irmovl $2,%ecx       # Should not be executed
0x030:    irmovl $3,%edx       # Should not be executed
0x036:    irmovl $4,%ebx       # Should not be executed
0x100:    .pos 0x100
0x100:    Stack:              # Stack: Stack pointer

```

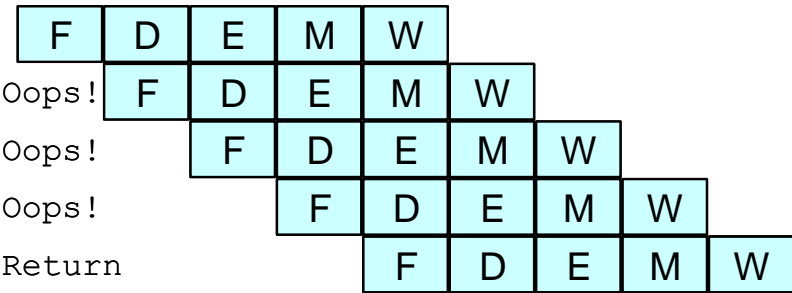
- Require lots of nops to avoid data hazards

Incorrect Return Example

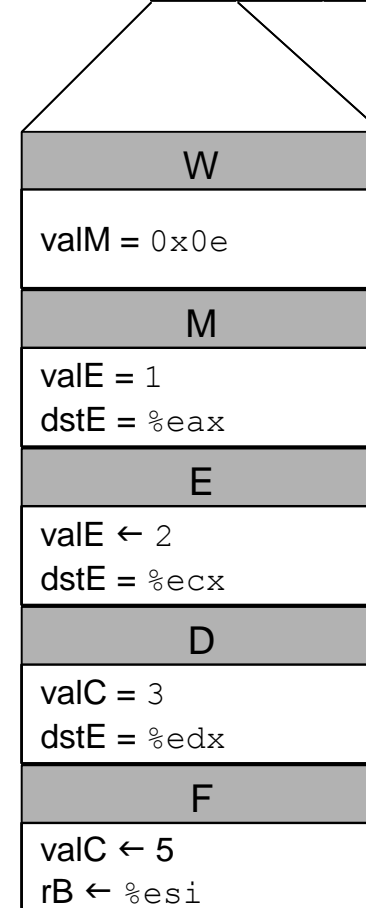
demo-ret

```

0x023:    ret
0x024:    irmovl $1,%eax # Oops!
0x02a:    irmovl $2,%ecx # Oops!
0x030:    irmovl $3,%edx # Oops!
0x00e:    irmovl $5,%esi # Return
  
```



- Incorrectly execute 3 instructions following `ret`



Pipeline Part 1: Summary

■ Concept

- Break instruction execution into 5 stages
- Run instructions through in pipelined mode

■ Limitations

- Can't handle dependencies between instructions when instructions follow too closely
- Data dependencies
 - One instruction writes register, later one reads it
- Control dependency
 - Instruction sets PC in way that pipeline did not predict correctly
 - Mispredicted branch and return

■ Fixing the Pipeline

- We'll do that next time

Pipeline Part 2: Overview

Make the pipelined processor work!

■ Data Hazards

- Instruction having register R as source follows shortly after instruction having register R as destination
- Common condition, don't want to slow down pipeline

■ Control Hazards

- Mispredict conditional branch
 - Our design predicts all branches as being taken
 - Naïve pipeline executes two extra instructions
- Getting return address for `ret` instruction
 - Naïve pipeline executes three extra instructions

■ Making Sure It Really Works

- What if multiple special cases happen simultaneously?

Pipeline Stages

■ Fetch

- Select current PC
- Read instruction
- Compute incremented PC

■ Decode

- Read program registers

■ Execute

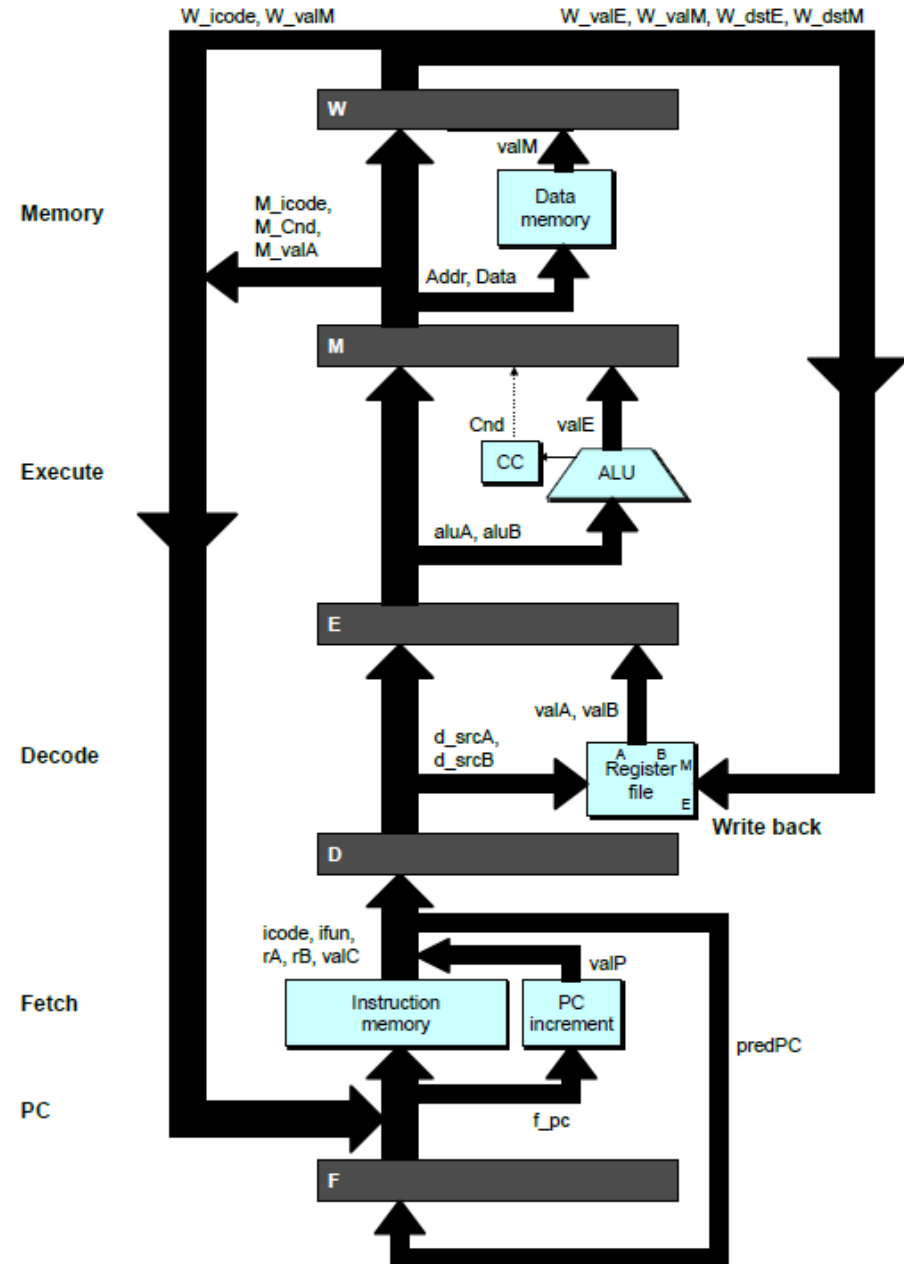
- Operate ALU

■ Memory

- Read or write data memory

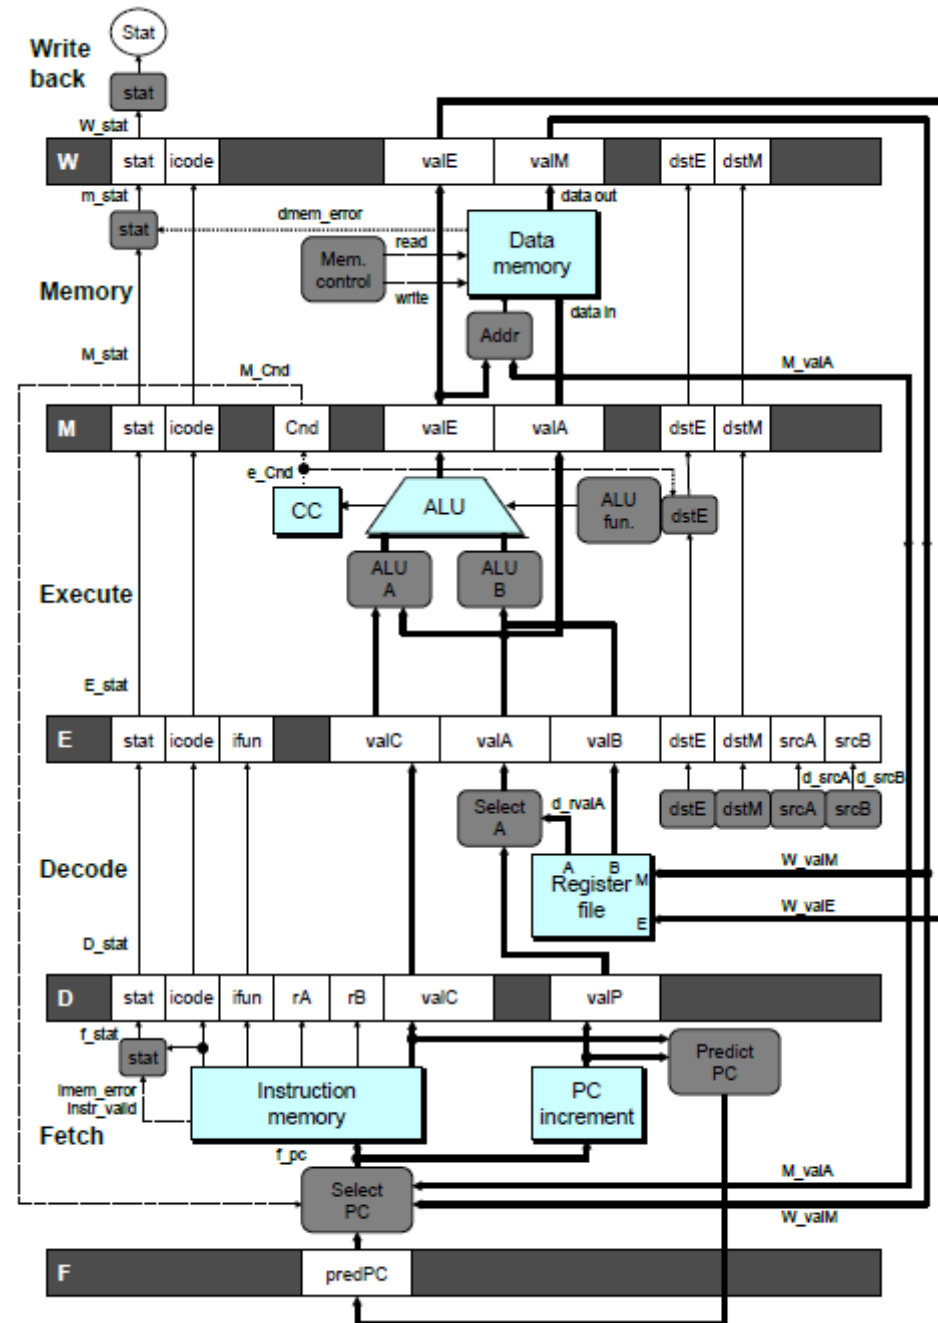
■ Write Back

- Update register file



PIPE- Hardware

- Pipeline registers hold intermediate values from instruction execution
- **Forward (Upward) Paths**
 - Values passed from one stage to next
 - Cannot jump past stages
 - e.g., valC passes through decode



Data Dependencies: 2 Nop's

```
# demo-h2.y
```

```
0x000: irmovl $10,%edx
```

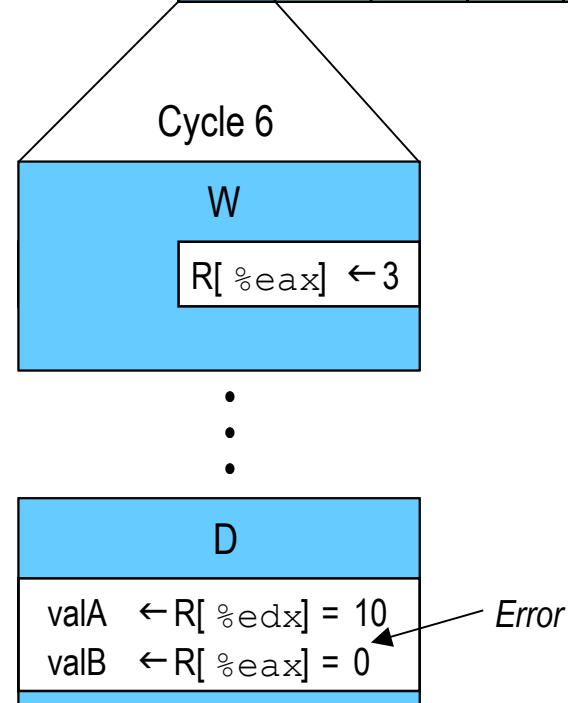
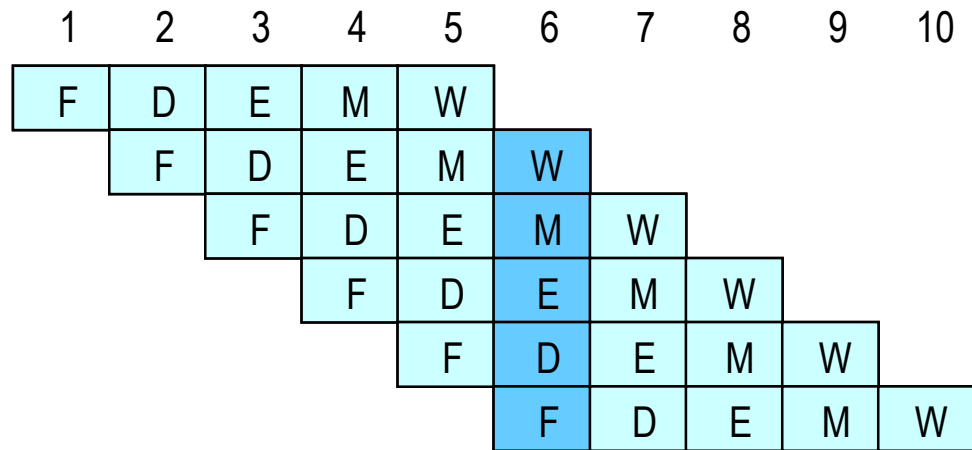
```
0x006: irmovl $3,%eax
```

```
0x00c: nop
```

```
0x00d: nop
```

```
0x00e: addl %edx,%eax
```

```
0x010: halt
```



Data Dependencies: No Nop

```
# demo-h0.y
```

```
0x000: irmovl $10,%edx
```

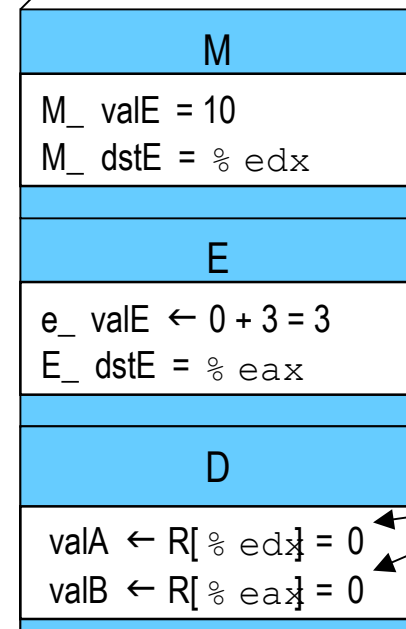
```
0x006: irmovl $3,%eax
```

```
0x00c: addl %edx,%eax
```

```
0x00e: halt
```

1	2	3	4	5	6	7	8
F	D	E	M	W			
	F	D	E	M	W		
		F	D	E	M	W	
			F	D	E	M	W

Cycle 4



Error

Stalling for Data Dependencies

demo-h2.y

0x000: irmovl \$10,%edx

0x006: irmovl \$3,%eax

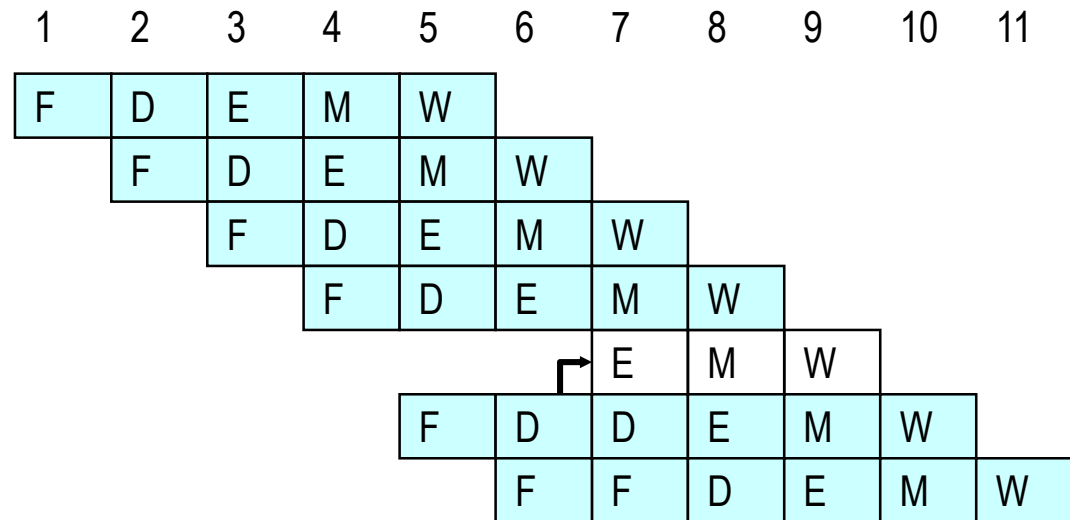
0x00c: nop

0x00d: nop

bubble

0x00e: addl %edx,%eax

0x010: halt



- If instruction follows too closely after one that writes register, slow it down
- Hold instruction in decode
- Dynamically inject nop into execute stage

Stall Condition

■ Source Registers

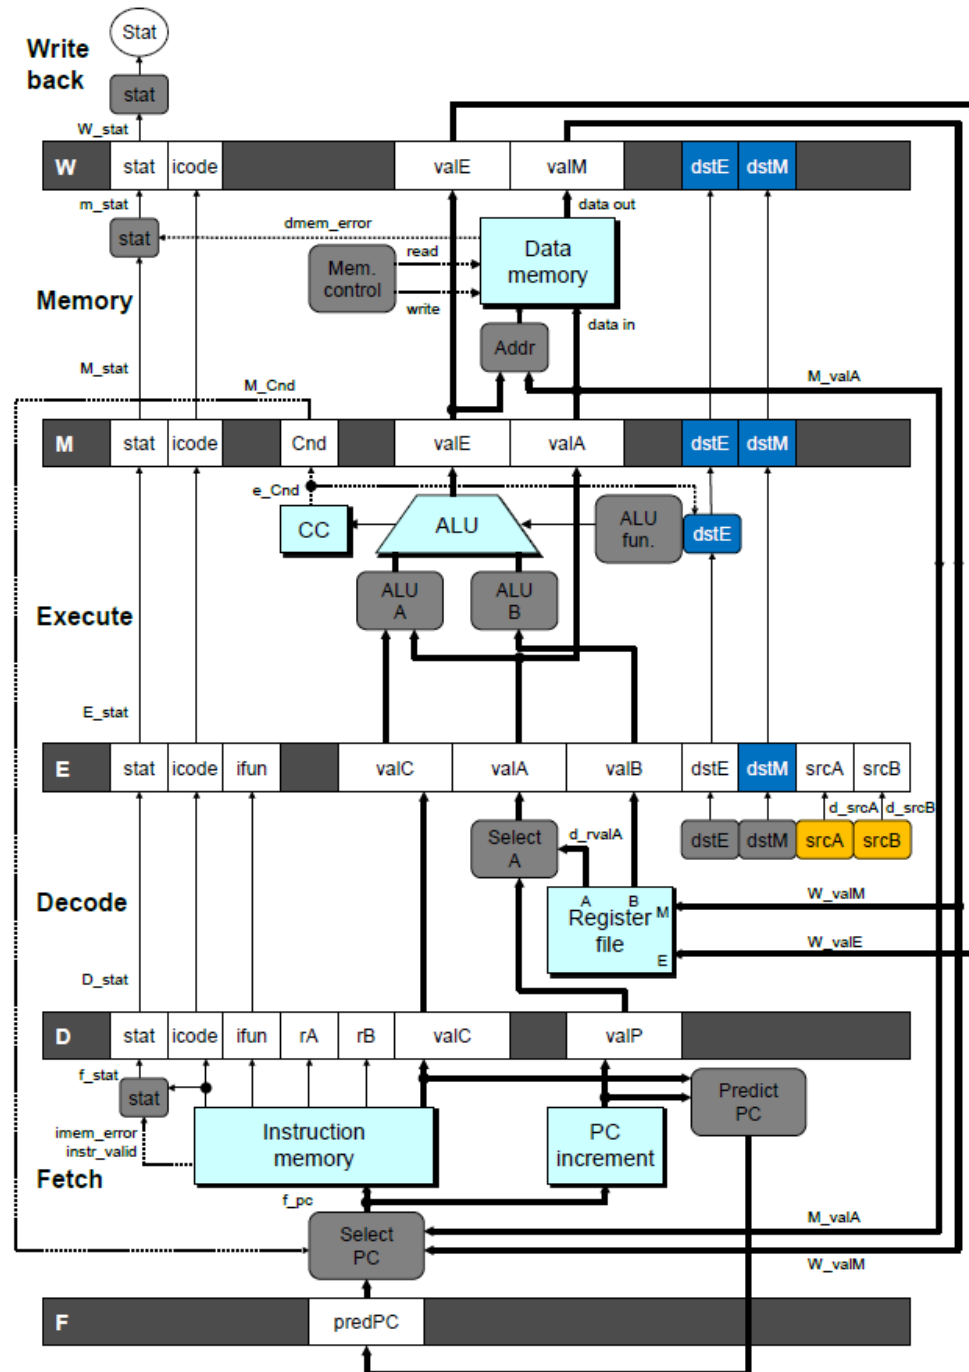
- srcA and srcB of current instruction in decode stage

■ Destination Registers

- dstE and dstM fields
- Instructions in execute, memory, and write-back stages

■ Special Case

- Don't stall for register ID 15 (0xF)
 - Indicates absence of register operand
- Don't stall for failed conditional move



Detecting Stall Condition

```
# demo-h2.y
```

```
0x000: irmovl $10,%edx
```

```
0x006: irmovl $3,%eax
```

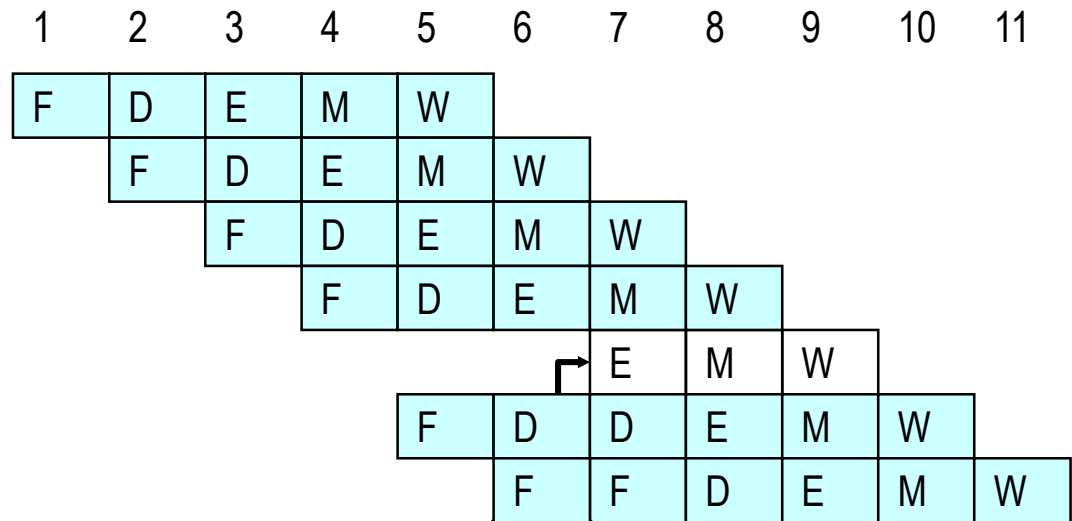
```
0x00c: nop
```

```
0x00d: nop
```

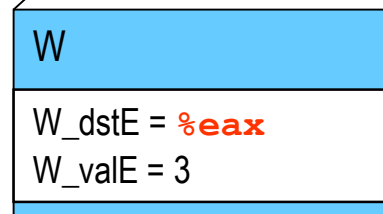
bubble

```
0x00e: addl %edx,%eax
```

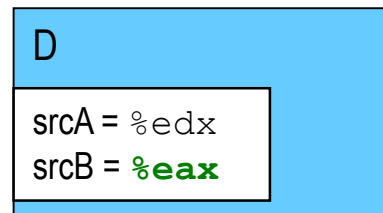
```
0x010: halt
```



Cycle 6



•
•
•



Stalling X3

```
# demo-h0.y
```

```
0x000: irmovl $10,%edx
```

```
0x006: irmovl $3,%eax
```

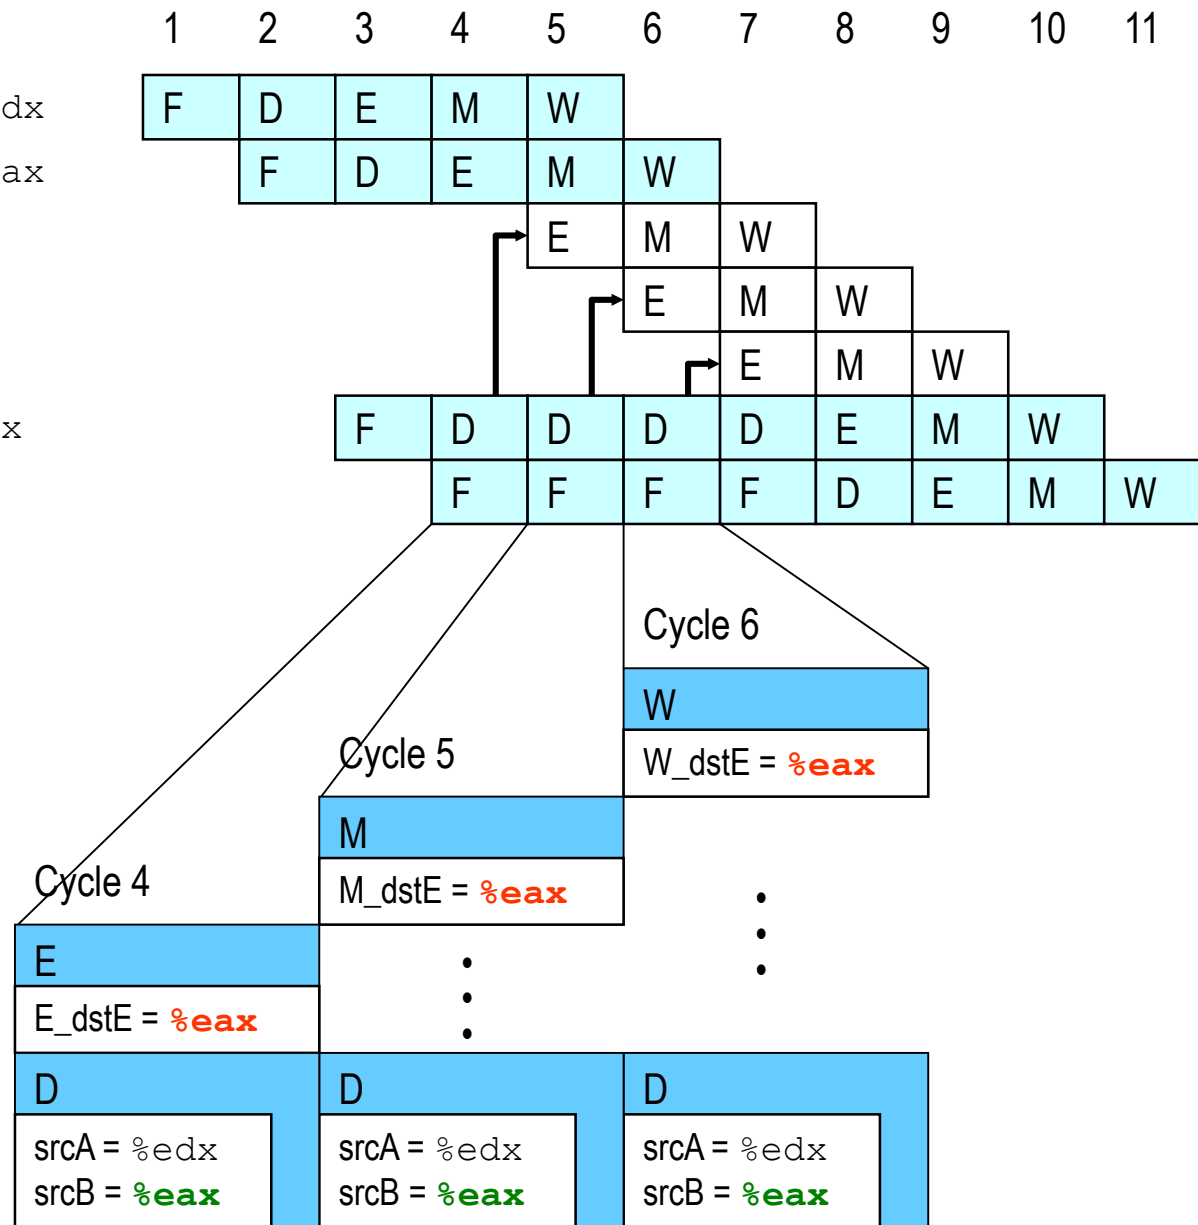
bubble

bubble

bubble

```
0x00c: addl %edx,%eax
```

```
0x00e: halt
```



What Happens When Stalling?

```
# demo-h0.y
```

```
0x000: irmovl $10,%edx
```

```
0x006: irmovl $3,%eax
```

```
0x00c: addl %edx,%eax
```

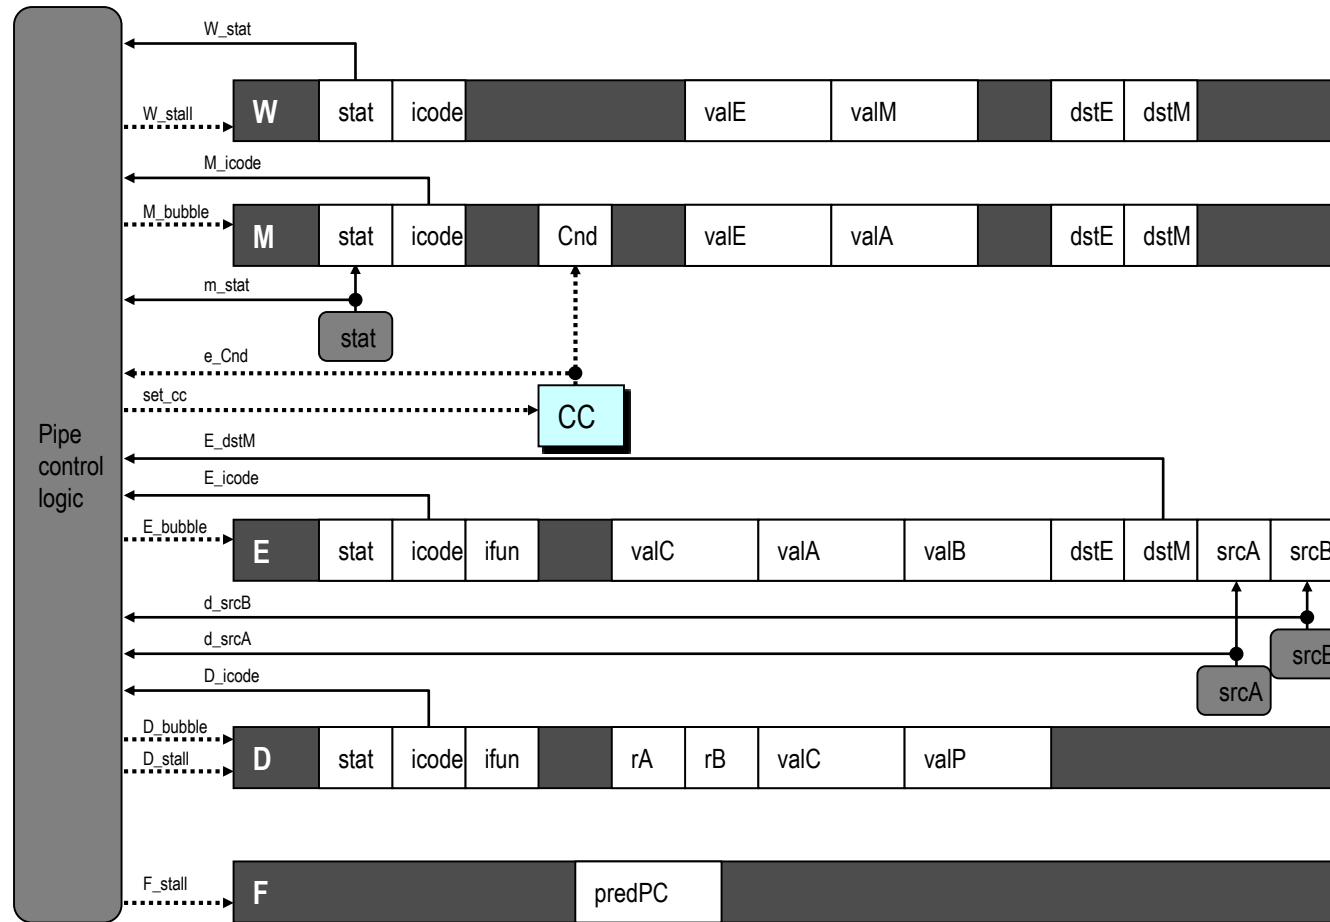
```
0x00e: halt
```

Cycle 8

Write Back	<i>bubble</i>
Memory	<i>bubble</i>
Execute	0x00c: addl %edx,%eax
Decode	0x00e: halt
Fetch	

- Stalling instruction held back in decode stage
- Following instruction stays in fetch stage
- Bubbles injected into execute stage
 - Like dynamically generated nop's
 - Move through later stages

Implementing Stalling

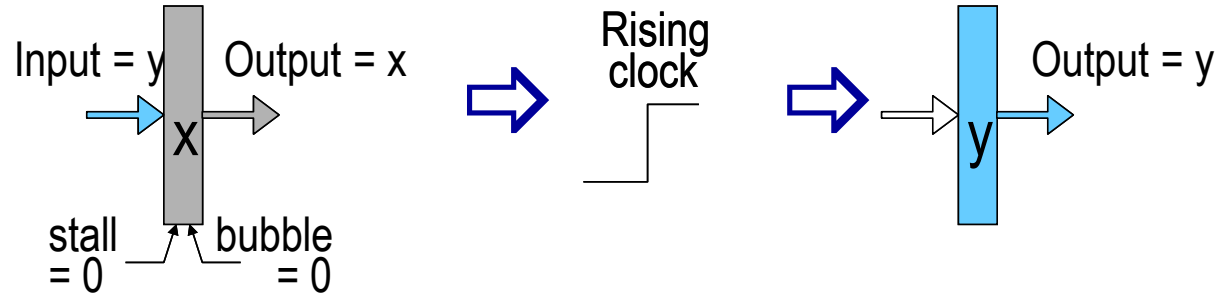


■ Pipeline Control

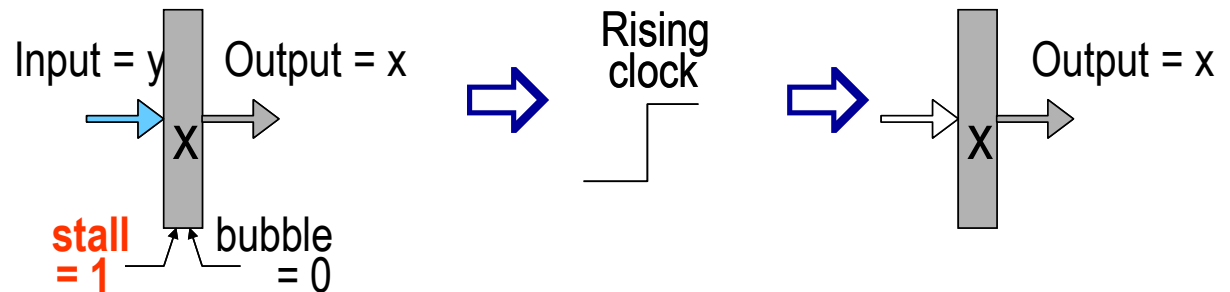
- Combinational logic detects stall condition
- Sets mode signals for how pipeline registers should update

Pipeline Register Modes

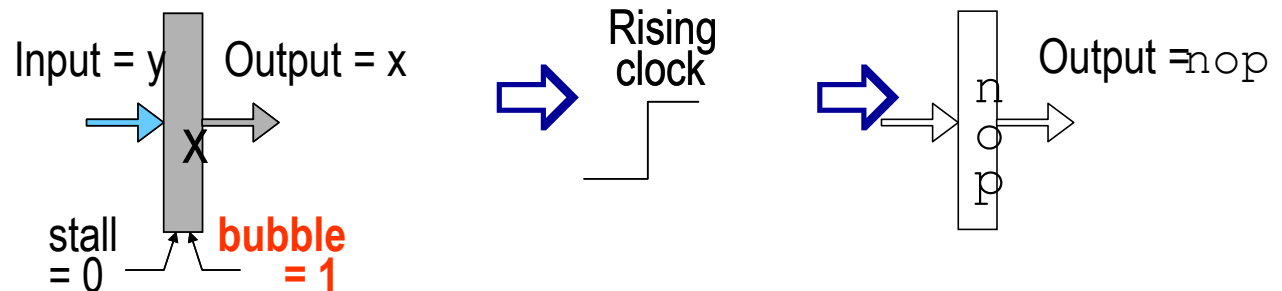
Normal



Stall



Bubble



Data Forwarding

■ Naïve Pipeline

- Register isn't written until completion of write-back stage
- Source operands read from register file in decode stage
 - Needs to be in register file at start of stage

■ Observation

- Value generated in execute or memory stage

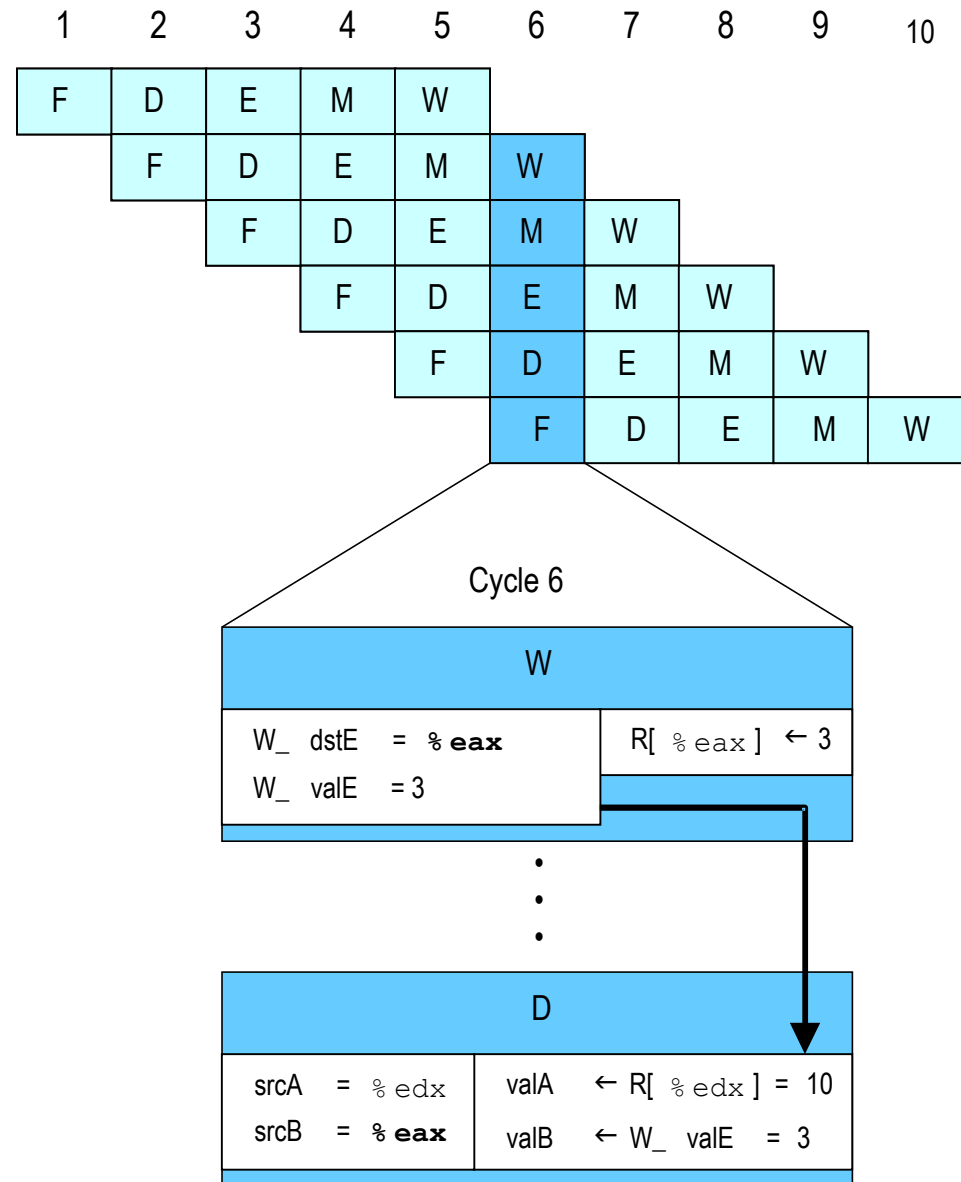
■ Trick

- Pass value directly from generating instruction to decode stage
- Needs to be available at end of decode stage

Data Forwarding Example

```
# demo-h2.y
0x000:    irmovl  $10,%  edx
0x006:    irmovl  $3,%   eax
0x00c:    nop
0x00d:    nop
0x00e:    addl   %edx,%  eax
0x010:    halt
```

- **irmovl** in write-back stage
- Destination value in W pipeline register
- Forward as valB for decode stage



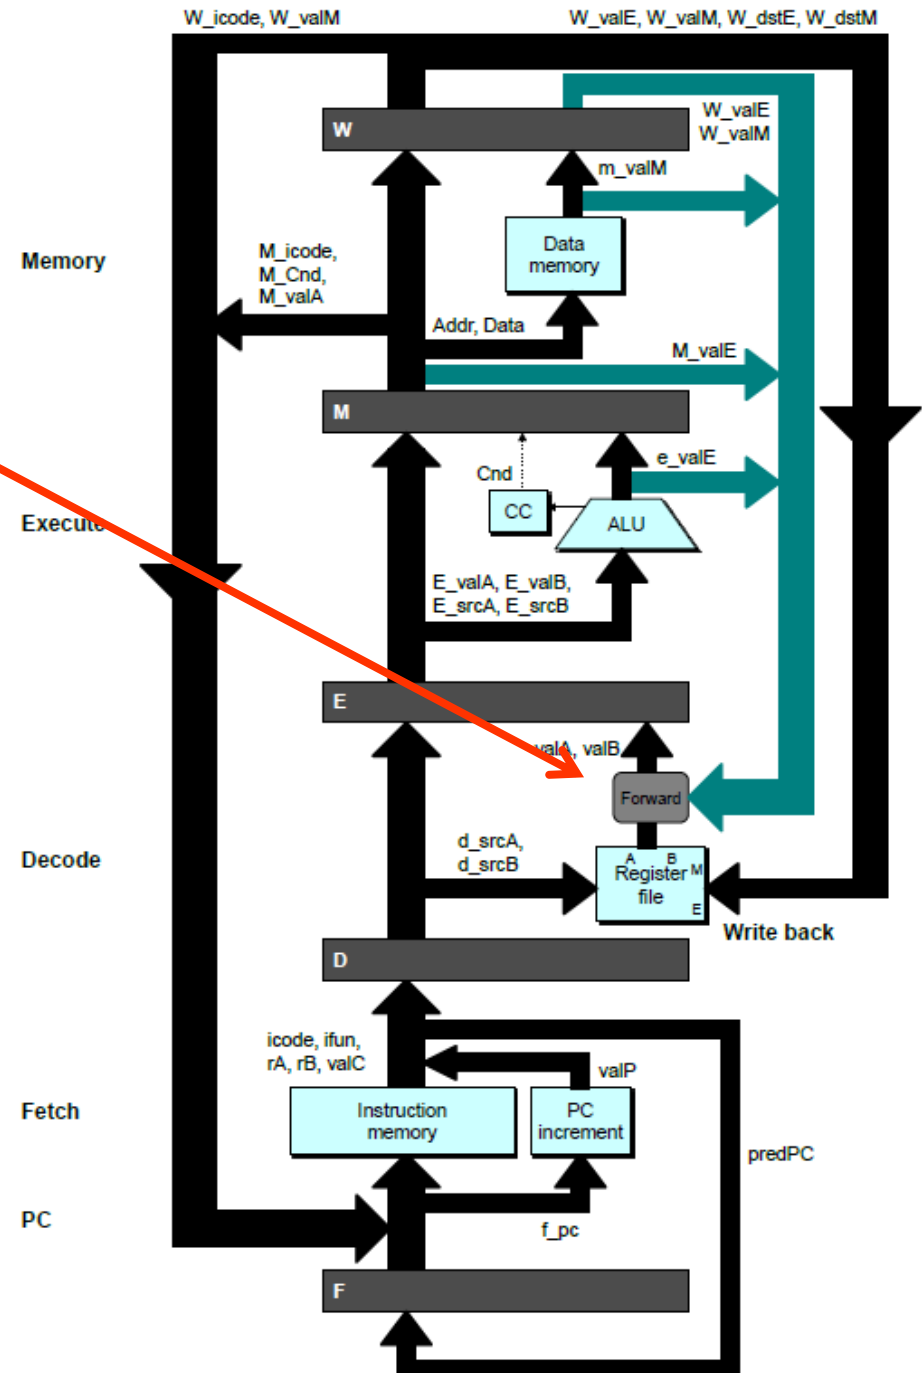
Bypass Paths

■ Decode Stage

- Forwarding logic selects valA and valB
- Normally from register file
- Forwarding: get valA or valB from later pipeline stage

■ Forwarding Sources

- Execute: valE
- Memory: valE, valM
- Write back: valE, valM



Data Forwarding Example #2

```
# demo-h0.ys
```

```
0x000: irmovl $10,%edx
```

```
0x006: irmovl $3,%eax
```

```
0x00c: addl %edx,%eax
```

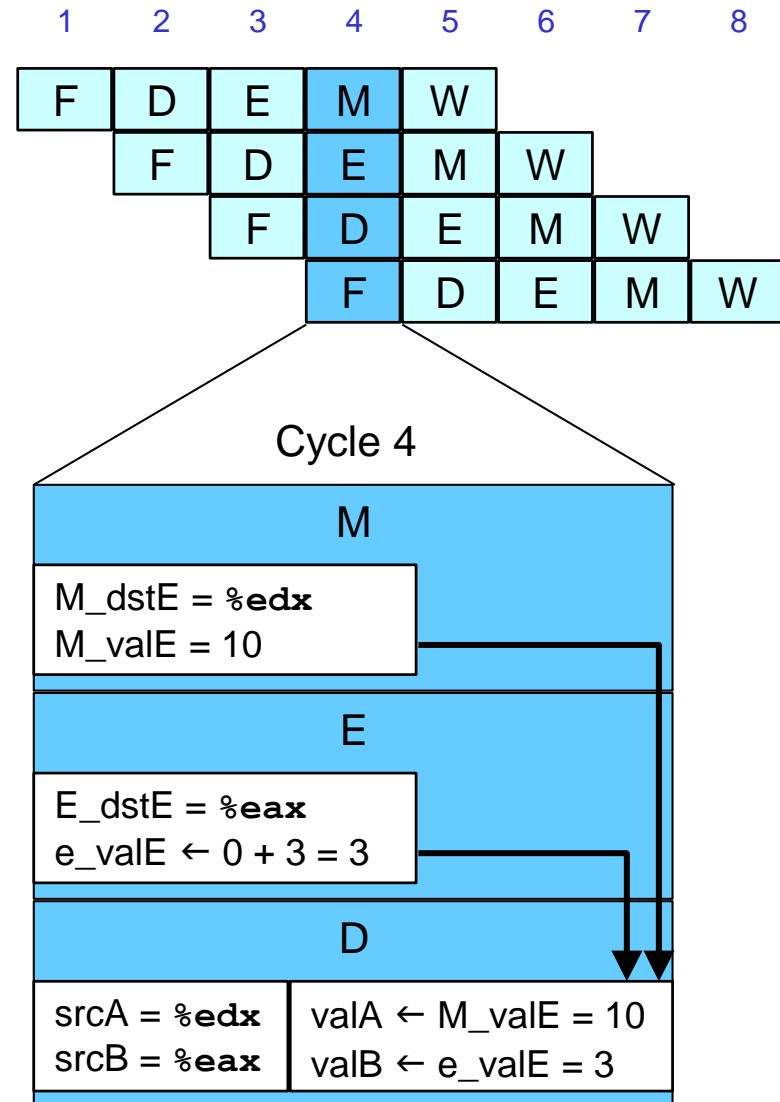
```
0x00e: halt
```

■ Register **%edx**

- Generated by ALU during previous cycle
- Forward from memory as valA

■ Register **%eax**

- Value just generated by ALU
- Forward from execute as valB



Forwarding Priority

```
# demo-priority.js
```


```
0x000: irmovl $1, %eax
```

```
0x006: irmovl $2, %eax
```

```
0x00c: irmovl $3, %eax
```

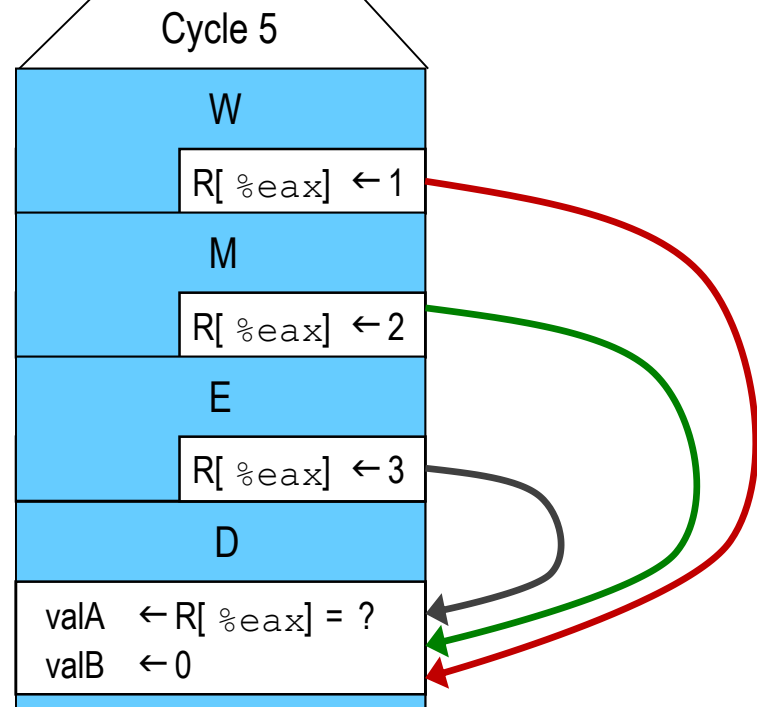
```
0x012: rrmovl %eax, %edx
```

```
0x014: halt
```

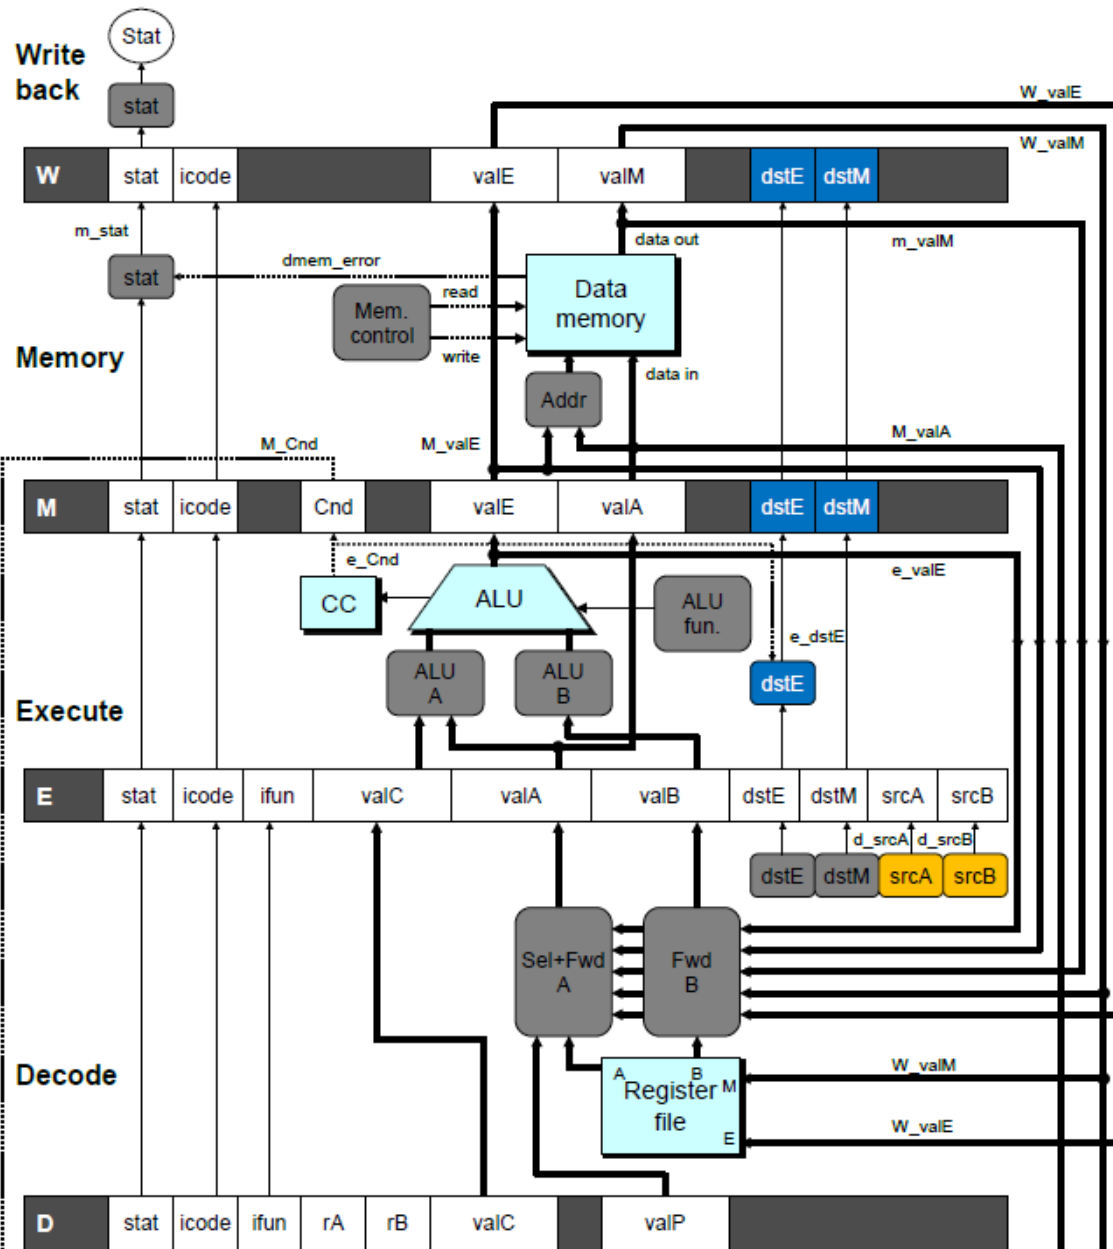


■ Multiple Forwarding Choices

- Which one should have priority?
- Use matching value from earliest pipeline stage
- Match sequential semantics

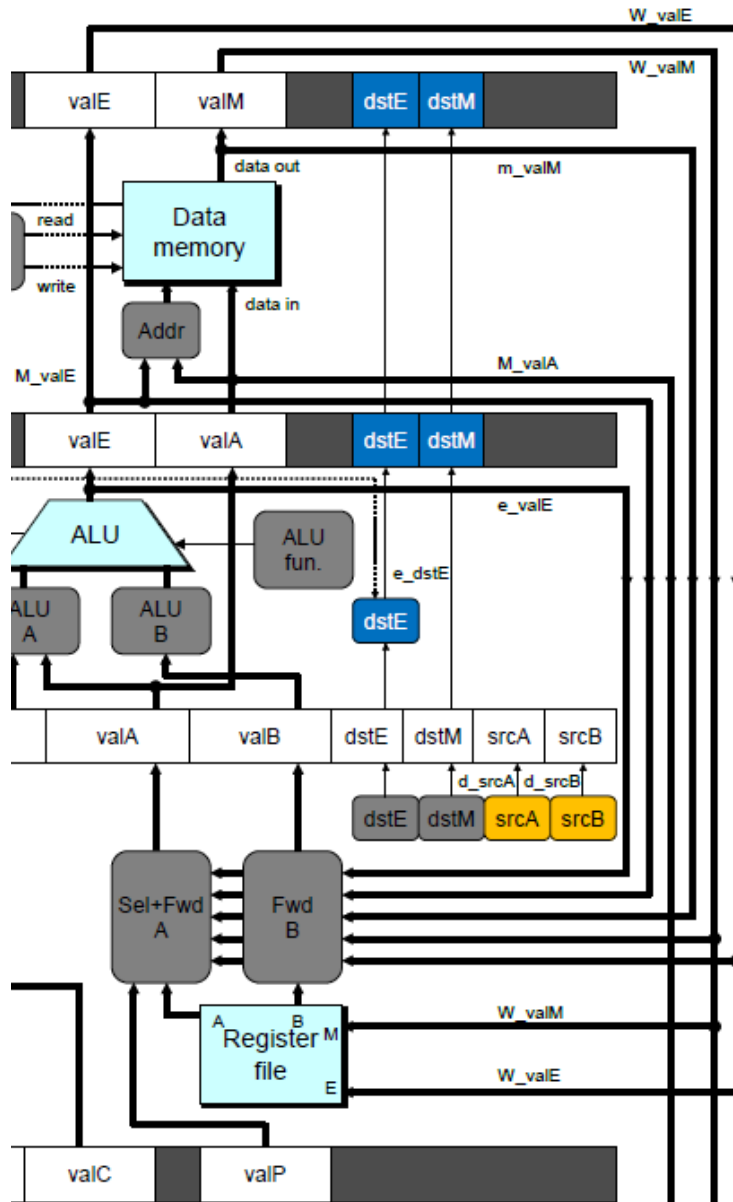


Implementing Forwarding



- Add additional feedback paths from E, M, and W pipeline registers into decode stage
- Create logic blocks to select from multiple sources for `valA` and `valB` in decode stage


Implementing Forwarding



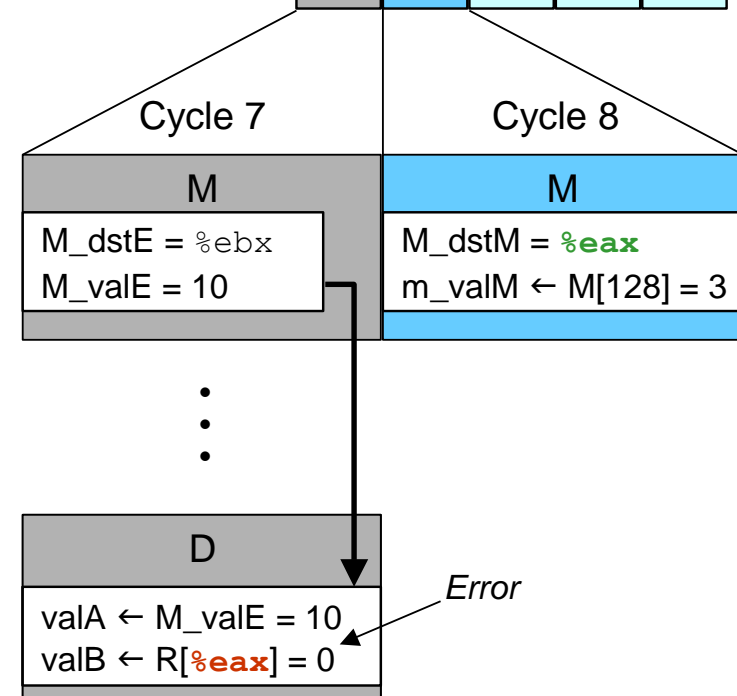
```
## What should be the A value?
int new_E_valA = [
    # Use incremented PC
    D_icode in { ICALL, IJXX } : D_valP;
    # Forward valE from execute
    d_srcA == e_dstE : e_valE;
    # Forward valM from memory
    d_srcA == M_dstM : m_valM;
    # Forward valE from memory
    d_srcA == M_dstE : M_valE;
    # Forward valM from write back
    d_srcA == W_dstM : W_valM;
    # Forward valE from write back
    d_srcA == W_dstE : W_valE;
    # Use value read from register file
    1 : d_rvalA;
];
```

```
0x020: halt
```

1 2 3 4 5 6 7 8 9 10 11



- Value needed by end of decode stage in cycle 7
- Value read from memory in memory stage of cycle 8



Avoiding Load/Use Hazard

demo-luh.js

0x000: irmovl \$128,%edx

0x006: irmovl \$3,%ecx

0x00c: rmmovl %ecx, 0(%edx)


0x012: irmovl \$10,%ebx

0x018: mrmovl 0(%edx),%eax # Load %eax

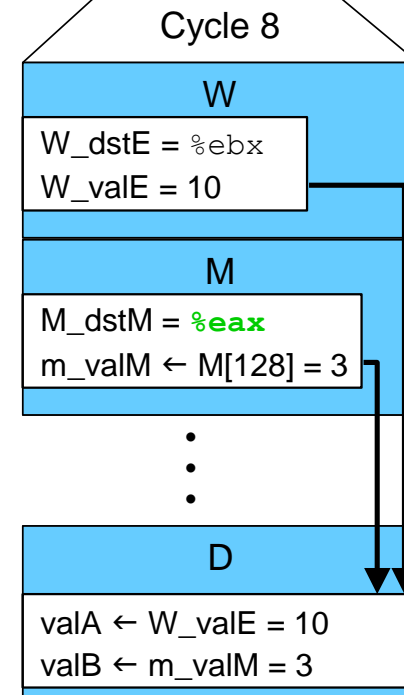
bubble

0x01e: addl %ebx,%eax # Use %eax

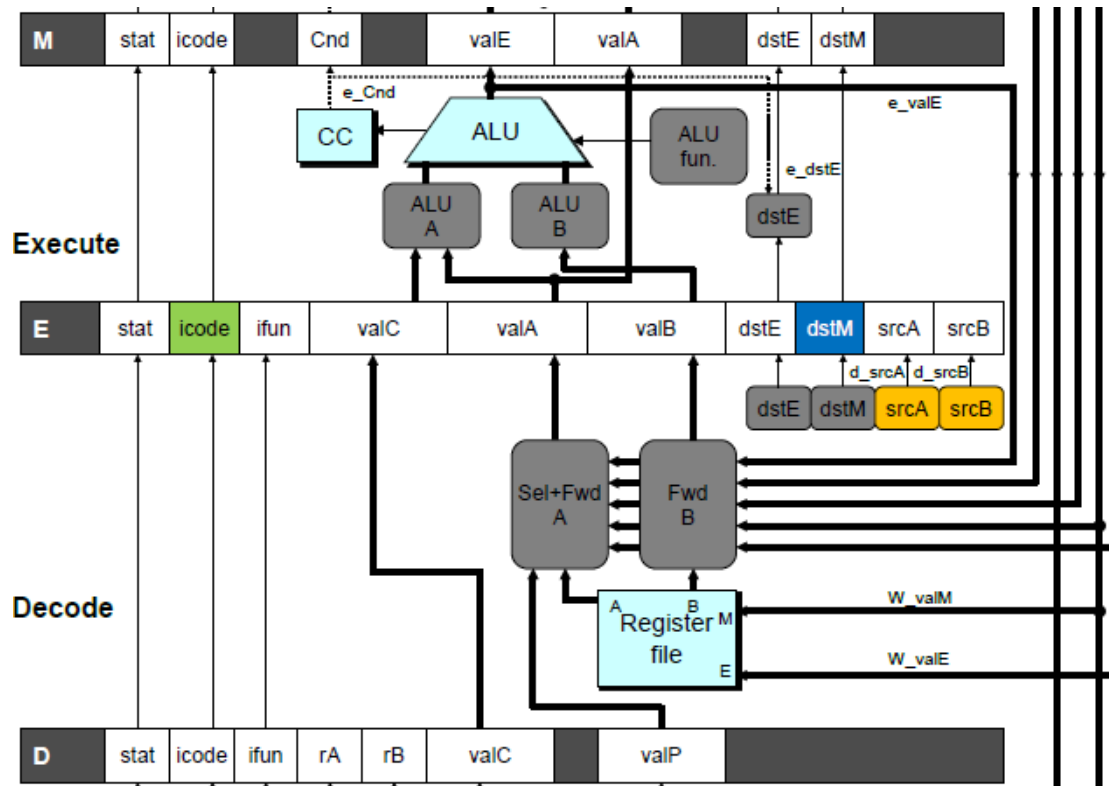
0x020: halt



- Stall using instruction for one cycle
- Can then pick up loaded value by forwarding from memory stage



Detecting Load/Use Hazard



Condition	Trigger
Load/Use Hazard	$E_icode \in \{ IMRMOVL, IPOPL \} \ \&\&$ $E_dstM \in \{ d_srcA, d_srcB \}$

Control for Load/Use Hazard

demo-luh.js

0x000: irmovl \$128,%edx

0x006: irmovl \$3,%ecx

0x00c: rmmovl %ecx, 0(%edx)

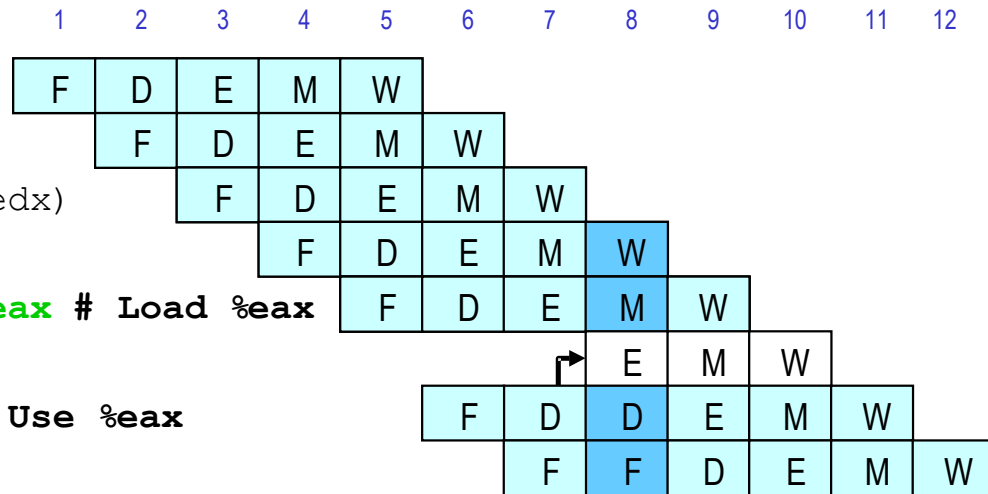
0x012: irmovl \$10,%ebx

0x018: mrmovl 0(%edx),%eax # Load %eax

bubble

0x01e: addl %ebx,%eax # Use %eax

0x020: halt



- Stall instructions in fetch and decode stages
- Inject bubble into execute stage

Condition	F	D	E	M	W
Load/Use Hazard	stall	stall	bubble	normal	normal

Branch Misprediction Example

demo-j.y

```

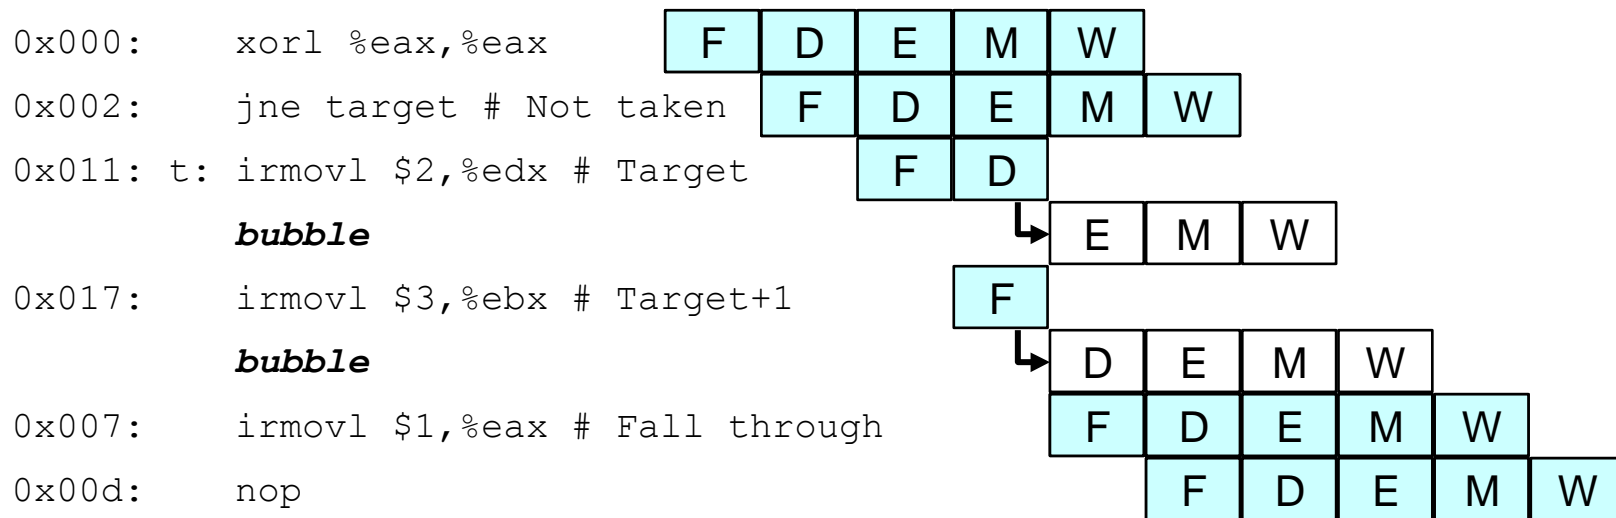
0x000:      xorl %eax,%eax
0x002:      jne  t                # Not taken
0x007:      irmovl $1, %eax      # Fall through
0x00d:      nop
0x00e:      nop
0x00f:      nop
0x010:      halt
0x011:  t:  irmovl $3, %edx      # Target (Should not
execute)
0x017:      irmovl $4, %ecx      # Should not execute
0x01d:      irmovl $5, %edx      # Should not execute

```

- Should only execute first 8 instructions

Handling Misprediction

demo-j.js



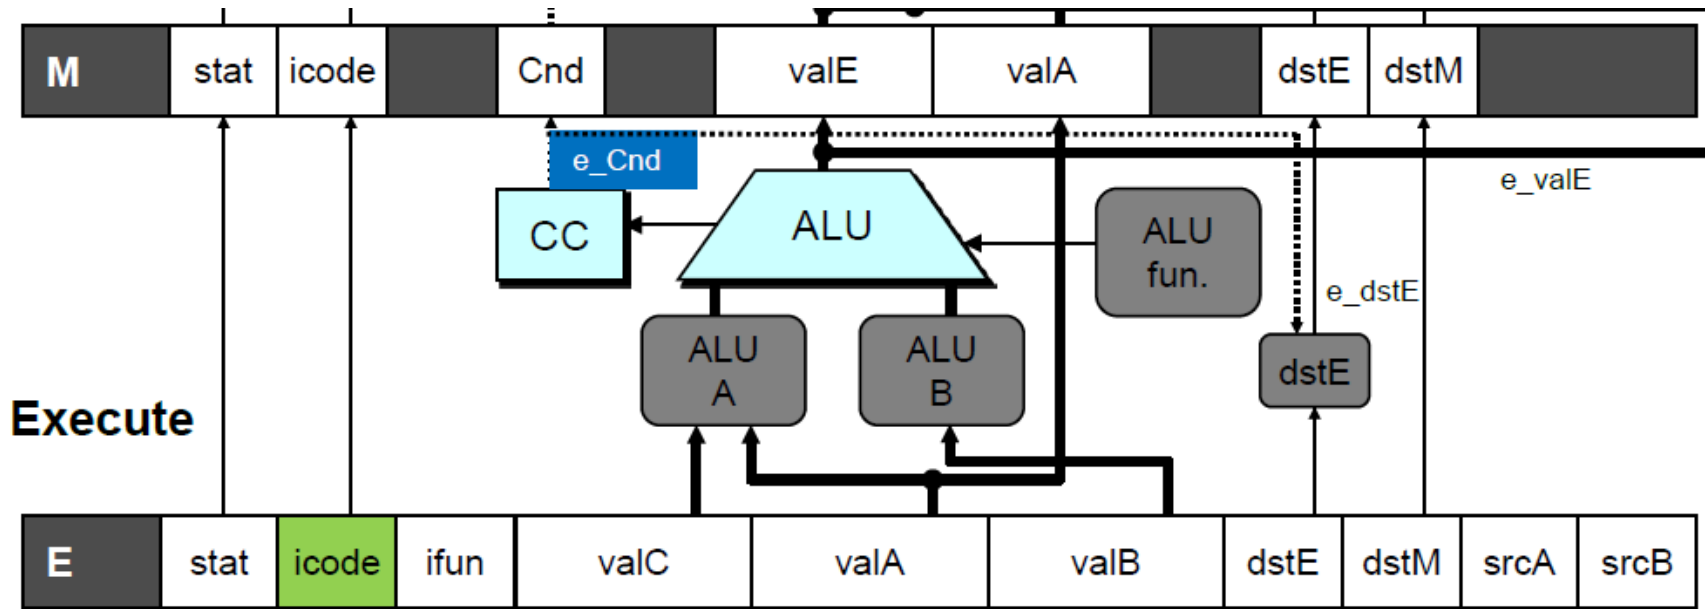
Predict branch as taken

- Fetch 2 instructions at target

Cancel when mispredicted

- Detect branch not-taken in execute stage
- On following cycle, replace instructions in execute and decode by bubbles
- Key point: No side effects have occurred yet

Detecting Mispredicted Branch



Condition	Trigger
Mispredicted Branch	$E_icode = IJXX \ \& \ !e_Cnd$

Control for Misprediction

demo-j.js

0x000: xorl %eax,%eax

0x002: jne target # Not taken

0x011: t: irmovl \$2,%edx # Target

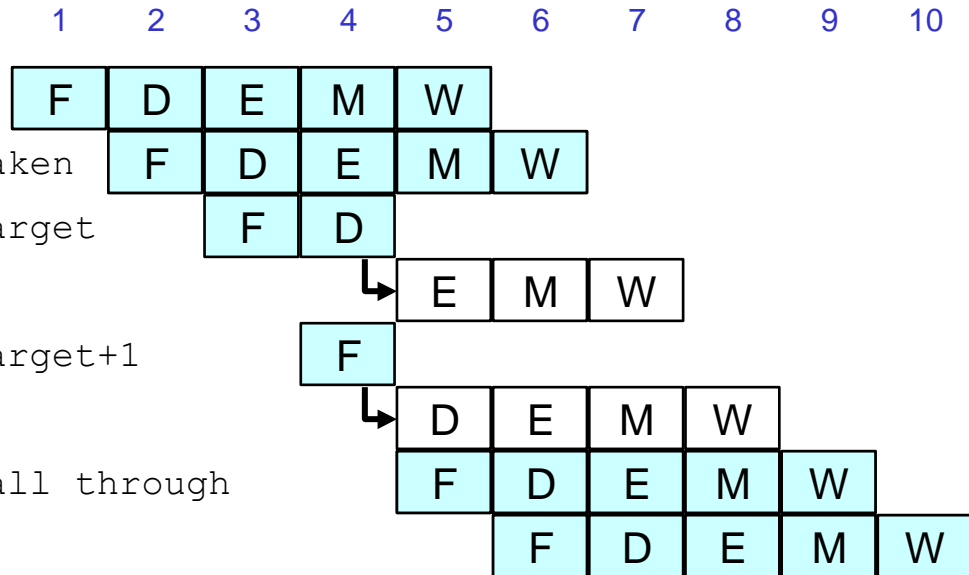
bubble

0x017: irmovl \$3,%ebx # Target+1

bubble

0x007: irmovl \$1,%eax # Fall through

0x00d: nop



Condition	F	D	E	M	W
Mispredicted Branch	normal	bubble	bubble	normal	normal

demo-retb.ys

Return Example

```

0x000:      irmovl Stack,%esp    # Initialize stack pointer
0x006:      call p              # Procedure call
0x00b:      irmovl $5,%esi      # Return point
0x011:      halt
0x020:      .pos 0x20
0x020: p:   irmovl $-1,%edi      # procedure
0x026:      ret
0x027:      irmovl $1,%eax       # Should not be executed
0x02d:      irmovl $2,%ecx       # Should not be executed
0x033:      irmovl $3,%edx       # Should not be executed
0x039:      irmovl $4,%ebx       # Should not be executed
0x100:      .pos 0x100
0x100:      Stack:              # Stack: Stack pointer

```

- Previously executed three additional instructions

Correct Return Example

```
# demo- retb
```

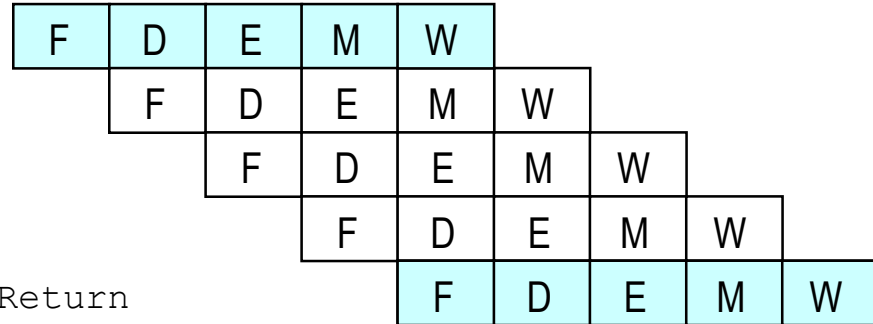
```
0x026:    ret
```

```
    bubble
```

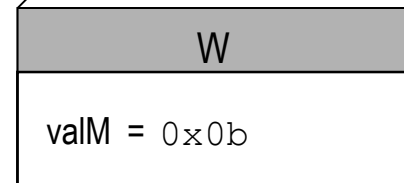
```
    bubble
```

```
    bubble
```

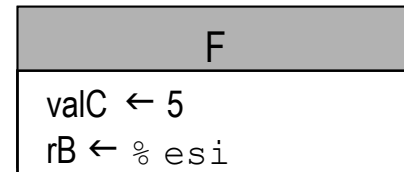
```
0x00b:    irmovl $5,%esi # Return
```



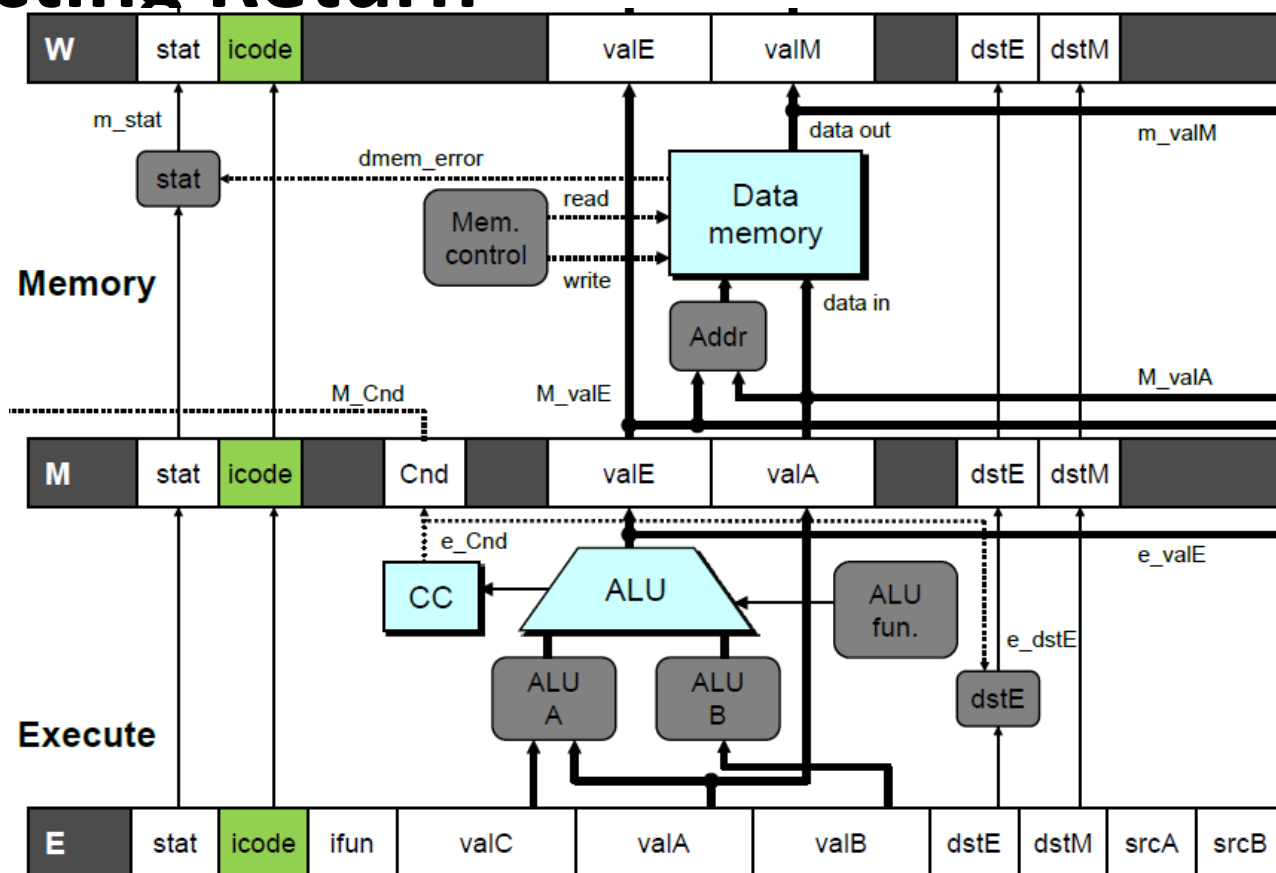
- As `ret` passes through pipeline, stall at fetch stage
 - While in decode, execute, and memory stage
- Inject bubble into decode stage
- Release stall when reach write-back stage



⋮



Detecting Return



Condition	Trigger
Processing ret	IRET in { D_icode, E_icode, M_icode }

Control for Return

demo-retb

0x026: ret

bubble

bubble

bubble

0x00b: irmovl \$5,%esi # Return

F	D	E	M	W					
	F	D	E	M	W				
		F	D	E	M	W			
			F	D	E	M	W		
				F	D	E	M	W	

Condition	F	D	E	M	W
Processing ret	stall	bubble	normal	normal	normal

Special Control Cases

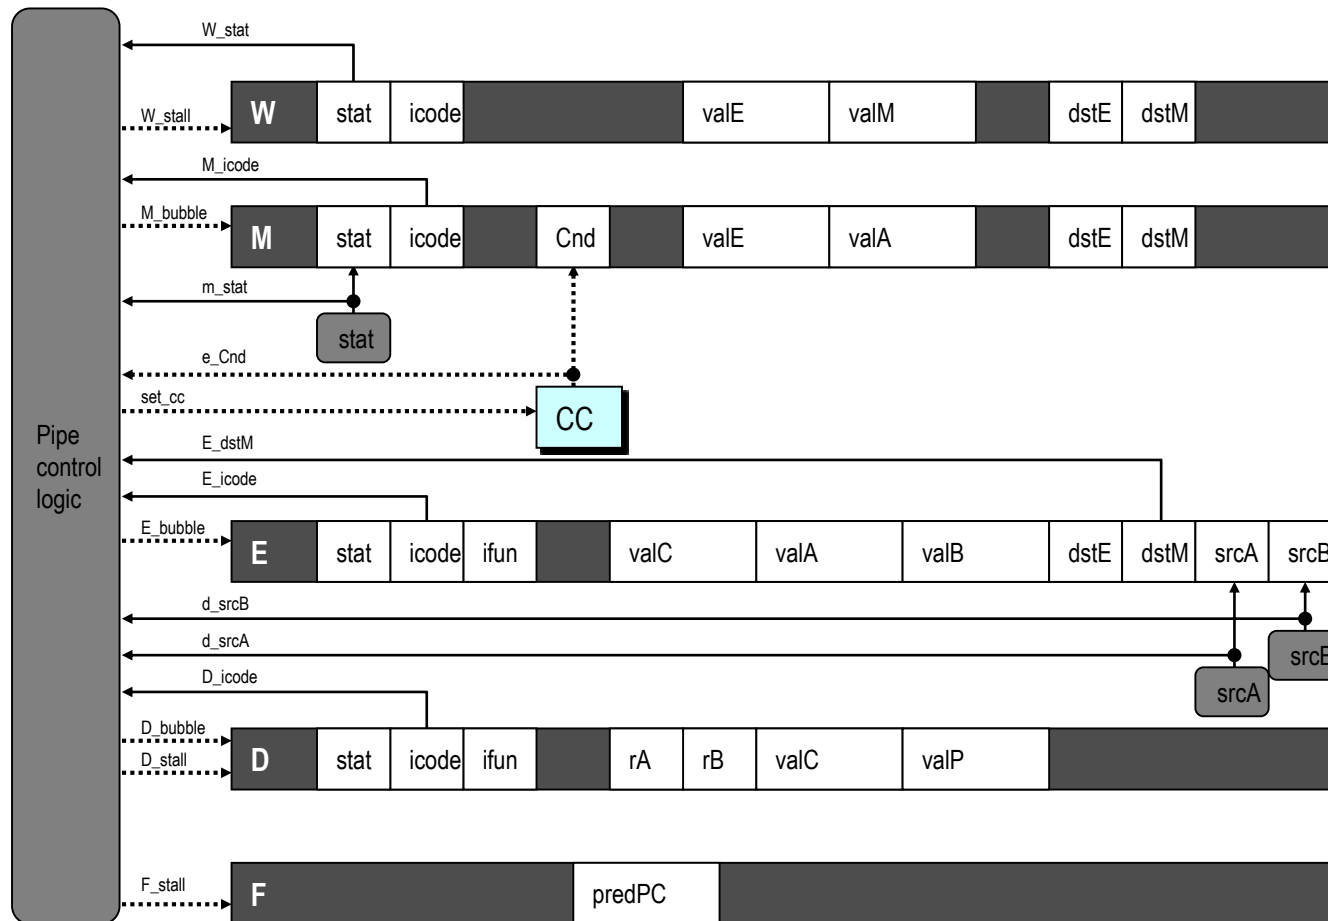
■ Detection

Condition	Trigger
Processing ret	IRET in { D_icode, E_icode, M_icode }
Load/Use Hazard	E_icode in { IMRMOVL, IPOPL } && E_dstM in { d_srcA, d_srcB }
Mispredicted Branch	E_icode = IJXX & !e_Cnd

■ Action (on next cycle)

Condition	F	D	E	M	W
Processing ret	stall	bubble	normal	normal	normal
Load/Use Hazard	stall	stall	bubble	normal	normal
Mispredicted Branch	normal	bubble	bubble	normal	normal

Implementing Pipeline Control



- Combinational logic generates pipeline control signals
- Action occurs at start of following cycle

Initial Version of Pipeline Control

```

bool F_stall =
    # Conditions for a load/use hazard
    E_icode in { IMRMOVL, IPOPL } && E_dstM in { d_srcA, d_srcB } ||
    # Stalling at fetch while ret passes through pipeline
    IRET in { D_icode, E_icode, M_icode };

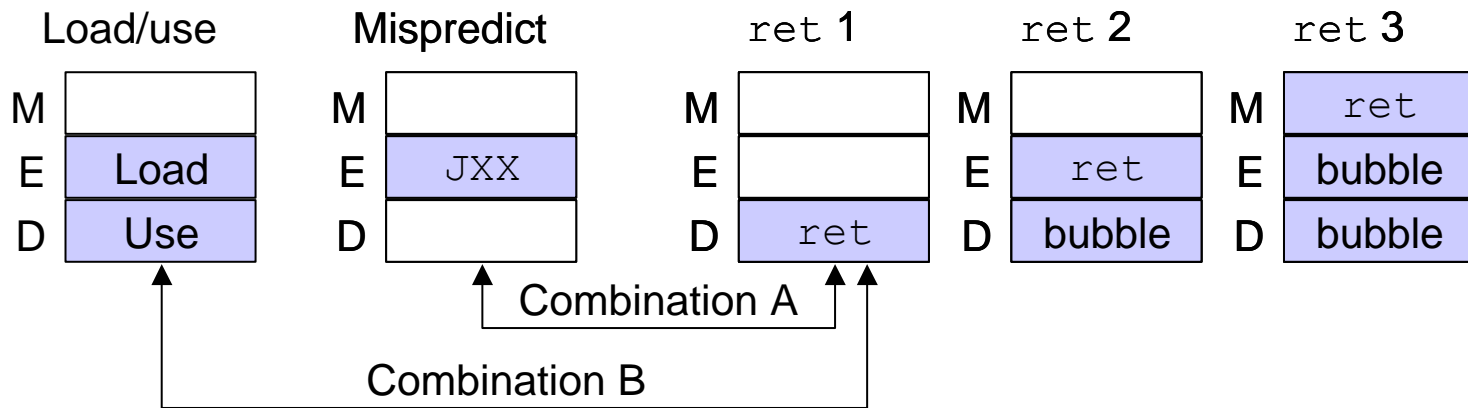
bool D_stall =
    # Conditions for a load/use hazard
    E_icode in { IMRMOVL, IPOPL } && E_dstM in { d_srcA, d_srcB };

bool D_bubble =
    # Mispredicted branch
    (E_icode == IJXX && !e_Cnd) ||
    # Stalling at fetch while ret passes through pipeline
    IRET in { D_icode, E_icode, M_icode };

bool E_bubble =
    # Mispredicted branch
    (E_icode == IJXX && !e_Cnd) ||
    # Load/use hazard
    E_icode in { IMRMOVL, IPOPL } && E_dstM in { d_srcA, d_srcB };

```

Control Combinations



- Special cases that can arise on same clock cycle

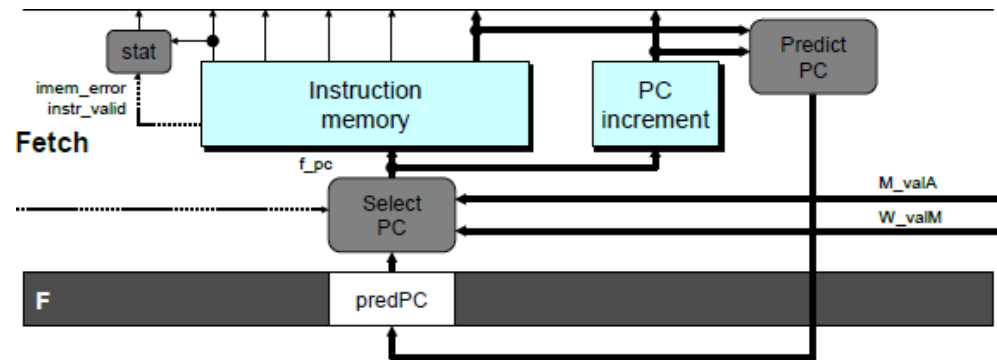
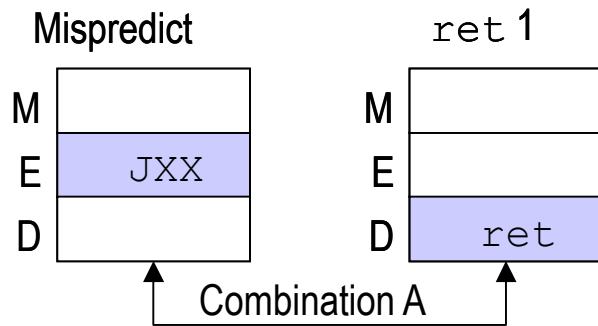
■ Combination A

- Not-taken branch
- `ret` instruction at branch target

■ Combination B

- Instruction that reads from memory to `%esp`
- Followed by `ret` instruction

Control Combination A



Condition	F	D	E	M	W
Processing ret	stall	bubble	normal	normal	normal
Mispredicted Branch	normal	bubble	bubble	normal	normal
Combination	<i>stall</i>	<i>bubble</i>	<i>bubble</i>	<i>normal</i>	<i>normal</i>

- Should handle as mispredicted branch
- Stalls F pipeline register
- But PC selection logic will be using M_valM anyhow

Control Combination B



Condition	F	D	E	M	W
Processing ret	stall	bubble	normal	normal	normal
Load/Use Hazard	stall	stall	bubble	normal	normal
Combination	<i>stall</i>	<i>bubble + stall</i>	<i>bubble</i>	<i>normal</i>	<i>normal</i>

- Would attempt to bubble *and* stall pipeline register D
- Signaled by processor as pipeline error

Handling Control Combination B



Condition	F	D	E	M	W
Processing ret	stall	bubble	normal	normal	normal
Load/Use Hazard	stall	stall	bubble	normal	normal
<i>Combination</i>	<i>stall</i>	<i>stall</i>	<i>bubble</i>	<i>normal</i>	<i>normal</i>

- Load/use hazard should get priority
- `ret` instruction should be held in decode stage for additional cycle

Corrected Pipeline Control Logic

```

bool D_bubble =
    # Mispredicted branch
    (E_icode == IJXX && !e_Cnd) ||
    # Stalling at fetch while ret passes through pipeline
    IRET in { D_icode, E_icode, M_icode }
    # but not condition for a load/use hazard
    && !(E_icode in { IMRMOVL, IPOPL }
        && E_dstM in { d_srcA, d_srcB }));
  
```

Condition	F	D	E	M	W
Processing ret	stall	bubble	normal	normal	normal
Load/Use Hazard	stall	stall	bubble	normal	normal
<i>Combination</i>	<i>stall</i>	<i>stall</i>	<i>bubble</i>	<i>normal</i>	<i>normal</i>

- Load/use hazard should get priority
- `ret` instruction should be held in decode stage for additional cycle

Pipeline Part 2: Summary

■ Data Hazards

- Most handled by forwarding
 - No performance penalty
- Load/use hazard requires one cycle stall

■ Control Hazards

- Cancel instructions when detect mispredicted branch
 - Two clock cycles wasted
- Stall fetch stage while `ret` passes through pipeline
 - Three clock cycles wasted

■ Control Combinations

- Must analyze carefully
- First version had subtle bug
 - Only arises with unusual instruction combination