

Dynamic Memory Allocation: Garbage Collectors

Introduction to Computer Systems

Today

■ Garbage collection

- Basic concepts
- Mark and Sweep
- Reference Counting
- Copying
- Treadmill
- Generational

■ Memory-related perils and pitfalls

Implicit Memory Management: Garbage Collection

- ***Garbage collection***: automatic reclamation of heap-allocated storage—application never has to free

```
void foo() {  
    int *p = malloc(128);  
    return; /* p block is now garbage */  
}
```

- **Common in many dynamic languages:**
 - Python, Ruby, Java, Perl, ML, Lisp, Mathematica
- **Variants (“conservative” garbage collectors) exist for C and C++**
 - However, cannot necessarily collect all garbage

Garbage Collection

- **How does the memory manager know when memory can be freed?**
 - In general we cannot know what is going to be used in the future since it depends on conditionals
 - But we can tell that certain blocks cannot be used if there are no pointers to them

- **Must make certain assumptions about pointers**
 - Memory manager can distinguish pointers from non-pointers
 - All pointers point to the start of a block
 - Cannot hide pointers
(e.g., by coercing them to an `int`, and then back again)

Classical GC Algorithms

■ Basic GC

- Mark-and-sweep collection (McCarthy, 1960)
 - Does not move blocks (unless you also “compact”)
- Reference counting (Collins, 1960)
 - Does not move blocks
- Copying collection (Minsky, 1963)
 - Moves blocks

■ Advanced GC


- Tricolor Marking and The Treadmill
- Generational Collectors (Lieberman and Hewitt, 1983)
 - Collection based on lifetimes
 - Most allocations become garbage very soon
 - So focus reclamation work on zones of memory recently allocated

- For more information: Jones and Lin, “*Garbage Collection: Algorithms for Automatic Dynamic Memory*”, John Wiley & Sons, 1996.

Memory as a Graph

■ We view memory as a directed graph

- Each block is a node in the graph
- Each pointer is an edge in the graph
- Locations not in the heap that contain pointers into the heap are called **root** nodes (e.g. registers, locations on the stack, global variables)



A node (block) is **reachable** if there is a path from any root to that node.

Non-reachable nodes are **garbage** (cannot be needed by the application)

Today


■ Garbage collection

- Basic concepts
- **Mark and Sweep**
- Reference Counting
- Copying
- Treadmill
- Generational

■ Memory-related perils and pitfalls


Mark and Sweep Collecting

- Can build on top of malloc/free package
 - Allocate using `malloc` until you “run out of space”
- When out of space:
 - Use extra **mark bit** in the head of each block
 - **Mark**: Start at roots and set mark bit on each reachable block
 - **Sweep**: Scan all blocks and free blocks that are not marked




Mark-phase


Mark all objects as “unused”



Mark the objects pointed by root set as “used”

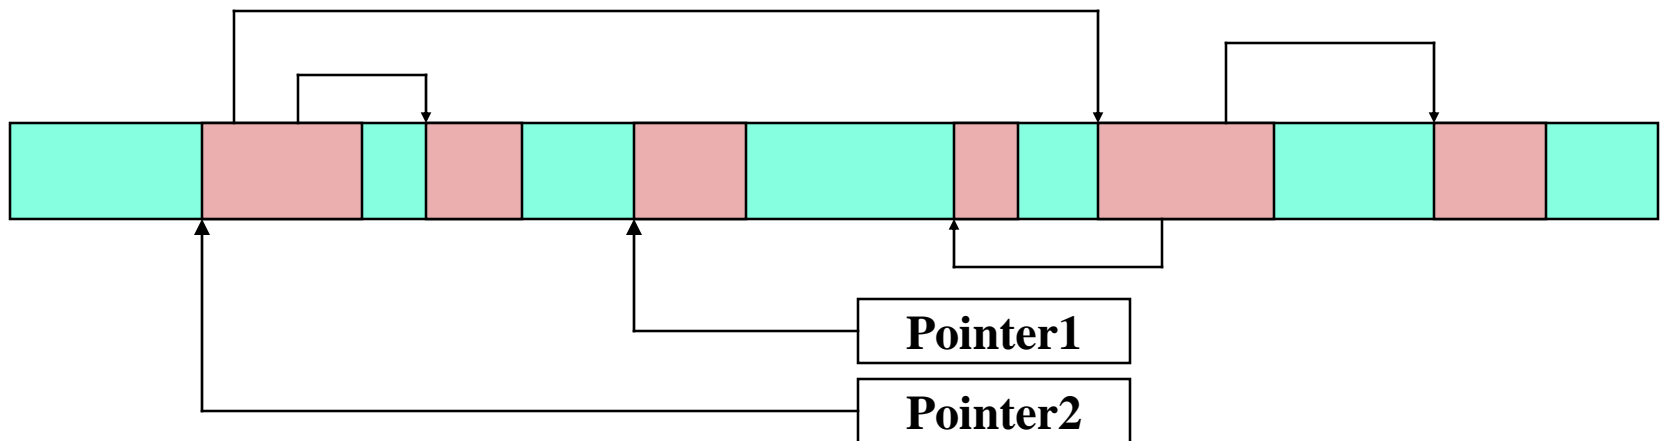


Mark the objects reached by root set as “used”




Mark-phase and Sweep-phase

Repeat this marking process until no new objects be marked



- Scan all heap space
- Reclaim the objects marked as “unused”



Assumptions For a Simple Implementation

■ Application

- `new(n)`: returns pointer to new block with all locations cleared
- `read(b,i)`: read location `i` of block `b` into register
- `write(b,i,v)`: write `v` into location `i` of block `b`

■ Each block will have a header word

- addressed as `b[-1]`, for a block `b`
- Used for different purposes in different collectors

■ Instructions used by the Garbage Collector

- `is_ptr(p)`: determines whether `p` is a pointer
- `length(b)`: returns the length of block `b`, not including the header
- `get_roots()`: returns all the roots

Mark and Sweep (cont.)

Mark using depth-first traversal of the memory graph

```
ptr mark(ptr p) {  
    if (!is_ptr(p)) return;           // do nothing if not pointer  
    if (markBitSet(p)) return;       // check if already marked  
    setMarkBit(p);                   // set the mark bit  
    for (i=0; i < length(p); i++)    // call mark on all words  
        mark(p[i]);                 // in the block  
    return;  
}
```


Sweep using lengths to find next block

```
ptr sweep(ptr p, ptr end) {  
    while (p < end) {  
        if markBitSet(p)  
            clearMarkBit();  
        else if (allocateBitSet(p))  
            free(p);  
        p += length(p);  
    }  
}
```

Conservative Mark & Sweep in C


■ A “conservative garbage collector” for C programs

- `is_ptr()` determines if a word is a pointer by checking if it points to an allocated block of memory
- But, in C pointers can point to the middle of a block



■ So how to find the beginning of the block?

- Can use a balanced binary tree to keep track of all allocated blocks (key is start-of-block)
- Balanced-tree pointers can be stored in header (use two additional words)



Left: smaller addresses
Right: larger addresses

Mark-compact & Incremental Mark-sweep

■ Mark-compact

- Two phases
 - Mark-phase: just like the mark-sweep collector
 - Compact-phase: move all the live objects to one end and leave all the garbage to the other
- Eliminate the external fragmentation and improve spatial locality

■ Incremental Mark-sweep

- Application processing may continue while garbage collection is taking place

Today

■ Garbage collection


- Basic concepts
- Mark and Sweep
- **Reference Counting**
- Copying
- Treadmill
- Generational

■ Memory-related perils and pitfalls

Reference Counting


- **Each object has an associated counter of the references**
 - When the object is first created, the counter is set to one
 - Whenever a reference is copied, the counter is incremented
 - Whenever a reference is overwritten, the counter is decremented
 - When the counter reaches zero, the object becomes garbage

Example of RC



- Create a new object
`a = new Object();`
- The counter is incremented when reference is copied
`b = a;`
- The counter is incremented when another pointer points to the object
`c.something = a;`

Example of RC (cont.)



- The counter is decremented when one reference is dead

```
{
    Object a, b;
    .....
    a = new Object();
    b = a;
    .....
} // when program executed to here
```

- The counter is decremented when the pointer points to another object
 - `c = some_thing_else;`
- The object will be reclaimed when its counter reaches zero
- The pointers in the reclaimed objects should be taken into account
 - Recursive reclamation

Discussion and Deferred RC

■ Discussion

- Incremental nature
- The problem with cycles
- Usually used between distributed nodes

■ Deferred Reference Counting

- Short-lived objects make considerable cost

Today

■ Garbage collection


- Basic concepts
- Mark and Sweep
- Reference Counting
- **Copying**
- Treadmill
- Generational

■ Memory-related perils and pitfalls


Copying

- **Originally “stop-and-copy” scheme**
- **The heap is divided into two contiguous semi-spaces:**
 - From-space: initially empty
 - To-space: allocation area
- **If the allocation requirement can't be satisfied**
 - Exchange from-space and to-space
 - Copy all reachable objects from from-space into one end of to-space
 - Begin new allocations in new to-space

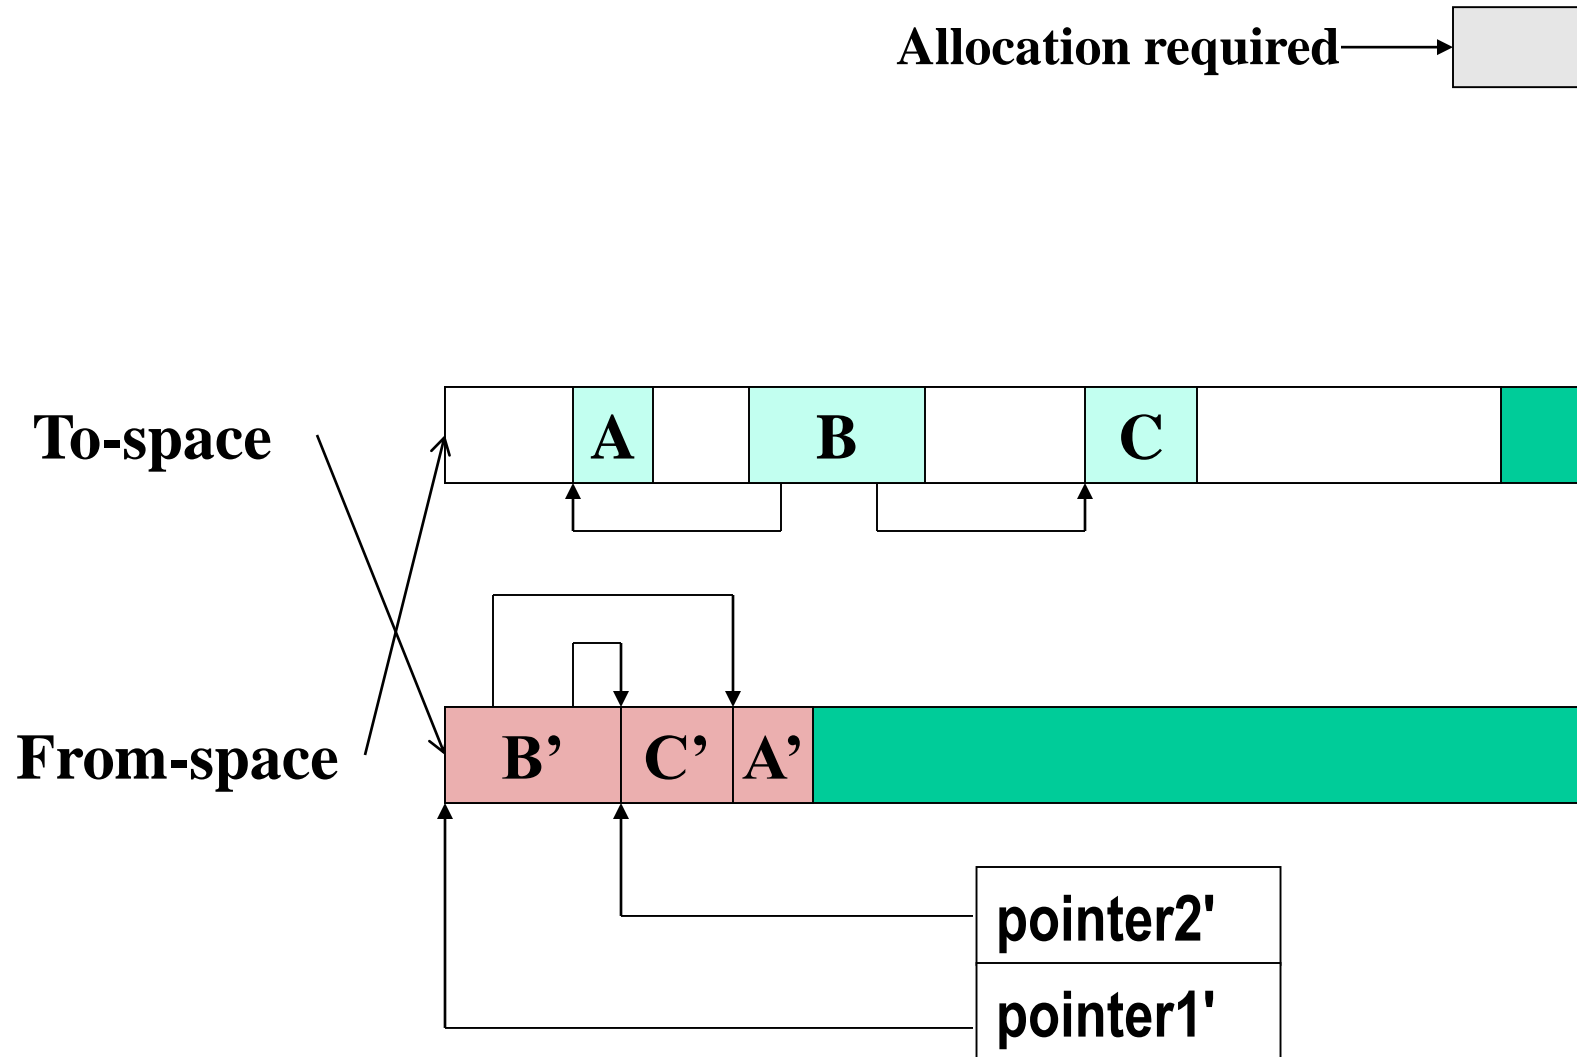
Example of Copying



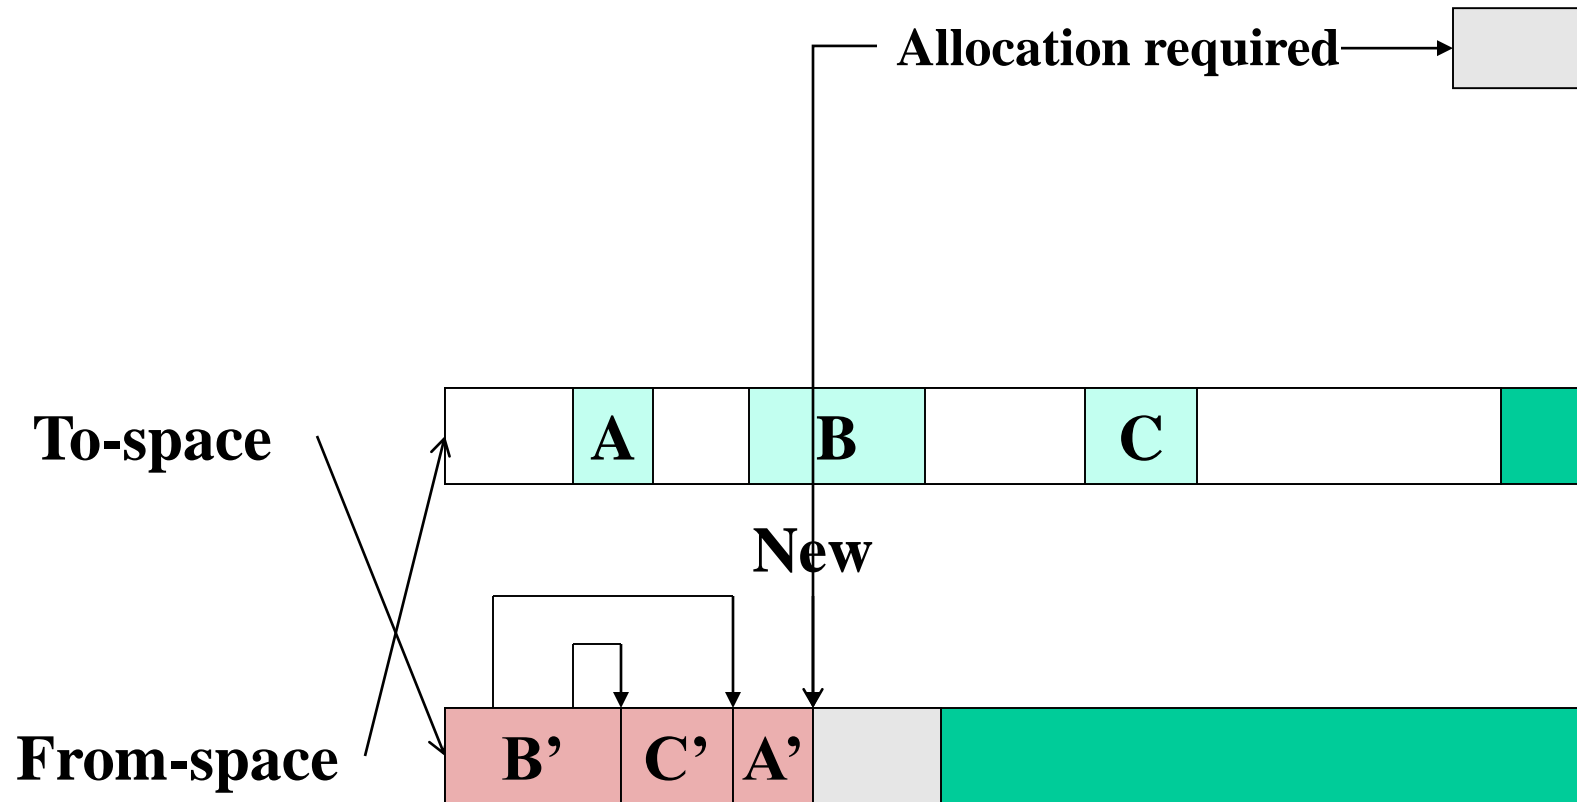
Example of Copying



Example of Copying



Example of Copying



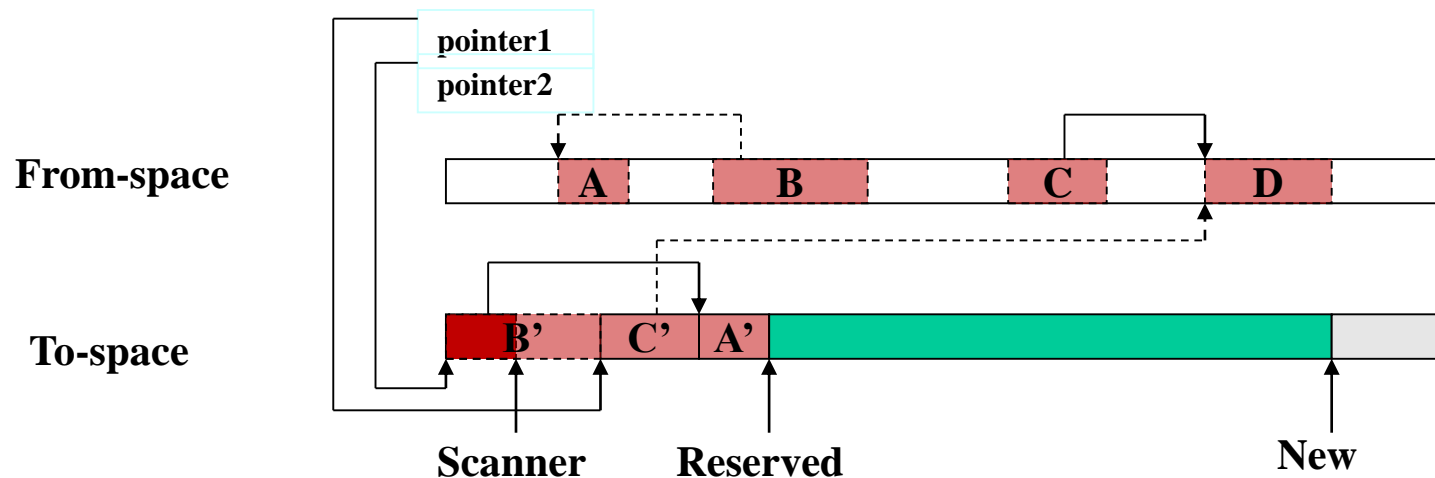
Discussion and Variants

■ Discussion about Copying Collector

- Stop-and-wait method suspends applications
- Difficult support conservative collection
- Half used, half wasted

■ Baker's Real-time Copying

- Incrementally copying all live objects out of from-space into to-space while new memory is allocated from to-space



Today

■ Garbage collection

- Basic concepts
- Mark and Sweep
- Reference Counting
- Copying
- **Treadmill**
- Generational

■ Memory-related perils and pitfalls

Tricolor Marking

■ Meaning of Incrementality:

- Incremental garbage collectors allow application processing to continue while garbage collection is performed

■ Abstraction of incremental GC: Tricolor scheme


- White objects -- garbage
- Black objects -- reachable objects that already processed
- Gray objects -- reachable objects that haven't been processed

■ For the incremental reason


- the black objects and gray objects may not be reachable objects
- the white objects may not be garbage

Principle of Tricolor Marking

- Initially, all objects are assumed to be garbage and are colored white
 - The purpose: identify all living objects and color them black
- The traversal phase proceeds in a wave-front of grey objects which separates the white objects from the black ones.
- The abstraction of tricolor marking is helpful in understanding incremental tracing phase of garbage collector.





Before



After


The Treadmill

- **A New Scan: All are white**
 - Free-list: free blocks
 - Available for allocation
 - From-list: allocated blocks
 - All objects are white (garbage)
- **During Scanning:**
 - New-list: All black
 - Newly allocated during GC
 - To-list: Black or grey
 - Reachable objects
 - From-list: All white
 - May not be garbage
 - Free-list: All white



The Treadmill

- **When to complete:** *No grey objects in To-list*
- **All are white again**
 - New-From-list = Old-To-list + Old-New-list
 - New-Free-list = Old-Free-list + Old-From-list
 - New New-list: empty
 - New To-list: empty



Today

■ Garbage collection

- Basic concepts
- Mark and Sweep
- Reference Counting
- Copying
- Treadmill
- **Generational**

■ Memory-related perils and pitfalls

Generational GC

■ Infant mortality

- Most objects live a very short time, while a small percentage of them tend to live much longer
- 80% to 98% of newly-allocated objects die within a few million instructions

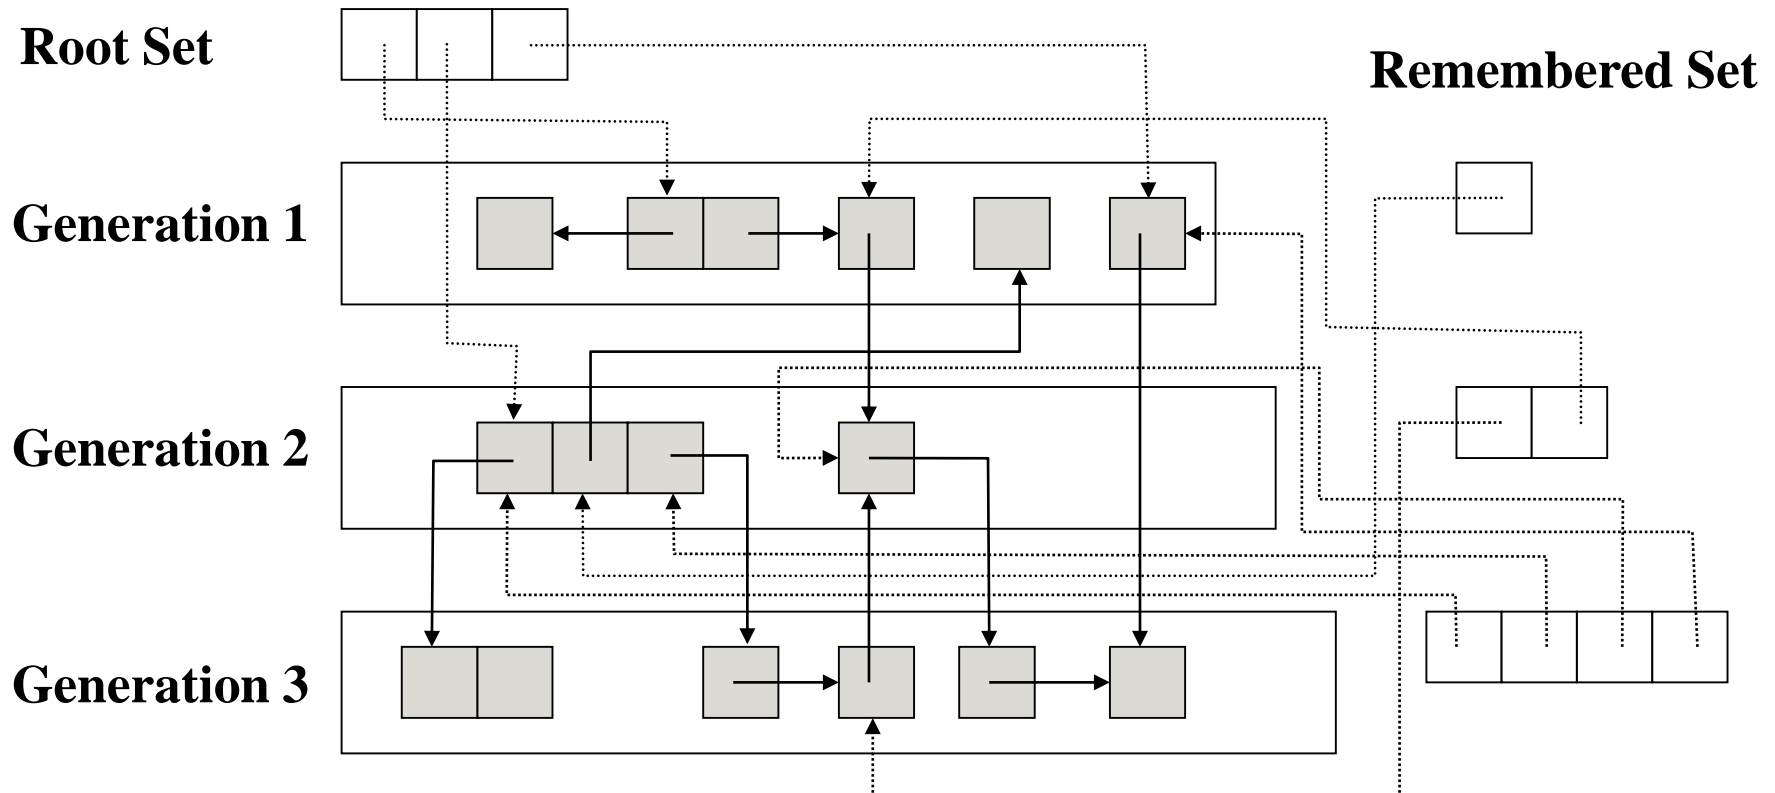
■ Generations

- Each holding objects of different age categories
 - New objects are allocated in the youngest generation
 - Whenever an object has lived for some period of time, it is moved to the next generation
- Each designed to be collected individually

Example of Generational GC

■ Inter-generational references

- *Remembered set*: containing the identity of all references into the generation from heap objects residing in other generations
- *Directionality of reference*: Newer objects tend to point to Older objects



Train Algorithm

- An incremental garbage collection scheme for achieving non-disruptive reclamation of the oldest generational area, the mature object space.
- The algorithm achieves its incrementality by dividing mature object space into a number of fixed-sized blocks and collecting one block at each invocation.
 - The blocks are referred as cars
 - The set of blocks to which a car belongs as its train
 - Mature object space can then be thought of as a giant railway station with trains lined up on its tracks

Today

■ Garbage collection

- Basic concepts
- Mark and Sweep
- Reference Counting
- Copying
- Treadmill
- Generational

■ Memory-related perils and pitfalls

Memory-Related Perils and Pitfalls

- Dereferencing bad pointers
- Reading uninitialized memory
- Overwriting memory
- Referencing nonexistent variables
- Freeing blocks multiple times
- Referencing freed blocks
- Failing to free blocks

C operators

Operators

```
( )  [ ]  ->  .
!  ~  ++  --  +  -  *  &  (type)  sizeof
*  /  %
+  -
<<  >>
<  <=  >  >=
==  !=
&
^
|
&&
||
?:
=  +=  -=  *=  /=  %=  &=  ^=  !=  <<=  >>=
,
```

Associativity

```
left to right
right to left
left to right
left to right
left to right
left to right
left to right
left to right
left to right
left to right
right to left
right to left
left to right
```

- `->`, `()`, and `[]` have high precedence, with `*` and `&` just below
- Unary `+`, `-`, and `*` have higher precedence than binary forms

C Pointer Declarations: Test Yourself!

<code>int *p</code>	p is a pointer to int
<code>int *p[13]</code>	p is an array[13] of pointer to int
<code>int *(p[13])</code>	p is an array[13] of pointer to int
<code>int **p</code>	p is a pointer to a pointer to an int
<code>int (*p)[13]</code>	p is a pointer to an array[13] of int
<code>int *f()</code>	f is a function returning a pointer to int
<code>int (*f)()</code>	f is a pointer to a function returning int
<code>int (*(*f()) [13])()</code>	f is a function returning ptr to an array[13] of pointers to functions returning int
<code>int (*(*x[3])()) [5]</code>	x is an array[3] of pointers to functions returning pointers to array[5] of ints

Dereferencing Bad Pointers

■ The classic scanf bug

```
int val;  
  
...  
  
scanf("%d", val);
```


Reading Uninitialized Memory

- Assuming that heap data is initialized to zero

```
/* return y = Ax */  
int *matvec(int **A, int *x) {  
    int *y = malloc(N*sizeof(int));  
    int i, j;  
  
    for (i=0; i<N; i++)  
        for (j=0; j<N; j++)  
            y[i] += A[i][j]*x[j];  
    return y;  
}
```

Overwriting Memory

- Allocating the (possibly) wrong sized object

```
int **p;  
  
p = malloc(N*sizeof(int));  
  
for (i=0; i<N; i++) {  
    p[i] = malloc(M*sizeof(int));  
}
```

Overwriting Memory

■ Off-by-one error

```
int **p;  
  
p = malloc(N*sizeof(int *));  
  
for (i=0; i<=N; i++) {  
    p[i] = malloc(M*sizeof(int));  
}
```

Overwriting Memory

- Not checking the max string size

```
char s[8];  
int i;  
  
gets(s);  /* reads "123456789" from stdin */
```

- Basis for classic buffer overflow attacks

Overwriting Memory

■ Misunderstanding pointer arithmetic

```
int *search(int *p, int val) {  
    while (*p && *p != val)  
        p += sizeof(int);  
  
    return p;  
}
```

Overwriting Memory

- Referencing a pointer instead of the object it points to

```
int *BinheapDelete(int **binheap, int *size) {  
    int *packet;  
    packet = binheap[0];  
    binheap[0] = binheap[*size - 1];  
    *size--;  
    Heapify(binheap, *size, 0);  
    return(packet);  
}
```

Referencing Nonexistent Variables

- Forgetting that local variables disappear when a function returns

```
int *foo () {  
    int val;  
  
    return &val;  
}
```

Freeing Blocks Multiple Times

■ Nasty!

```
x = malloc(N*sizeof(int));  
    <manipulate x>  
free(x);  
  
y = malloc(M*sizeof(int));  
    <manipulate y>  
free(x);
```


Referencing Freed Blocks

■ Evil!

```
x = malloc(N*sizeof(int));  
  <manipulate x>  
free(x);  
  ...  
y = malloc(M*sizeof(int));  
for (i=0; i<M; i++)  
    y[i] = x[i]++;
```

Failing to Free Blocks (Memory Leaks)

- Slow, long-term killer!

```
foo() {  
    int *x = malloc(N*sizeof(int));  
    ...  
    return;  
}
```

Failing to Free Blocks (Memory Leaks)

■ Freeing only part of a data structure

```
struct list {
    int val;
    struct list *next;
};

foo() {
    struct list *head = malloc(sizeof(struct list));
    head->val = 0;
    head->next = NULL;
    <create and manipulate the rest of the list>
    ...
    free(head);
    return;
}
```

Dealing With Memory Bugs

■ Debugger: `gdb`

- Good for finding bad pointer dereferences
- Hard to detect the other memory bugs

■ Data structure consistency checker

- Runs silently, prints message only on error
- Use as a probe to zero in on error

■ Binary translator: `valgrind`

- Powerful debugging and analysis technique
- Rewrites text section of executable object file
- Checks each individual reference at runtime
 - Bad pointers, overwrites, refs outside of allocated block

■ `glibc malloc` contains checking code

- `setenv MALLOC_CHECK_ 3`